Seldon Manifolds — Social and Psychological Application

Contact Comfort, Delay of Gratification, and What a Substrate Produces Without Being Fitted To It

Matthew Parslow
Independent Researcher

Draft — draft-2026-08-30.4
1 September 2026 • document revision 1

DRAFT — draft-2026-08-30.4 r1 (1 September 2026). Circulated for comment, not as a finished result. This is the applied companion to *Seldon Manifolds — Psychological Grounding.* That paper asks whether a nervous system implements anything of the manifold's shape; this one reports what happens when the substrate is run instead, and what the outputs do and do not license. Sections marked as falsifiability commitments are the intended points of attack.

Abstract

Two classical experiments were put to a substrate built for other reasons. Contact comfort (Harlow, 1958) reproduces, with an ablation control establishing which structure carries it. The delay-of-gratification task does something more useful than reproduce: the substrate exhibits, constructively, why the measure is confounded — identical dispositions produce different wait times under different beliefs about reliability, and dispositions differing fourfold produce identical wait times under compensating beliefs. That agrees with what the empirical literature already found (Kidd et al., 2013; Watts et al., 2018), which is the point rather than a discovery: an account that had to be told this would be worth less than one in which it falls out. *Nothing here is evidence about children or monkeys.* It is evidence about what this substrate does, and the paper is careful throughout about which of the two is being claimed.

1 What this paper is, and what it is not

The theory paper says what a Seldon manifold is; the grounding companion asks whether a biological nervous system implements anything of that shape and answers in pieces. This paper runs the third question, which is the one the substrate can settle on its own: put the substrate in the position of a classical experimental subject and report what it does.

The results below are about a substrate and not about people. A simulation that reproduces an experimental finding has established something about the simulation. It becomes evidence about the world only through an argument this paper does not make and the grounding companion makes only in pieces. What a reproduction does establish is narrower and still worth having: that the finding is *available* from the substrate's own machinery without being put there, which is a fact about what that machinery entails.

Two disciplines are enforced by the runs themselves rather than promised here. Each records the commit it ran against and the seed it ran with, so a claim below is either reproducible or withdrawn. And each carries a *source audit* — an automated check that the decisions were taken by graph rules rather than by the host language, reported with the run and reproduced in the results tables.

2 Contact comfort

Harlow (1958) found infant macaques preferring a cloth surrogate offering no food to a wire surrogate offering it, and fleeing to the cloth under stress. The substrate was given a contact drive and the same two-surrogate arrangement.

$n = 200$ per arm	**contact drive present**	**ablated**
cloth-directed fraction	0.797 [0.792, 0.803]	0.000 [0.000, 0.000]
flee-to-cloth under stress	1.000 [1.000, 1.000]	0.000 [0.000, 0.000]
infants preferring cloth	200/200	0/200

The intervals are separated and the preference is above half, which is the finding. The second column is what makes it worth reporting: *ablating the contact drive removes the preference entirely rather than reducing it*, so the effect is carried by that structure and not by the arrangement. A glass-box readout taken from the graph rather than from behaviour gives cloth proximity 53 against wire 8 at equal bond strength.

Remark 2.1 (The ethics gate, and why a simulation needs one). Harlow's procedure is among the most criticised in the discipline's history, and a substrate that reproduces it is running a deprivation protocol on something. The run therefore carries a gate: the stressed state must be restored *byte-identically* afterwards, verified by comparison of the serialised graph. It is: arousal moves $63 \to 347 \to 63$ across deprivation and restoration with the restoration non-oriented, and the graph returns byte-identical.

This is offered as a discipline rather than as a settlement. *Byte-identical restoration establishes that nothing persisted, not that nothing happened.* Whether a state that is entered and perfectly reversed is thereby of no account is a question this paper raises and does not answer, and a reader who thinks the gate insufficient is disagreeing with something the paper also does not claim to have settled.

3 Delay of gratification

The delay task is the more interesting case, because the substrate does not merely reproduce it.

The arrangement is a factorial grid: three dispositions (a drive weight of 2, 4, 8) crossed with three beliefs about how reliable the promised second reward is (60, 80, 100 on a scale of 100), $n = 60$ per cell, grid fixed in advance of the run.

Remark 3.1 (Both factors are real, and each has a null control). Holding the belief distribution fixed and varying disposition alone gives mean waits of 12.05 [10.40, 13.67] against 4.13 [3.68, 4.58], intervals separated. Holding disposition fixed and varying belief alone gives 2.92 [2.60, 3.27] against 13.23 [11.63, 14.97], separated.

Each has a matched null control in which the factor is held *constant* and the intervals are required to overlap, which they do in both cases. A run in which the machinery manufactured an effect would fail there, and the controls are reported because a positive result without them says less than it appears to.

The measure is not identifiable for disposition, and the grid shows it in both directions. One disposition ($w = 4$) under three beliefs (60, 80, 100) waits 3, 8, 13 ticks: a naive score reads one subject as three levels of self-control. Three dispositions differing fourfold ($w = 2, 4, 8$) under compensating beliefs (60, 68, 85) wait 5, 5, 5: the same score reads three subjects as identical.

Remark 3.2 (Impatience as the rational output of a correct belief). The sharpest form holds disposition fixed at its *most patient* setting and varies only the belief. Under an unreliable prior it waits 5.12 [4.90, 5.32]; under a reliable one, 23.82 [23.60, 24.03]. The intervals are separated and the deterministic anchors are 5 and 25.

So a patient subject waits briefly, and does so because waiting briefly is what its belief recommends. *The short wait is attributable to the prior and not to a deficit*, and no disposition needs to be posited to explain it. What a naive reading scores as poor self-control is here a correct inference from an environment that does not keep its promises.

Remark 3.3 (This agrees with the literature, which is the point). Kidd et al. (2013) showed children's wait times moving with experimentally manipulated reliability of the environment, and Watts et al. (2018) found the long-run predictive power of the original measure much reduced once background is controlled. The finding above is theirs and is not claimed here.

What is claimed is smaller and different in kind. *A substrate assembled for unrelated reasons, carrying no term for self-control and no term for the task, produces the confound as a consequence of its ordinary machinery.* An account that had to be told about the confound would be worth less than one in which it falls out, and this is the second. That is a fact about the substrate's entailments and not additional evidence for the psychology, which already has better.

4 What the runs enforce

Both experiments carry automated gates whose results are reported with them, because a claim that the decisions were taken in the substrate rather than in the host language is not one a reader should take on trust.

The delay run records zero occurrences of the host rule-engine entry points, and zero occurrences of reinforcement vocabulary — no reward signal, no punishment, no reinforcement, no backpropagation, no gradient — so the waiting is not shaped. The decisions are named: valuation by two graph rules, choice by a third taking an argmax over candidate values, with the host confined to setting the disposition, supplying the belief, sweeping elapsed time and reading the chosen action. World constants are integers and are listed. The canonical seed set is checked untouched by digest.

None of this establishes that the model is right. It establishes that the model is the thing being reported, which is a different and lower bar, and one that a great deal of published simulation work does not clear.

5 The wider record

Two experiments are reported here in full. They are drawn from a larger body of pre-registered runs in the same substrate — at the time of writing, one hundred and five pre-registrations and seventy-one result documents, covering social conduct, institutional decision, bond formation and the dynamics of coercive relationships.

That record is not summarised here and should not be read as supporting anything above. *A result is reported in this programme when it has been written up against a pre-registration, not when it exists*, and the two below are the two that have. The remainder is named so that a reader knows the selection was made from something, and so that a later paper reporting more of it is understood as continuing this one rather than announcing a new result.

6 Falsifiability commitments

1. [DETAIL] **Contact preference is carried by the contact drive.** Fails if the preference survives ablation of that drive at comparable magnitude, in which case the arrangement rather than the drive is producing it.

2. [SHAPE] **The delay measure is not identifiable for disposition.** Fails if some assignment of beliefs recovers disposition uniquely from wait time in this substrate — that is, if the two directions exhibited above can both be blocked without adding structure for the purpose.

3. [SHAPE] **Impatience does not require a deficit.** Fails if a short wait cannot be produced from a patient disposition by belief alone, which is the arm reported in Remark 3.2 and would be settled against by its non-replication.

4. [DETAIL] **The waiting is not shaped.** Fails on any occurrence of reinforcement machinery in the path that produces the choice; the audit is reported with each run and can be re-run.

5. [SHAPE] **Reproduction is not evidence about organisms.** This paper commits to the negative: no claim below the level of the substrate is made, and a reader who finds one has found an error rather than a result.

References

Harry F. Harlow. The nature of love. *American Psychologist*, 13(12):673–685, 1958.

Celeste Kidd, Holly Palmeri, and Richard N. Aslin. Rational snacking: Young children's decision-making on the marshmallow task is moderated by beliefs about environmental reliability. *Cognition*, 126(1):109–114, 2013.

Tyler W. Watts, Greg J. Duncan, and Haonan Quan. Revisiting the marshmallow test: A conceptual replication investigating links between early delay of gratification and later outcomes. *Psychological Science*, 29(7):1159–1177, 2018.