Assessing an Account of This Kind

What a Universal Framework Can Risk, What It Earns Instead, and What the Record Shows

Matthew Parslow
Independent Researcher

Draft — draft-2026-08-30.4
30 August 2026 • document revision 25

DRAFT — draft-2026-08-30.4 r25 (30 August 2026). This document is under active revision and is circulated for comment, not as a finished result. Definitions, numbering and claims may change between drafts; cite the version string above if referring to it. Sections marked as falsifiability commitments are the intended points of attack.

Abstract

The Law of Graphic Equalisation states its commitments and marks each for whether failing it costs a mechanism or costs range. It does not argue about what kind of assessment such an account admits, and that argument had accumulated inside its falsifiability section until the section was a quarter of the paper. This companion carries it: what a framework claiming to hold of everything can and cannot risk, why usefulness rather than falsification is what it earns, why implementability supplies the surface on which it can fail, what the implementation record currently shows and does not, and how the choice of question bounds the value of any answer. It ends by holding itself to the same standard, with three commitments of its own, a paper arguing about what an account may be held to being subject to the same. Nothing here is required to read the substrate account. It is required to argue with it.

1 What a framework can and cannot risk

A list of falsifiers invites a reading this section should refuse before making it. *A universal framework that is only a vocabulary cannot be falsified, and that is not the same defect as being unfalsifiable* — the qualifier is load-bearing, and is discharged in Section 8, where the difference between a vocabulary and a claim with content turns out to be the whole question. Stated without it, the sentence is quotable in a form the rest of the section spends its length withdrawing. Something claimed to hold of everything has no outside from which a counterexample could arrive: every observation is expressible in it, so no observation contradicts it. That is a property of frameworks rather than a failing of this one — nobody demands a falsifier for arithmetic, or for the concept of a group — and the distinction is old enough to be named. Carnap (1950) separates questions asked *within* a linguistic framework, which have answers, from the question whether the framework itself is true, which he argues is not factual but practical and assessed by fruitfulness; Lakatos (1970) makes the same structure methodological, with a hard core protected by decision while an auxiliary belt takes the empirical hits.

This account can say something sharper than either, in its own terms. A falsifiable element is a *contingent* commitment. The tempting next step is that a contingent commitment carried by a framework is therefore structure doing no framework-work, so a maximally well-formed

framework has *no* falsifiable content — and that step is too quick, because this account's own write-economics is contingent and does a great deal of work. The claim that survives is narrower: a contingent commitment does no framework-work *where the framework would stand without it*, and such commitments are what the substrate paper's definition of well-formedness prices as idle. What remains falsifiable in a framework is accordingly the part that is load-bearing and could have been otherwise, which is a smaller set than the contingent one and is not empty — and because it is not empty, the strong reading does not hold either: *this account is not unfalsifiable-by-construction*, and an earlier version of this paragraph drew that conclusion after removing the premise supporting it. What survives is narrower. The part of a framework that is pure framework carries no falsifiable content, and that part is not a concession but what well-formedness looks like at this level; it is why the commitments the substrate paper identifies as derivations rather than wagers are the framework showing through rather than tests that could not be found. The rest of the account is load-bearing, could have been otherwise, and is answerable — which is what that paper's list of commitments is for.

That defence is defeasible and should not be left to do more than it can, since it does not distinguish a well-formed universal framework from a vacuous one — both are unfalsifiable. What separates them is available in the substrate paper's definition of well-formedness and requires reading both of its halves rather than one. Well-formedness is paying for nothing *that does no work*. A framework is held structure; by its own criterion it must therefore earn its retention by doing work, and one that affords nothing is inert structure at the top level — the fourth cost of that definition, applied to the account itself. *The framework earns its keep by being useful, not by being falsifiable*, and that is not a weaker discipline substituted for a stronger one. It is the discipline this account applies to everything else, turned on itself, and refusing to exempt itself from it is the whole of the argument.

Usefulness has a property falsification was meant to supply and cannot supply here: it is *auditable*. What a framework affords is a record rather than a promise — distinctions that do work the previous vocabulary could not do, disputes dissolved rather than adjudicated, results transferring between domains without rework, predictions that could have failed and did not. Each item is a specific claim, checkable by someone who does not accept the framework, which is what the falsifiers were for.

One limit is imposed by the substrate paper's standard and is not evaded. A test whose stopping condition lies with the assessor is not a test, and "is it useful enough?" has exactly that shape. So usefulness is offered as a *record* and never as a threshold: instances, enumerable and individually checkable, rather than a verdict on the total. And a secondary axis remains for the case the record does not settle, since two frameworks may both do work: there, the one covering the same ground with less structure wins, which is minimum description length rather than Popper and is comparative rather than empirical.

2 Where an unfalsifiable framework becomes falsifiable

One item of the record differs from the others in kind and completes the argument above. This framework is *implementable, forwards*, and has been implemented (the constructive companion). That is not a claim about elegance: the substrate is at its core an analysis-by-synthesis engine, in which recognition proceeds by generating candidates and comparing them, so generation and recognition are one operation run in two directions rather than two theories requiring separate justification — and where the two directions are exact inverses the round trip is checkable rather than asserted.

What that does and does not establish should be separated carefully, because the temptation is to claim too much from it. It is *not* evidence that the account is true: a framework can be perfectly implementable and describe nothing real, and by that standard an implementation built by the same hand shares vocabulary and personnel and is disqualified as independent

support exactly as the computing case is. Nor does it validate any particular design, since an implementation is free to add rules to suit and the framework does not determine one.

What it establishes is that the account is *specified enough to be built from*, which is a property most frameworks lack and which cannot be faked — one discovers that a framework is underspecified only by attempting to build it, which is the substrate paper's result that a boundary is known only by crossing it at the scale of the account itself. And it supplies the thing Section 1 conceded was missing. *The framework's purely framework-level part cannot be falsified; its contentful claims can, and any implementation of it can.* A running system makes the commitments concrete enough to break, and breaks, and is corrected — so the unfalsifiable core acquires a falsifiable surface, which is where the empirical discipline actually operates. Nothing here is unprecedented: the free-energy and active-inference programme is implementable in the same sense and is the nearest neighbour (Friston, 2010). The claim is not that this is the only such framework but that being one is a different kind of warrant from coherence, and this account has it.

3 The same demand, turned on this work

If a capability change is a quantity with an interval, then *this account's own capability differential is one too*, and having required that of everyone else it cannot exempt itself. What follows is what such an estimate would take and what the present record supports, which is not the same thing.

The design is specifiable. Choose tasks in advance on which the framework claims a difference — stating something a control vocabulary cannot state, resolving a dispute the field leaves open, locating a defect where the account says defects sit — assign readers to conditions, and have the outputs scored by assessors blind to which condition produced them. The estimate is then a differential with an interval, per frame class, because the denominator is a frame's own and pooling across kinds of reader estimates something belonging to none of them.

One instrument is ruled out by this account's own prediction, which is worth stating because it is the instrument anyone would reach for first. *Self-report is invalid here.* A description that is right reads as obvious (the substrate paper's treatment of carvings), so a reader who could already do the thing and merely could not say it will report a large gain — and a reader who genuinely gained nothing will report the same. Agreement measures recognition, and recognition is what this account predicts regardless of whether any capability moved. What has to be measured is *performance on a task*, not assent.

For a reader, no such interval exists: what there is is an author's enumeration of what the framework did for the author, without control, blinding or independent assessment, which is not a wide interval but the absence of one.

A second self-application, carrying the same limit. The production of this account exhibits a mechanism the account prices elsewhere, and saying so requires the care the previous paragraph demands. The substrate paper's definition of well-formedness charges search on every operation over what is held, scaling with how much is held. *Decomposition is search*: to ask what a claim says about some case it was not written for, one has to hold the claim. So a claim narrow enough to be held entire is a claim that can be tried anywhere, and a claim that must be looked up is one whose application to a new case will not occur to anybody. Narrowing and breaking down are therefore not competing directions but one and then the other, and the first is the condition of the second.

What the record supports and what it does not are different here in exactly the way they are above. The *sequence* is documented — a ledger driven down to four entries that can be read at once, and the decompositions that followed by taking one measure to cases it had not been written for — and a sequence is not a cause. No condition exists in which the claim stayed wide and the decomposition was attempted anyway, so what is available is an ordering, not a

differential. *This is the same absence of an interval named above, and it is not improved by being about method rather than about capability.*

One consequence separates from the rest because it can be read off the documents rather than reported. *The account grows while what it claims shrinks.* A decomposition that succeeds converts a posit into a consequence, which lengthens the text by the derivation and shortens the ledger by an entry, so pages and postulates move in opposite directions. A reader taking length as a proxy for how much is being asked will therefore read the trend backwards — which is the companion's point about a short generator with long output, met a second time in the making of the thing rather than in its structure.

The direction has the floor that every other use of it has. Narrowing helps while it stays lossless; past that it yields a small account of nothing, which is the limit §8 reaches from the other side — the most well-formed question asks and receives nothing, and arriving there would be dissolution rather than arrival.

And one thing review cannot do, established by its failing to. Six review passes over this corpus, each running some twenty-odd lenses, did not find that describing the substrate as a *typed* graph asserted something the substrate does not have. Nor that a companion had reinstated *substance* in the sentence claiming to have decomposed it away, with the word already on that companion's own list of things removed. Neither was hidden. Both sat in prominent sentences through every pass.

The reason is the one the substrate paper gives for numbering its axiom where it does: *the import is free.* A lens checks claims — against each other, against their sources, against what they were said to replace. A commitment arriving as vocabulary is not a claim in the text; it is a claim in the reader, supplied on reading and never written down. So the lens has nothing to stop on, and a reviewer finding nothing is reporting accurately about the wrong thing. This is not a lapse of care that more care would fix, and it should not be recorded as one: *a procedure that examines what a text asserts is constitutionally unable to examine what a text lets its reader assume.*

There is a sharper reason than blindness, and it decides which imports survive. *The ones that last are the true ones.* A graph described as typed is typed — kinds arise wherever a plurality is operated over, which is wherever anything runs — so every lens asking whether the description holds returned yes, correctly. What was wrong was never the content but the *ordering*: a derived thing taken as given. A false import is refuted by a counterexample and does not last a single pass; a true-but-derived one is refuted by nothing, and passes every check that asks whether it is so. So review does not merely fail to catch these — it *selects* for them, retaining exactly the imports no counterexample will ever remove. The same holds of the second instance: there is one graph-data stuff in the obvious sense, which is why the sentence read as harmless for as long as it did.

What did find them was not review but removal — the companion's one operation, applied to a word rather than to a posit, which is a candidate class its stated list did not have and now does. That gives a division of labour rather than a defeat. Review checks assertions and catches the defect this record shows most of, a correction made at its site and not swept to what points at it. Removal catches the free import, and nothing else here does. *An account revised only by review will drift in exactly one direction*, accumulating imported vocabulary while its explicit claims grow more careful, and the two are easily mistaken for one another from inside.

4 What kind of contribution this is

This material was carried in the substrate paper's scope section, where it sized the section and answered questions that are this companion's rather than that one's. It is reproduced here unchanged in substance.

At best a correct account of a substrate; at worst a useful but incomplete description. Which

of those it is cannot be settled from inside it, and the more useful statement is about the *form* of what is offered rather than its correctness.

What kind of contribution this is should be said first, because it determines what counts as a defect against it. *The claim is explanation and unification, not discovery.* Almost everything this account describes has been observed and named already, in its own field, by people who were not looking for a substrate — stasis and exaptation and niche construction in biology, opacity between levels in engineering, boredom's function in psychology. That those phenomena occur is not this paper's finding and is not offered as one. What is offered is an account of *why* they are the same phenomenon, which is a claim about mechanism rather than about observation.

Two kinds of prior-art finding have to be told apart, and running them together absorbs both.

A *phenomenon* already described in its own field is the material rather than a refutation: nothing can be unified that nobody has found. A *mechanism* already published is a different matter — it is a straightforward loss of novelty, and the honest response is to withdraw the claim and say so at the size it was made. That has happened here more than once, most sharply where a mechanism this account had claimed turned out to be a result from 1989.

One tempting move is ruled out here rather than left available: that a unification whose cases were all novel would be *unfalsifiable for want of anything to check it against.* It is a non sequitur — novel cases would be more checkable, not less, since each would be an independent prediction — and it converts prior art into a virtue, which is the move this passage most needs not to make.

That is exactly the exposure the substrate paper's falsifiability section's composition commitment already carries, so no separate defence is added for it. And one further limit follows from this paper's own machinery and is stated here: its compression-convergence proposition predicts that unifications of this kind will be *available*, because structures compressed under shared constraints come to resemble one another. So finding one is not itself evidence that the account is right. The unification is the contribution; it is not the confirmation.

Stated that way the position risks being a ratchet in which nothing can confirm and nothing can refute, so the exposures are named here. Two are definite. *A case the vocabulary covers while the mechanism does no work* damages the claim, and what counts as doing work is the economics: if the account cannot say what a case costs, or says it wrongly, the vocabulary is decorating. And *a domain under weak compression exhibiting the arrangement* damages its compression-convergence proposition directly — ritual, ornament and systems maintained past their purpose are where cost pressure is not binding and the arrangement should therefore be absent. That test is developed where the proposition is stated, and it belongs here too, since this is where a reader forms the objection.

Beyond those, this paper carries ordinary predictions that can fail on their own: the substrate paper's located-degradation hypothesis specifies where classification failures should concentrate, and the commitments in the substrate paper's falsifiability section marked [SHAPE] are refutations rather than revisions if they fail. Understating that exposure would be a defect of the same kind as overstating the contribution.

What is claimed is a shape rather than a set of details. The details in this account have changed repeatedly and are expected to change again: definitions have been renamed, propositions weakened, citations corrected, and quantities that looked primitive have turned out to be roles played across a relation. Those revisions are ordinary and the account is written to make them cheap — which is why the falsifiability commitments are stated separately from the constructions, and why corrections are left in the text rather than smoothed out of it.

The word *shape* is doing real work there and must be pinned down, because "the details may be wrong but the shape is right" is the classic form of a theory retreating from its tests. Two things keep it from being that. The commitments in the substrate paper's falsifiability section are stated at the level of the shape rather than the details — whether a primitive beyond the

triple is required, whether a non-discrete quantity is needed, whether the arrangement recurs where its constraints do not hold — so a detail failing is a revision while a commitment failing is a refutation, and the substrate paper's falsifiability section marks each commitment [SHAPE] or [DETAIL] so that the distinction is a fact about the list rather than a claim about it. And the shape is specified enough to be reported wrongly: several of the corrections in this paper were shape-level, not cosmetic.

The honest form of the claim is therefore a *rate* and not an invariance. Details change fast; the shape changes slowly; neither is fixed. Over the work behind this account the shape has itself shifted and wobbled — more slowly than the details, and not negligibly — so anyone treating the shape as the stable part should understand that as a comparison rather than a guarantee. That is checkable against a revision record rather than merely assertable, so the record is kept and the current reading is reported here: over this account's most heavily revised single day, many revisions to this paper produced roughly a dozen corrections to claims it had already made, of which four were shape-level — a primitive that turned out to be a role, two components that turned out not to be required, and a property that turned out to be well-formedness rather than necessity. Roughly one shape correction to three detail corrections, on that day. A single day is not a rate, and it is offered as the beginning of a series rather than as a measurement.

The standard this should be held to, and where it currently stands

An account claiming to be about a substrate rather than a domain invites one demand, and it is worth stating in its strong form rather than a convenient one: it must handle a rock, the observable universe, cognition, computers, and quantum physics *at once*, and failure on any of them is failure.

That list is not arbitrary, and its structure is stated here accurately rather than tidily. A rock is inert and models nothing; quantum physics is active and models nothing; computers and cognition are both active and both model, and are separated from each other on a third axis — one is built and open to inspection, the other is selected for and is not. So three occupied positions rather than a spanned 2×2: the inert-and-modelling cell is empty, and nothing is claimed for it. The observable universe is the container the other four sit in and is scored separately below as a domain, which is a double duty, noticed rather than defended. What makes the list a test is that an account can pass any one position by being built for it.

What the demand asks for here is narrower than a unified field theory and should not be confused with one. This account claims that the same arrangement underlies all five, not that any of them is derived from it — which is exactly the composition claim of this section, and the reason that claim carries a falsifier rather than an argument.

Scored against the demand, the account reaches the level of shape in all five cases and the level of derivation in none. Two of the five entries also disqualify themselves as evidence in their own wording, and the summary should say so rather than let the count carry an implication it has not earned: *no entry below supplies independent support*, since the rock is the definitional reference point rather than an application of the account, and the computing case is explicitly ruled out for borrowed vocabulary. What the five establish is coverage, which is what the demand asks; what none of them establishes is confirmation:

Rock.
Strongest. It is g with r and Γ absent at the grain of individuation (its remark that only g is required), and it is the definitional reference point rather than an application.

Computers.
Strong as a worked instance and explicitly disqualified as independent evidence, since the vocabulary is borrowed; what survives is the forced convergence, which is separable (its compression-convergence proposition).

Cognition.
Substantial but grounding-level: aperture, self-models, the closure of primitives, and convergence under a metabolic budget. It explains shapes; it derives no psychology.

Observable universe.
Shape-level only. The physics companion recovers the *form* of relativity from one change-budget with c as a frame's own unit. No metric is derived, no dimensionless constant fixed, and it is not a quantum theory of gravity.

Quantum physics.
Weakest, and the one on which the demand bites hardest. The Born rule is *recovered* — reproduced within the account's structures given substrate-supplied axioms and cited reconstruction theorems — rather than derived from necessity, and complex-over-real comes out contingent rather than forced.

On a strict reading of the demand that is a failure, and the strict reading should not be argued away: an account that recovers rather than derives has not met the standard it has just adopted. What is claimed instead is the weaker thing the composition falsifier states, and the value of holding the strong standard in view is that it names where the work is unfinished — which, on the above, is chiefly the last two entries and specifically the gap between recovering a rule and requiring it.

It is worth being explicit about the standard that scorecard is written against, since a reader could take breadth-at-shape as an excuse for shallowness and it is not offered as one. The measure is Heinlein's competent man (Heinlein, 1973) applied to theories: whether one account handles problems of different *shape* — inert and active, modelling and not — rather than how far it goes on any one. A theory deriving quantum mechanics from necessity and saying nothing about a rock or a mind scores badly on that axis, which is the axis this paper occupies, and a claim about breadth is not a demotion of depth.

That is not a licence. Each competence has to be real — the competent man can actually butcher the hog, and breadth with no depth anywhere is dilettantism rather than generality — which is why every entry above is tied to a commitment in the substrate paper's falsifiability section that can fail on its own. And the evidence that counts is not this account covering five domains, which it could do by being flexible enough to cover anything. It would be those domains' own practitioners arriving at the same arrangement independently, in their own idiom, without reference to this account — an engineer forced to abandon a global clock by finite signal speed rather than any number of paragraphs in which the substrate is shown to accommodate one.

That clause is qualified rather than withdrawn. its compression-convergence proposition predicts exactly such independent re-derivations as a consequence of shared constraints, which calibrates what they show rather than abolishing it, and the second reading over-corrects. A re-derivation does discriminate against the artefact-of-description reading, since someone who reached the arrangement without this vocabulary cannot have been led there by it. What it does not do is separate the composition claim from the weaker reading that the constraints are simply shared, which this account predicts anyway and which was never in doubt. So the positive side is real and partial, and the burden stays on the negative side because that is the closeable one.

5 The implementation as the place to look

The implementation is the other place to look, and it is the better one, because there the differential has been measured repeatedly and with intervals. One thing must be said before any of those numbers are, since this section demands checkability of everyone else: *the measurements below are not in this corpus.* They come from the engineering record — commit history, per-run

result documents, an audit of capability acquisition — which is not published alongside these papers, and the constructive companion is a normative specification containing none of them. Until that record is available a reader cannot check any figure quoted here, which by this section's own standard makes them assertions rather than evidence, and they are marked as such throughout. The frame they concern is one whose capabilities are enumerable, whose before-and-after is a control, and whose tasks are scored rather than assented to — so a framework-motivated change either moves a measured number or does not. What that record says is worth reporting in the deflationary direction, since it is mostly null: *most framework-motivated changes made no difference.* On one recognition task the arms differing by a few points — analysis-by-synthesis, dwell scheduling, appearance, higher-order relations, a finer sub-stroke grain — all sat inside one another's confidence intervals at $n = 400$ and were noise. Two did not: composition, which is arrangement treated as the object rather than the parts, at $+21$ points and interval-separated; and the substitution of a strong continuous verdict for a weak one at $+51.5$ points, $[67.0, 79.1]$ against $[22.6, 31.3]$, which is this account's in-frame/cross-frame discipline applied to a comparator.

Those numbers are the wrong evidence even where they are real, and the reason is this account's own. *Absolute capability is the wrong quantity; the trajectory is the right one.* A score at a moment conflates the framework's contribution with everything else in the system and invites comparison against arrangements fitted to one benchmark, which is the substitution of a number for a shape that this programme is meant to avoid. What the framework claims concerns how a thing *grows*: the decode-and-mint loop is a growth mechanism, Narrativium is capacity for change, and momentum is direction of change. The derivative is the object, not the value.

The starting point makes that unavoidable rather than merely preferable. *The substrate began with no capabilities at all.* Every one it holds was acquired, so there is no inherited baseline to subtract and the entire record is trajectory — and comparing it against a system that began with fitted parameters or a hand-built solver compares a trajectory against a starting position, which is a category error rather than an unfavourable result.

That yields the claim actually worth testing, and it has a shape rather than a value: *the cost of the Nth capability should fall as N grows*, because structure is reused instead of rebuilt. The record carries instances of the right kind — a sequence decoder written for speech running unmodified on ink, a codebook loader that turned out to be grain-agnostic when a new grain arrived, a transform reader that handled rotated input with no change — and each is a capability added without paying again for what it stood on.

There is better evidence than either the scores or the cost curve, and it is better because of what the author did not control: a change of frame forces a learning, and the forced learning is then available to frames that were not being worked on. That claim has since been audited against the history, and the audit removed most of what had been offered for it. What follows is what survived.

The clean case is this. A dynamic-time-warping comparison and its analysis-by-synthesis decision procedure were built for speech, and were picked up sixteen days later by the commit that *created* the ink frame's first files — the beneficiary did not exist when the structure was made, which is as close to an uncontrolled counterfactual as this record contains. The example previously given in its place, a lattice decoder said to have been written for speech and later run on ink, is mis-attributed: speech and ink took it up on the same day nine hours apart with both frames under active development, which is co-development rather than transfer. *The origin attributed to it in an earlier statement of this paragraph — printed-letter recognition — is not supportable either*: no printed-letter or optical-character work exists anywhere in the record, and the seed's own description is grain-agnostic and declares no dependencies. A second example, a writer-style dimension described as the ink twin of a speech speaker dimension, is withdrawn entirely — the two share no code, and where the construction was finally evaluated it *harmed* the receiving frame against its own guard.

What replaces those is a stronger and narrower claim about portability, which is verifiable rather than argued: the decoder seed is byte-identical — same hash — from its creation to the present, across five dated points and seven consuming frames, never modified. Structure built once and consumed unaltered by frames that did not exist at the time is what the account predicts, and it is what the record shows.

The qualification given earlier was too generous and must be replaced by the measurement. It said most instances are components reused without modification and *fewer* are cases where the receiving frame's measured behaviour improved. The true count of the second kind is *zero*: every measured before-and-after improvement in the history is within a frame, and no cross-frame transfer has a measured benefit attached to it at all. The account's central evidential claim is therefore supported at the level of portability and unsupported at the level of benefit.

Two corrections in the account's favour. Failed transfers *are* logged — seventeen are recorded, including one that measurably harmed the receiving frame and one explicitly retracted — so the success rate is auditable and the failures are in the record rather than absent from it. And one transfer meets the strict date test cleanly: a defect census produced under an embodied control problem led, eight days later, to a fix in a conversation module that had been untouched since its creation.

5.1 What the substrate is, and what carries the weight

A reader of the companions that report results has so far been given almost nothing about the thing producing them, which makes those papers read as arriving from nowhere. What follows is a reference point, and two facts in it do the work. The inventory does not.

The architecture. A typed property graph g; rewrite rules r with *match*, *where* and *effects* clauses run to fixpoint; and an admissibility grammar Γ deciding which edges may exist. A Rust core of 89 files exposes a term language of 239 variants. The agent's cognition is *graph data*, not code: 805 seed manifests carry its rules, concepts, productions and reflexes, loaded by dependency closure at boot. The host language translates external input into graph mutations and reads state back out, and does nothing else.

The breadth is the least of it. The substrate has been pointed at perception, ink, speech, spatial navigation, several vehicles, a humanoid body, a first-person shooter, program synthesis, a benchmark's answer keys, four classical psychology paradigms, a simulated town, a video codec and a cipher. *That list proves nothing on its own and this paper does not offer it as proof* — the companion is explicit that graphs are expressive and almost anything can be drawn as one, and a catalogue of domains is exactly the evidence that objection predicts. Two things are not.

Transfer, which is what the composition claim actually rests on. The substrate paper's headline commitment fails if the recurrence of one structure across domains is an artefact of description rather than a property of those domains. Breadth cannot answer that. What can is structure built once for one frame and consumed *unmodified* by another, and the record's strongest instance is stated above: a decoder seed byte-identical by hash from its creation to the present, across five dated points and seven consuming frames. The cleanest case is stronger still, because it removes the author from the loop — a comparison procedure built for speech was picked up sixteen days later by the commit that *created* the ink frame's first files. *The beneficiary did not exist when the structure was made*, which is the closest thing in this record to an uncontrolled counterfactual, and it is why the transfer claims survived an audit that removed most of what had been offered beside them.

Hardware, which is what makes it a mechanism rather than a description. An account that only ever runs as an analysis cannot fail on a clock. This one closes control loops:

a graph decision at 46 μs at the ninety-ninth percentile, a fault detected in 8.33 ms, a recovery loop held at 240 Hz.

The portability boundary was tested across two microarchitectures and, usefully, *found* — and it does not sit where a reader would guess. The Rust engine is generic `x86-64` and needs no recompile: verified here independently by disassembling the shipped binary and finding *zero* AVX-512 registers in it, against 24,669 AVX2 ones. What breaks byte-identity between the two boxes is a *numerical library in the host language* dispatching on CPU features at runtime, and what does not transfer at all is timing, which is specific to the microarchitecture it was measured on. *So the substrate is the portable part and the host layer is not*, which is the architecture's own claim showing up somewhere it was not arranged to.

One thing must not be read into this. A companion analysis concludes that the digit recogniser's int16 form would fit a commodity microcontroller. *That is arithmetic over model size and a real-time factor, not a run*, no such device has been used, and it is a projection rather than a result.

5.2 What is wrong with that record

The preceding subsection would be advertising without this one.

Most of the evidence is untracked. The results directory is excluded from version control: seventeen of 1,216 files are in it. The external corpora are excluded. Most measured evidence therefore exists in exactly one working tree, and several arcs with real numbers have log files and no results document.

Pre-registration is real and incomplete. Eighty-four pre-registration stems have no matching results file. In one domain, thirty-eight have seeds but neither gate tests nor results, and the honest tag for those is seeded-but-unmeasured.

The grammar is declared more often than it is used. Rule effects write edges directly rather than through the admissibility check, and a majority of rule-carrying seeds declare no productions at all. This is named in the engineering record as traced, sized and not fixed, and it is repeated here because a substrate whose distinguishing commitment is generative admissibility should not have that fact reported only internally.

The acquired knowledge is inert. A large corpus of extracted encyclopedic material exists and is never loaded at boot. It is a data-production result and not a knowledge result, and the internal audit says so of its own sampled files.

The gate is materially red. The enumerated failure baseline — a list of node identifiers rather than a count, for reasons this programme had to learn — stands at 216 red node identifiers out of 1,092 files swept, under stated conditions, and is now some weeks old.

None of that is offered as a reason to discount the results the companions report, each of which carries its own controls and its own seed. It is offered because a reader given the first subsection and not the second would have been given a brochure.

6 The author's own estimate, and its bounds

The author's own estimate of the differential should be recorded, because withholding it while demanding intervals from everyone else would be inconsistent. *It starts not far above zero — interesting but not consequential — and it runs very high.* The lower bound is a real possibility and not a formality: everything here may be a competent practitioner's ordinary work, with the framework doing no lifting and the transfers being the software reuse any well-factored system exhibits. The upper bound is also real, since if the sequencing and transfer claims hold the implication is a general growth mechanism. The data do not currently separate those, and *the width is the finding* rather than an apology attached to one.

That contrast generalises, and it is a rule rather than a preference. *An account of this kind cannot apologise.* An apology is an instruction about how its own text should be read — *this looks like X, take it as Y* — issued from a standpoint outside the reading, and there is no such standpoint: the writer is in a frame and the reader is in another, and no sentence legislates across that gap. It is also held structure doing no work, which the substrate paper's definition of well-formedness prices as idle at the top level. So the prohibition is not decorum. It is the account's own axiom applied to its prose, and an apology fails on both counts at once.

The point is not that the impulse behind an apology is wrong, since it is usually a real weakness correctly felt. *The impulse has an admissible form and the apology is the inadmissible one.* What an apology wants is for a reader not to misjudge; what does that job is an *exposure* — the statement of what would settle the matter against the account — because an exposure asks for nothing and hands over something. Where this record shows apologies removed, none of the underlying worries were dropped with them: each was already carried, or came to be carried, as a commitment in a falsifiability list. That is the translation, and it is available every time.

Stated concretely rather than as a range of numbers: at the low end it is a working start whose consequences are local and modest — and the stronger phrasing available here, that it is a unification which did not previously exist with its sharpest consequences for the foundations of quantum mechanics, fails twice over: it smuggles a priority claim into the floor of the interval where it carries the least hedging, and it contradicts the substrate paper's own scorecard, which rates the quantum entry *weakest* and the one on which the demand bites hardest; at the high end it is a general framework of the kind that changes how a civilisation does things. The author's weight sits nearer the upper bound.

Three considerations put the weight lower. The reference class of unifying frameworks announced as transformative has a poor record, and the exceptions were identified as exceptions only after decades of work by many hands rather than by their authors — so the prior is low and it is a prior about the class, not a judgement of this instance. This programme's own recent record is a calibration signal pointing the same way, since repeated prior-art checks have found the neighbouring work closer than expected, always in that direction. And nothing in the evidence discriminates *within* the upper region: forced cross-frame transfer and a buildable specification are hard to obtain by luck, and they separate a useful framework from an empty one without saying anything about the distance between a useful framework and a civilisational one.

The account supplies a stronger reason than any of those, and it is the one that should govern. *The upper bound is a claim about other frames, not about this work.* Impact is a capability change measured against the absorption capacity of whoever receives it — a magnitude relative to frames the author is not in, gated by their apertures and paid out of their Narrativium. By this substrate paper's Axiom -1 there is no position from which one's own impact is visible, so an author's weighting of their own upper bound is not evidence about the world and cannot be, whatever its merits as a prediction.

One thing must be established rather than assumed, the paper having already ruled the assumption out. It read the upward weight as *functional* — as an expectation of mattering that the builder needs in order to justify the spend. That is not what happened here. *It was built, and then it mattered*, which is the found-then-explained order the substrate paper insists on when it says explanation is not what licenses a structure but what is done with it once held. Assuming the rationalist order for an author's motivation, immediately after denying it for everything else, was the same error one level up.

The consequence runs the other way from the dismissal it replaces. A weighting formed *before* the work, to justify committing to it, is motivated and discountable. A weighting formed *after*, by building the thing and observing what it then did, is a posterior — an observation rather than a bet, and better evidence than the prior would have been, though still bounded by one hand and no external assessor. Nor did the expenditure require the expectation: the

account already has most search paid by whoever finds nothing, so building without a view about mattering is not an anomaly needing a motivational story but the ordinary case.

The upper bound also has an argument behind it that should be stated, because it is a conditional with a threshold rather than an unbounded hope, and that is a better epistemic object than the bound alone. It runs: take the system as built, supply enough compute for it to become generally capable enough to improve itself substantially, and *even if the only result is an order of magnitude in reasoning and efficiency over current language models, that suffices* — given the magnitude of change those have already produced, a tenfold improvement on them is not a modest outcome. The threshold is set deliberately low so that the bound does not depend on anything exotic.

The form is right and the risk is all in one clause. "Capable enough to improve itself substantially" is not a setup step; it is the entire distance, and every historical claim of this shape has failed precisely there. What this substrate has demonstrated is growth at small grain — a residual mints a quantum, a codebook grows, the loop closes and terminates — and the gap between that and substantial self-improvement is unbridged and is not a matter of scale alone.

The audit reported above bears on that clause, and its reading requires a distinction this paper had been running together. *Self-direction and self-implementation are different things, and only the second is absent.* Deciding what to build next — decoding the residual, identifying where the largest unabsorbed thing sits, choosing the cheapest structure that would absorb it — has been exercised extensively and is what set the acquisition order discussed above. Writing the resulting structure has been performed by a language model rather than by the substrate's own rule-writing, and that half is *deliberately* not built: below a threshold of general coding and problem-solving competence, a system modifying itself mostly churns, or efficiently makes itself worse in a domain where it is weak.

So the finding that *nearly half the primitives introduced have no consumer in any seed*, with the anti-unification and substitution and transfer operations having essentially none, is not composition having been attempted and failed. It is the executor being elsewhere. That reading is more favourable than the one this paragraph first gave, and it should be corrected to it.

It does not, however, relieve the antecedent, and it sharpens two problems rather than resolving them. If direction is the substrate's and implementation is the model's, then the falling structural cost measured above is attributable to either, and the confound the audit named — assistance improving over the same period — becomes harder to remove rather than easier. And self-direction exercised extensively is the claim with the thinnest ledger behind it — though that statement needs its own correction first, because the audit producing it searched about a quarter of the available record. The working history runs to some four thousand commits and nineteen hundred documents over five months, including several hundred pre-registrations with results documented against them, and the sweep covered the smaller of the two repositories holding it. Re-checked against the whole, the finding survives: explicit sequencing-error language is close to absent, and material comparing an estimated cost against an actual one runs to a handful of documents rather than to none. It now rests on a search adequate to support it rather than on one that was not.

7 The scale of the record, and what an absent error log shows

The scale is worth recording for its own sake, since Section 5 referred to "the record" throughout without saying what it is. Five months of one person's work, pre-registered at that rate, is a substantial evidence base for a single-hand programme, and the qualifications above bound what it shows rather than how much of it there is.

What remains is the part that matters, and the obvious reply to it should be taken seriously before it is answered. *A near-absence of sequencing errors is the shape a working procedure predicts*: if the criterion really does locate the largest unabsorbed residual and the cheapest

structure that absorbs it, then following it ought to mis-sequence rarely, and a high error rate would be the falsifying observation instead. That is correct as far as it goes.

It does not go far enough, because an unaudited process predicts the same observation. Few errors is what a good procedure produces and also what a procedure whose errors are not recorded produces, so the observation does not discriminate between them — and treating it as confirmation is precisely the move the substrate paper spends its falsifiability section refusing.

This account's own machinery states the position more sharply than the ambiguity does. A recognizer that always answers can never produce a residual, and a grain with no residual cannot grow; that was measured on this project's own meta-shape faculty, where deleting the tiebreak — making failure possible — is what made growth reachable at all. *A sequencer that never errs is in the same position*, and the remedy is the same: not to celebrate the rate but to arrange for the procedure to be capable of visibly failing.

The framing above imports a category this account does not have, and correcting that dissolves most of the puzzle. *Errors are not errors here; they are residual.* An "error" presupposes a correct answer that was available and missed. The substrate has no such object — there is what was admissible, what was realised, and what did not fit — and the last of those is residual, which is not a failure to be logged apart from the work but the input the next mint runs on.

So the search that found nothing was looking for the wrong word. Sequencing choices that turned out expensive do not appear as *built too early* or *wrong order*; they appear as supersessions, retractions, registered negatives and rungs whose results overturned an earlier rung — and of those the record holds a great deal, including dozens of enumerated negatives and a set of transfers that failed with one measured harm among them. The material is there in quantity, under the name the account actually uses. An absent error log is what a system with no error category looks like, not what an unaudited one looks like.

One part of the objection survives the correction and should be kept, because residual being logged does not grade the sequencing. *Whether a choice produced more residual than a different choice would have is a counterfactual*, and no amount of recorded residual answers it. Grading a sequence against its alternative does require running the alternative, which is the expenditure the sequencing existed to avoid and is again the identity of the test and the change. But the counterfactual is not the only grading available, and treating it as though it were concedes too much, because *residual comes in kinds and not only in sizes.*

A choice producing a *large* residual is frequently a good one. Much did not fit, but something did, and the part that did locates the boundary — so a large partial residual is a specification for what to mint next, and is the most informative outcome a write can have. A choice producing *complete* residual is different in kind rather than worse, and calling it uninformative gets it backwards. Nothing cohered and no continuation was realised, but that outcome *eliminates* — with a qualification the first version of this passage omitted and which the rest of the paper requires. What is eliminated is the region *in the searching frame's own space*: this frame cannot hold it, and that is a fact about this frame rather than about what can be held anywhere. Residual is mismatch and not diagnosis, and it does not become diagnosis by being complete. Within a search that is enough, since a search is conducted in a frame and removing a region from it is real work — and *the amount removed scales with the scope of the question asked.* The largest question that returns full residual removes the largest region of the solution space, which is the most informative negative result available and is not a small thing to have bought with one write. The substrate paper already relies on that elsewhere, where the composition commitment is deliberately negative-only because the negative side is the closeable one. A choice producing *no* residual is the genuinely uninformative case — absorbed, cheap, and telling the frame nothing it did not already hold.

So the value of a choice has two dimensions rather than one: the *scope* of the question, and the *kind* of residual it returns. A large partial residual localises within a large space and is what construction wants; a complete residual on a large question eliminates a large space and is what

refutation wants. Both are good writes and they do different work. What is weak is a question of small scope, which neither localises much nor eliminates much whatever kind of residual it returns.

That gives a criterion the counterfactual is not needed for. *A sequencing choice is assessable by the scope of the question it asked and by whether its residual localised or eliminated*, both of which are visible in the record without knowing what a different choice would have produced. It is this programme's own standing rule one level down — a null that localises the cause and generalises the lesson is a better outcome than a pass that does not.

The criterion applies to a review as readily as to a build, and applying it corrects the obvious way of reading one. A review is a write; what it returns that does not cohere is residual; and the residual is what the next revision runs on. So the thing to ask of a review is not which direction its findings moved the account — a finding that costs the account something is not thereby worth more than one that does not, and grading them that way is scorekeeping with the sign flipped. What to ask is whether it left residual with content, which is the same question asked of any other write. A review returning nothing is the expected case and costs little; a review returning something localised is the one worth having, and its value does not depend on whom it embarrasses.

8 The question bounds the answer

Which is the general statement this whole treatment has been circling, and in this account it is a consequence rather than a sentiment. *The question is often more important than the answer.* An answer is a realisation within a frame; the question is what fixes the frame — its aperture, its change grammar, and therefore what will count as residual at all. So the value of a write is *bounded above* by the scope of what was asked: a small question neither localises nor eliminates much however well it is answered, while a large one raises the ceiling on every answer it could receive. Improving the question multiplies; improving the answer moves within a limit the question already set.

That is also the sequencing criterion in ordinary language. Choosing what to build next is choosing what to ask, and the account grades it by scope and by the kind of residual returned — which is why the choice of question is the part of the procedure that carries the weight.

Taken to its limit this yields a result worth stating, because it inverts the ordinary expectation that a good question should *produce* something. *The most well-formed question is one that asks and receives nothing.* It costs a single write; it returns no structure into the frame, so there is nothing further to hold, to search, or to maintain; and at maximum scope it removes the largest region of possibility available. Maximum work done, minimum retention incurred — which is the substrate paper's definition of well-formedness at its limit, and the same shape as this programme's observation that the empty limit is not a frame and does nothing, so reaching it would be dissolution rather than arrival.

The abuse it converts into must be marked, because the sentence reads as a licence for never answering anything. *Asked and received nothing* is not *not asked*. The first is a write that was made and returned nothing into the frame, which is complete residual and eliminates; the second is no write at all, produces no residual of any kind, and is the uninformative case in the strictest sense — nothing localised, nothing eliminated, nothing paid and nothing learned. It is the write that converts a question into residual, and the question bounds the answer's value without substituting for it.

A second abuse is nearer and has a name. A maximum-scope question that *answers* nothing is *god* — and the distinction from the case above is the whole of the guard. *Asked and received nothing* is a write that returned a complete residual: the largest region eliminated, and the elimination is what was learned. *Answers nothing* is a question at the same scope that eliminates nothing at all, because whatever is offered fits every case — and a term that fits every case

localises none, which is what makes it the maximally general explanation and the emptiest. The two look alike from outside, since both leave the frame with no new structure. They differ in what was removed, and that is checkable: name the region the answer eliminates, and if there is none, the answer is god by another word.

This bears on the account making the argument and should be said there rather than left implicit. Its own claim is not the second case. *To be is to stand in relation* is an answer with content — it eliminates every structure whose parts stand apart from nothing — and what returns nothing is the attempt to refute it, not the claim itself. A reader who reads the scoring rule above as licensing the claim's unfalsifiability has run the two together: the rule scores the residual a *counterexample* leaves, not the emptiness of an answer, and an account whose answer eliminated nothing would score worst by it rather than best.

And a programme of such questions alone would eliminate forever and build nothing. Narrowing by exclusion leaves a region and does not occupy it; occupying it is a write that produces retained structure, with everything that entails in holding and search. The account's own economics is what says so: elimination is the cheap half and construction is the half that is paid for afterwards, and a research programme needs both for exactly that reason.

The other half has a limit of its own, and the two limits are what this paper is built between. Minimum description length asks the maximally scoped question in the opposite direction: not what can be removed, but *what is the shortest structure that accounts for all of it.* That question is maximal by construction, since answering it requires taking everything at once rather than a piece — and its answer, pushed to the limit, is the smallest thing there is. *The biggest possible question has the smallest possible answer*, and this account's name for that answer is g: what exists, relationally, and nothing further (the substrate paper's minimal case).

That is not a coincidence but what compression means. The more a description must account for, the more the structure shared across it dominates the structure peculiar to any part, and at the limit the only thing shared by everything is that it stands in relation at all. So the two maximal questions meet at opposite ends of the same write: at maximum scope with nothing cohering, the largest region is eliminated; at maximum scope with the shortest sufficient account, the least structure is retained. This paper begins at the second of those, and the reason it begins there is the first.

That also explains the shape of its own testability, and corrects Section 1's opening in the process. Asking what exists is a claim of the form *everything is like this*, and its two directions are not symmetric: establishing it would require enumerating everything, and the extensive frontier is unbounded by this account's own argument — always one more step outward, never held complete — so the enumeration cannot terminate even in principle. Bounding it requires *one extant thing that does not fit.* The asymmetry is therefore not a methodological preference imported from elsewhere but a consequence of what is being asked and of a property this paper has already argued for.

The word matters. A counterexample does not *refute* an account of this kind; it *locates its boundary*, which is what residual does everywhere else in the substrate account and there is no reason it would behave differently when the frame in question is this one. That is not a softening. An account whose range turns out to be the vocabulary rather than the world has been deflated to almost nothing, which is exactly what the composition commitment threatens — but the finding is a boundary, stated as one, and not a verdict of falsity delivered from nowhere.

Which sharpens the claim made at the head of Section 1, that a universal framework cannot be falsified. That is true of a framework which is only a *vocabulary*: if everything is expressible in it, nothing can fail to fit, and there is no counterexample to be had. It is false of a universal claim with *content*, and this one has content — that what exists is relational, discrete at operative resolution, and governed by admissibility — so a thing could exist and fail to be any of those. The counterexample is conceivable, which is what makes it a claim rather than a language.

And that supplies what Section 1 said falsification could not: the discriminator between a

well-formed universal framework and a vacuous one. It is not that one is more useful or cheaper, though both hold. *It is whether attempting a counterexample has an outcome.* A framework nothing could fail to fit is a vocabulary and should be assessed as one; a framework something could fail to fit is a theory about what exists, and is answerable in the only direction such a theory can be.

The test has to be stated as *attempt* rather than as *conceive*, because the conceivability reading gets the most basic case backwards. Put a counterexample to the relational claim and the attempt does not produce a coherent object the account then has to accommodate; it produces nothing, with the whole frame as residual. A reader who requires the counterexample to be coherently statable before it counts will conclude the claim risks nothing — when what has actually happened is that the claim is at maximum scope, so the only thing that could refute it is the one thing it says cannot be assembled. That is not the vocabulary case and it is distinguishable from it by inspection: putting a counterexample to a vocabulary returns another expression in the vocabulary, which is not a failure of anything, while putting one here returns a determinate nothing. *I tried and got nothing back* is a different result from *there was nothing to try*, and the difference is what this section's discriminator is for.

What a question does, stated as the operation it is. This section and the companion on mind hold the two halves of one thing and have not put them together. Here: a question fixes a frame's aperture and its change grammar, and therefore what will count as residual at all, while the answer is a realisation within what was fixed. There: a realisation is not the receipt of new information but the clearing of information already held and unread, and suddenness is the signature of an aperture that was excluding rather than of evidence that has just arrived. *Those are one operation.* A question is an aperture change. It supplies nothing — which is why the most well-formed one asks and receives nothing, above — and it changes what the graph can no longer decline to hold, after which the change follows without anything further being added.

Two things follow that are worth having. The first is why questions are cheap and productive at once: the residual a good question exposes was already there and already paid for, so what the question costs is one write and what it returns is everything that write makes readable. The second is that the *am I the one in the wrong here* case of the companion is not an analogy for this but an instance of it, differing only in that the frame is asking about itself — which is why it arrives as a moment, the aperture having been excluding until it could not.

The corresponding failure is a question that widens nothing. It can be perfectly well formed as a sentence and return an answer, and it will move a frame no further than the aperture it was already using, which is the ordinary case and is why improving the question multiplies where improving the answer does not.

9 The objection this account expects most often

Which brings the objection this account expects to meet most often, and it bites exactly here. *Everything can be represented by a graph, so the claim is trivial.* The first clause is true and the conclusion does not follow, and the shortest answer is that the universality is the point rather than an embarrassment — but the objection deserves more than that, because on the weak reading it is correct.

Three things separate the readings. *Representability is not constitution*: that a thing can be encoded with nodes and edges is cheap and uninformative, whereas the claim here is that realised dynamics require nothing beyond (g, r, Γ), which a fourth necessary primitive would refute. The content sits in the *prohibitions* rather than in the encoding — no view from nowhere, no substrate time dimension, no frame-free identity, no consequence serving as its own antecedent — and a world encoded as a graph but carrying a global clock, or frame-independent sameness, or retroaction, satisfies "everything is a graph" while violating this account outright. That is where a counterexample would come from, and it is why one is conceivable. And the universality is

load-bearing in a specific and limited way: the transfer claims this programme makes are about its economics, so what they require is that the economics range widely, not that every clause of the account be a strict universal — which the conclusion denies of the write-economics itself. The objection has identified the feature the account is built on and called it a defect.

There is no need to argue the point further here, because it is already the substrate paper's primary commitment: the composition falsifier states that the account fails if its structure is recoverable across unrelated domains *only* because a sufficiently general vocabulary can be applied to anything. The objection is not a rebuttal to be deflected. It is the test, stated by the account, in the account's own words, and left open.

The triviality objection is one instance of something that has held of nearly every result in this programme, and stating the general form is more useful than answering the instances one at a time. *The same claim is both trivial and consequential*, and this is neither a paradox nor a rhetorical posture — it is what the account predicts of a correct compression at large scope, and it has three separate derivations already given above. A cheap cut looks like a joint once found, so a good compression reads as inevitable in retrospect. The unmarked half of a distinction goes unnamed because naming it does no local work for anyone. And a description that is right reads as obvious, recognition being the success condition for description rather than an embarrassment about it.

The sharper reading is that the two are one property seen from two places. *Trivial is an in-frame judgement and consequential is a cross-frame one*: held, the thing follows immediately and surprises nobody; transferred, it changes what other frames can represent and do. Those are measured in different frames, so a claim of conflict between them would have to be made from a third — which does not vindicate the pairing, and establishes only that its two halves are not each other's counterexample. That is the substrate account's in-frame/cross-frame discipline applied to the reception of its own claims, and is the same reason the value of a disclosure is relative to whoever receives it.

The guard is that this signature is necessary and not sufficient, and stating it requires care because the obvious phrasing imports a dichotomy the substrate paper's founding axiom forbids. A framework is not *right* or *wrong* — those would be verdicts from no frame. *Newton was not wrong.* Newtonian mechanics is exact within its range and is still used; what relativity did was locate its boundary, not refute it, and a framework that is superseded has had its edge found rather than been shown false.

So what the signature fails to tell you is not whether the account is correct but *what its range is.* A framework can read as obvious and consequential from inside and hold over a much narrower region than its author supposes, and the feeling is identical in both cases because the feeling is an in-frame judgement and range is not. Finding the range is the same operation as everything else here: meet something that does not cohere, and the residual locates the edge.

One hazard follows from all of this and should be named, since it is how frameworks of this kind usually die. The prohibitions are what keep the claim from being a vocabulary, so relaxing any of them to accommodate an awkward case does not cost a detail — it moves the whole account toward the trivial reading the objection describes. Every such relaxation in the substrate paper is therefore marked as one.

What this does not supply should still be said. Per-choice quality is assessable; *ordering* optimality is not, since a sequence of individually well-localising choices may still have been available in a better order. The position is therefore narrower than a bare appeal to the counterfactual and better than it: each choice can be graded from the record, the sequence as a whole cannot, and the part that is gradeable has not been tabulated.

Two further gaps, both specifiable. "An order of magnitude in reasoning and efficiency" has no operationalisation here — tokens per solved task on a shared benchmark would be one, and no such comparison has been run. And the conditional's own premise supplies compute, which means the counterfactual that matters is the same expenditure without this substrate, and that

control does not exist.

What survives unchanged is the argument above it, which never depended on the motivation. Whatever its status as a prediction, the upper bound remains a claim about other frames' absorption and is not assessable from here.

The asymmetry is characteristic and should not be mistaken for a licence. An interval that includes large values does not make the large reading available — it makes it *unsupported*, which is a different thing that is easy to elide when the numbers are flattering. What would narrow it is already specified above: pre-registered tasks with blind assessment, denominators conditioned per frame, effort costed per capability against acquisition order, and an audit of the transfers that failed. None of that is exotic and none of it has been done.

The standard held to here is therefore not that the differential is large, or that it is established. The limitations are stated with the claims.

The trajectory claim has since been quantified, and the result is that *it is true in one reading and false in another, with the answer flipping on a modelling choice this paper had not made.* Taking a capability to be a shipped seed with its own tests, and counting the acquisition order across two hundred and thirty-eight of them: the number of *new primitive kinds* required per capability falls from 0.24 over the first hundred and forty-two to *zero* over the last ninety-six, intervals separated, with the five most recently entered domains requiring none at all while composing more held structure than earlier ones did. That is the predicted shape, measured.

The volume reading does not agree. Total implementation code grows linearly throughout with no saturation, and the marginal code per capability returns to founding-era levels whenever a new domain is entered. So structural cost falls to zero and volume cost does not fall at all, and which of the two *is* "the cost of the Nth capability" is a choice rather than a finding. The honest report is both numbers and the statement that the account did not specify which it meant.

Three further results bound the favourable one. Nearly half the primitives introduced have *no consumer in any seed*, so the growth of the primitive set and the composition of it are not the same story — one exception is large and real, a decode primitive used by eighteen seeds across eight domains over fifty-four days. Rework *rises* with acquisition order rather than falling, with later capabilities revised more often than early ones despite less exposure. And the confound that cannot be removed from this data has a name now: the same period saw the author's assistance improve substantially, so a falling cost per capability is not separable from a falling cost of assistance without a second builder. The attribution is arguable, since a competent engineer might have reached the same two changes without the framework and the prereg records only that the prediction preceded the result. The frame measured is an implementation and not a reader, so nothing here transfers to what these papers do for a person. And it is the same hand throughout, which the substrate paper's own rule disqualifies as independent support exactly as it disqualifies the computing case.

Three debts, stated rather than deferred. The comparison against a named rival framework has not been run here. This paper does not itself compile the record described above; the groundings do that, severally. And the implementation's own failures — the sharpest evidence available under this section's argument — sit in an engineering history that is not published with these papers, so neither they nor the figures quoted above can currently be checked by a reader. Publishing that record is the single change that would move most of the estimate above from assertion to evidence. What follows is therefore two kinds of thing and the reader should hold them apart. The substrate risks little by design, which is the claim just made. The *groundings* carry the risk, and the substrate paper's commitments that bear on them are falsifiable in the ordinary sense.

10 Modal clothing, and how to check for it

One habit is worth naming because it is easy, invisible in the writing, and was caught three times in this corpus within a day. An empirical regularity gets written in modal clothing: *the substrate forbids* for *nobody has found*, *cannot* for *does not, in every case measured*, *impossible* for *unattested*. The cost is not stylistic. A regularity stated as a prohibition acquires a quantifier over everything, which in an account with an unbounded frontier and no view from nowhere is a quantifier nobody can discharge — so the claim becomes weaker as it is stated more strongly, and its falsifier goes from a finite inspection to an impossible enumeration.

The check is mechanical and worth running on any new section: for each *forbids*, *cannot*, *never*, *impossible* and *must*, ask whether the support is an argument or a record. If it is a record, the honest form names the record — *every frame anyone has measured* — and the modal form is an overreach that costs the claim its test. If it is an argument, the modal form is correct and the argument should be cited at the site.

The corpus has this right in most places and got it wrong in one cluster, which is the useful part of the observation. The physics companion had already made the correction independently and left it visible — *so "impossible" was the wrong word twice over* — while the three failures were all in economics material written in a single day. The lesson is not that the discipline was missing. It is that new material does not inherit a discipline established elsewhere unless someone carries it across, which is the same defect as a repair applied at the site of an objection and not swept.

11 This paper's own commitments

A paper arguing about what an account may be held to should be held to something, and this one carried no commitments at all while naming a claim as the one worth testing. That is the defect it spends its length describing, committed by the document describing it, so the list is owed here.

1. [SHAPE] **The Nth capability costs less than the first.** The claim named above as worth testing, stated so that it can fail. It fails if, over a record of capability additions sequenced without regard to this account, the structural cost of the Nth does not fall as N grows — where *structural cost* is the count of new primitives a capability requires that no existing one supplied, fixed here rather than chosen afterwards. The ambiguity reported above — cost falling to zero on one reading and not at all on the other — is resolved in that direction here, and the limit of that repair has to be stated or the repair claims more than it does. The numbers are on the preceding page, so this definition is fixed *after* seeing them and is prospective only: it covers the next record and not the one reported here, which remains a result whose reading was chosen with the answer visible. Three terms in it also need operationalising before the count is a measurement rather than a policy — *primitive*, *capability*, and above all *requires*, which is counterfactual while what gets counted is what a builder added. This paper's own audit shows the gap, reporting that nearly half the primitives introduced have no consumer in any seed; and a builder following the project's standing rule to compose rather than mint drives the count down by discipline rather than by substrate economy, which is a confound closer than the assistance one named earlier. Finally, the record demanded here — capability additions sequenced without regard to this account — is not the record this paper has, since that one is the same hand throughout and disqualified for exactly the reason stated above. The commitment is therefore live for a future record and not for this one.

2. [SHAPE] **Usefulness does not substitute for falsification.** Fails if this paper's own argument — that a framework earns its place by what it lets a builder do rather than by

surviving a test — is used anywhere in this corpus to decline a test that was available. The falsifier is a citation: any place where a grounding could have stated a closeable commitment and appealed to usefulness instead.

3. [DETAIL] **The record is the evidence, and it is not published.** Fails if the engineering record, when published, does not contain the measurements this paper attributes to it. Until then every figure here is an assertion, which the paper says throughout and repeats in the debts above.

The list of prohibitions above once carried *cost attaching to the write* as a fifth member and no longer does. The substrate paper excludes it from its single prohibition, on the ground that folding it in would be the over-unification that section warns against. It stays in this paper's argument as part of what keeps the account from being a vocabulary; it is not among the things the account forbids. Note that being *derived* would not by itself settle this — a derived constraint on what is possible is still a constraint, which is why the substrate paper keeps a falsifier for the write-economics regardless of where it sits in the ledger.

References

Rudolf Carnap. Empiricism, semantics, and ontology. *Revue Internationale de Philosophie*, 4(11):20–40, 1950.

Karl Friston. The free-energy principle: a unified brain theory? *Nature Reviews Neuroscience*, 11(2):127–138, 2010.

Imre Lakatos. Falsification and the methodology of scientific research programmes. In Imre Lakatos and Alan Musgrave, editors, *Criticism and the Growth of Knowledge*, pages 91–196. Cambridge University Press, Cambridge, 1970.

Robert A. Heinlein. *Time Enough for Love.* G. P. Putnam's Sons, New York, 1973.