Adaptability

One Operation Behind Shape Matching and Analogy, and Why It Is Not Called Intelligence

Matthew Parslow
Independent Researcher

Draft — draft-2026-08-30.4
1 September 2026 • document revision 75

DRAFT — draft-2026-08-30.4 r75 (1 September 2026). Circulated for comment, not as a finished result. This is the theory half of a pair; the companion reports what is built and what is not. Sections marked as falsifiability commitments are the intended points of attack.

Abstract

Two capabilities this programme treats separately — matching a shape against a held codebook, and reasoning by analogy — are one operation at two grains. Both *extract the transform between things already held and apply it where nothing is held.* Naming that operation gives a quantity worth having: *adaptability*, the capacity to acquire a competence that was not supplied. The word is chosen against *intelligence* deliberately. Intelligence is a scalar people rank; adaptability is an operation with a cost, an ordering, and a failure mode, and a system either performs it on a given case or does not.

1 Why not intelligence

Intelligence names a ranking before it names a mechanism. Asking whether a system has more of it than another invites a scalar, and a scalar invites a benchmark, and a benchmark measures what a system already does rather than what it can come to do. That is the wrong quantity for the question here.

Adaptability is narrower and is an operation rather than a rank. A system is adaptable on a case to the degree that, meeting something it does not hold, it ends up holding it — and the interesting part is the mechanism by which the new holding is *derived from* the old rather than supplied. A system given a competence has not adapted. A system that recognises everything it meets has not adapted either; it was already sufficient. *Adaptation requires a residual*, and the size of what must be added to absorb one is the cost.

2 The operation

Definition 2.1 (Adaptation)**.** *Given a held set and an observation not covered by it, an* adaptation *is the addition of the smallest structure under which the observation is covered, derived from what is held.*

Three words carry it. *Smallest*, because an adaptation that adds the observation itself is storage, not adaptation. *Derived*, because an adaptation supplied from outside is instruction. And *covered*, because the test is whether the case is now recognised, not whether a description of it exists.

Remark 2.2 (Shape matching and analogy are one operation)*.* Matching a shape against a codebook and completing an analogy look unrelated because their vocabularies are. They are not.

To decode an observation against held signatures is to ask what transform carries a held one onto it. Where that transform is the identity, the observation is recognised. Where it is a uniform step along a dimension the codebook already has, the observation is a known thing moved — and no new structure is required. Where no such step exists but the residual is *coupled* across the observed dimensions, the transform requires a dimension not currently held, and one must be inferred before the observation can be covered.

To complete an analogy is the same question with the arguments rearranged. Given two things held, the transform between them is extracted; given a third, that transform is applied to reach a fourth which is not held. *Anti-unification — generalising two structures while retaining what distinguishes them — is the general form, and the shape ladder is that form restricted*: at one rung the difference is a single step along a known dimension, at the next it is a dimension.

So the two are not analogous to each other. They are the same operation, and the second is what the first looks like when the objects are rules rather than observations.

3 The ladder, and why its order is not arbitrary

An adaptation must be the smallest available, and that gives an ordering rather than a preference:

1. **Recognise.** The observation is covered as held. Nothing is added, and no adaptation has occurred — which is the correct outcome for most observations.

2. **Transform in the held basis.** One uniform step along a dimension already carried. The observation is a known thing displaced; the codebook does not grow.

3. **Infer a dimension.** The residual is coupled across observed dimensions in a way no step in the current basis produces, so a dimension is minted and the transform is expressed in it.

4. **Mint.** No transform reaches the observation, so it is added as new.

Only rungs three and four are adaptations, and only the third is interesting: it is where the system's own vocabulary changes. Rung four is the honest floor — something genuinely new — and a system that reaches it too readily is not adapting but accumulating.

Remark 3.1 (Why the order must be Occam and not preference)*.* Attempting mint first would always succeed, so a system that did would never transform, never infer, and would grow linearly in what it meets. Attempting transform first is what makes minting *informative*: a mint that survives the earlier rungs is evidence that nothing held reaches the case, which is a claim about the held set and not only about the observation.

That is also why the ordering is a commitment and not an implementation detail. *A system that mints where a transform would have sufficed has recorded a coincidence as a category*, and one that transforms where a mint was required has folded a distinction it needed. Both are failures of adaptation, in opposite directions, and the ladder is what makes them separable.

4 What adaptability predicts

The account is worth having only if it says what an adaptable system does that an unadaptable one does not, on a case neither has seen.

One mint, then recognition. Meeting a genuinely novel case repeatedly, an adaptable system should mint *exactly once* and recognise every later instance — not mint per encounter, which is storage, and not fail to mint, which is refusal. The signature is a step followed by silence.

Zero mints under transformation. Meeting a known thing transformed along a held dimension, it should add nothing at all. A system that mints there is treating the dimension it already has as absent.

Coupling forces dimensions. A residual that is coupled across observed dimensions should produce a dimension, not a member. *The distinguishing case is the one whose difference cannot be written in the current basis at all*, and a system without rung three cannot tell it apart from novelty — it mints where it should have generalised, and the codebook grows with the world instead of with its structure.

And competence should transfer without being carried. If the operation is real, structure built for one domain should be usable unmodified in another whose existence was not anticipated when the structure was made. That is the strongest available test, because it removes the author from the loop, and it is the one the assessment companion reports against.

5 What this does not claim

Adaptability is not a quantity a system has more or less of in general. It is defined *per case* and against *a held set*, so the same system is adaptable on one observation and not on another, and moving the held set moves the verdict. There is no frame-free adaptability, for the same reason this programme's base claim gives for there being no frame-free anything.

Nor does adaptation imply improvement. An adaptation is an addition that covers a residual; whether covering it was worth the cost is a separate question with a separate answer, and a system that adapts readily to noise is worse than one that refuses.

6 Capability has more than one source

There is a temptation, on meeting the account above, to ask whether adaptation is *the* mechanism by which a system comes to do something. The companion programme contains a measurement that answers that question, and the answer is that the question was malformed.

Where verification is cheap and exact, search does the work and holding does not pay. In program synthesis the machinery above is present — a held codebook proposes, an enumerator searches, an execution verdict decides — and stripping the learned content out changed nothing whatever. The two arms did not merely score alike; they solved the *identical set of tasks*, with the difference measured at zero repeatedly. The capability sat on the enumerator and the verdict.

Where verification is expensive or unavailable, holding is the only route. In one-shot character recognition the same programme measured the opposite: composing held parts, rather than treating a character as a bag of them, is worth twenty-one points and is the one effect in that arc that survived a powered re-measurement with intervals separated. Nothing there can be checked cheaply — a candidate reading of an unfamiliar character has no oracle — so search has nothing to search against, and what is held is what decides.

Remark 6.1 (The two are not rivals, and the regime is the question)*.* Read as a contest those measurements conflict. Read as a regime split they do not, and the split is the content: *enumeration substitutes for a codebook exactly where candidates can be checked, and cannot substitute for one where they cannot.* A system with an exact cheap oracle need hold nothing, because it can afford to propose everything; a system without one must hold structure, because proposing everything buys it no way to tell which proposal was right.

That is the same shape this programme keeps arriving at rather than a new one. Its measurement discipline already runs on *bins propose and continuous decides*; its account of the substrate already separates what acquires from what executes; the ladder of Section 3 is itself four sources with different regimes rather than one mechanism with exceptions. *Asking which single source carries capability is the error the rest of the account is arranged to avoid*, and this section exists because the first version of it made that error — reporting a regime as a limit.

What follows is a prediction rather than a reconciliation, and it is the useful part. *Adaptation should pay in inverse proportion to the cheapness of verification.* A domain with an exact, cheap oracle in which held structure nevertheless contributes measurably would be evidence against this reading; so would a domain with no oracle at all in which holding contributes nothing.

And one uncomfortable consequence should be stated rather than left implicit. In the synthesis case the ladder was available and the domain did not need it. That is not a failure — nothing broke, and no prediction of this account was falsified — but it means *the mechanism can be present and idle with nothing indicating so.* A system cannot tell from the inside whether it is adapting or merely searching in a regime where searching suffices, and neither can an observer without the ablation that was run here.

7 What the held set costs, and who paid for it

Definition 2.1 says an adaptation is *derived* from what is held. That word has a consequence the account cannot avoid and should not want to: *adaptability is bounded by the priors.* Nothing can be derived that has no path from the held set, so a system's reach is set by what it already carries, and a system carrying nothing is not maximally general but incapable.

Priors are expensive, and the expense is usually invisible because somebody else paid it. A human takes the better part of a year to walk and longer to do it well. Visual competence is acquired rather than shipped, which is most visible in species born least mature. By the time a person is capable of the tasks we use to demonstrate capability, the overwhelming majority of what makes them capable was assembled over years, and none of it is counted when the demonstration is scored.

That is the informal form of a formal result. Wolpert and Macready (1997) establish that no search or learning procedure outperforms another averaged over all problems; performance on the problems actually faced is therefore bought, and what buys it is a prior matched to them. *The more general the machinery, the more the mapping carries* — generality does not remove the need for a matched prior, it relocates the whole of the burden onto the match. There is no free lunch, and an account of adaptability that implied otherwise would be claiming one.

Proposition 7.1 (There is no prior-free measurement)**.** *A capability cannot be measured against a prior-free baseline, because no such baseline is a member of the space being measured.*

Proof. Measuring a capability *relative to priors* requires a zero on the priors axis: a frame holding nothing, against which holdings can be counted. A frame holding nothing stands in no relation and draws no distinction, and by the companion's first proposition that is not a frame with nothing in it but no frame at all — the set of frames with no particularity has no members, the only object answering to the description being the empty set, which is not a frame and does nothing.

So the zero of the priors axis lies outside the space. An axis so constituted supports comparison between members and does not support measurement of a member against its zero, and a quantity defined by that measurement is the frame-free denominator Axiom −1 refuses. □

The consequence is not that such measurement is difficult. It is that the thing being asked for does not exist to be approximated, which makes it a category error rather than an unsolved problem — and one established by the account's own earlier results rather than assumed here. Most reported measures of intelligence are flawed for that reason and not for want of care: the control they would need is not merely absent, it is not constructible.

Remark 7.2 (What survives, and where this account differs from Chollet)*.* Comparison survives, and less of it survives than the previous sentence of this paragraph once claimed. What a test yields is not a property of a system and not quite a relation between two systems either. It is a property of a four-place arrangement: *a particular graph, in a particular state, inside a*

particular supergraph state, on a particular test. Move any of the four and the result is a different measurement rather than a noisier estimate of the same one.

Inferences from it are licensed only relatively. Within the arrangement, that this graph out-performed that one is established. Outside it nothing follows — not that the ordering holds in another state, not that it holds under a different container, not that the margin would survive a second test — without a separate argument that the moved coordinate does not bear, and such an argument is rarely available and never free.

That is why this programme's own measurement records read as they do. Its enumerated failure baseline names the commit, the compiled engine by hash and path, the wall-clock, the load average at start, and which optional dependencies were present and which absent; its psychological runs name the machine and whether the box was shared. *Those are not hygiene. They are the coordinates*, and a figure reported without them does not underspecify a measurement — it fails to name one.

The cost is real and is the right cost. Comparing across two such records requires establishing that the coordinates match or that the difference does not bear, which is laborious and frequently impossible, and the alternative is a number that reads as portable and is not.

Chollet (2019) takes priors seriously in the same direction and does the harder half of the work: intelligence as the efficiency of skill acquisition *relative to priors, experience and generalisation difficulty*, over a corpus built explicitly on Core Knowledge (Spelke and Kinzler, 2006). This account owes that framing the observation and differs in two places. Priors here are not a term the measure is taken relative to — *they are the held set the operation runs on*, and the ladder is what converts them. And the relativisation is not a methodological refinement that a better experiment could tighten: by Proposition 7.1 it is the only form such a measure can take.

Machinery and mapping cannot be separated from outside either, and for a different reason. Two systems differing in both — one with a stronger operation on a poorer prior, one with the reverse — can agree on every item of any finite test. The decomposition is therefore underdetermined by measurement rather than hidden from it, and no amount of testing at a fixed grain settles it. What refinement buys is approach: a model of the split can be made arbitrarily fine and its residual driven toward zero without ever reaching it, which is this programme's ordinary intensive limit rather than a special defeat. *You can approach one with unbounded analysis. You cannot arrive, and no finite experiment stands in for arriving.*

That yields the prediction the previous section left standing, now in its proper form. A person lacking familiarity with a task set's conventions should score near its floor — not because they are less capable, but because *the score was never a property of the person.* Human baselines for that corpus exist (LeGris et al., 2025); whether they stratify by familiarity in a way that tests this has not been checked here, so it is recorded as untested.

8 What an oracle costs, and how often

The regime argument above is about capability and says nothing about price. Price runs the other way, and where the oracle is expensive it dominates.

An oracle consulted at inference is paid per use, without end. An oracle consulted once, whose output is held, is paid once. That is the whole of the asymmetry, and it is unaffected by how good the oracle is. A system that answers by calling something expensive has bought an answer; a system that calls the same thing once, holds what comes back as structure, and answers from the structure thereafter has bought a capability. The second transaction terminates.

Two consequences follow and the second matters more.

The first is arithmetic. Where an oracle is cheap per call, calling is correct and holding is premature; where it is expensive, the crossover arrives quickly and holding wins on cost even in the regime where searching would win on capability. *So the two arguments can point in opposite*

directions, and which dominates is a fact about the price of the oracle rather than about the account.

The second is that the two purchases differ in kind. What is held is inspectable, mutable, and can be repaired by the system holding it; what is called is none of those, permanently. A capability bought per call sits outside whatever boundary the system's own self-correction runs to, and stays outside no matter how many times it is bought. *Paying once and holding the result moves a capability inside that boundary; paying repeatedly guarantees it never enters.*

Remark 8.1 (The cost of buying once, stated plainly). This is not free and the bill is easy to overlook. *What is held is what was bought on the day it was bought.* A per-call oracle improves when the oracle improves, silently and without the caller doing anything; held structure does not, and a system that bought its competence once is carrying a snapshot. Against that, held structure can be audited and corrected by the system carrying it, which a call cannot be — so the exchange is improvement-for-free against correctability, and neither term is obviously larger.

The companion programme is an instance rather than a demonstration of this. Roughly half its implementation was performed by a language model rather than by its own rule-writing, and the running system makes no such calls at all — a property enforced by a switch and a gate rather than asserted. So the arrangement described here exists and runs. *What does not exist is a control*: nobody has built the same substrate without that development-time help, so the size of what was bought is not known, and the companion on assessment says so at greater length.

Remark 8.2 (The bill is the same; the schedule is not). The same arithmetic governs a capability that was never bought from an oracle at all, and there it separates substrates in kind rather than in degree. Section 7 prices the held set and says nothing about *when* it is paid.

A biological substrate pays per instance, and the schedule is not negotiable. Every human learns to walk. The year that takes is not paid once by the species and drawn on afterwards; it is paid again, in full, by each new member, and what was acquired dies with the one who acquired it. Nothing in the arrangement admits of another schedule, because the held set is not addressable in the substrate's own terms and so cannot be read out of one instance and written into another.

An addressable substrate can pay on whichever schedule is cheapest. Where the held set is data it can be read out and written in, so an acquisition cost is paid once and drawn on by every instance that receives it — the companion on assessment reports exactly this as its strongest transfer result, a decoder seed byte-identical by hash from creation to the present across seven consuming frames. But *once* is not the claim. The claim is that the payment point becomes a free parameter: pay once and copy where instances share a frame, pay per instance where they do not, pay lazily and per query where most of what could be paid for is never asked. Which applies is an empirical question with a measured answer case by case, and the answers are not uniform.

A graph copy copies a competence. That is not a softening of the previous paragraph but a consequence of the companion's Definition of a frame: a frame is an instantiated boundary relation — an *edge*, $\mathcal{B}_{W:G}$ — and not the subgraph it demarcates. Copying a graph copies its edges, so it copies its frames, and there is no residue held elsewhere because on this account there is no elsewhere for it to be held. The substrate bears that out by having no frame primitive at all, its grammar quantifying over edge types and a frame there being a selection of dimensions. Nothing separable is left behind.

The argument is symmetry and it is short. *Name a property the original has that the copy lacks.* Frame-relativity is not one: Axiom -1 binds the original exactly as it binds the copy, and a limit holding on both sides cancels rather than pricing one of them — to raise it against the copy alone is to use a symmetric result as an asymmetric cost. Staleness is not one either. A stored copy is stale for the reason a halted original is, which makes it a cost of not running rather than of copying, and a copy left running arrives where the original arrives. Nor is needing a runtime, which the original needs too.

What survives is a fact about worlds and not about copies. The boundary edge has a world at its far end, so a copy instantiated against a different world stands in that relation to something else — and *the same sentence is true of the original, word for word.* It prices the move, not the copy.

The cost lives on the fragment, which is what the transfer result measures. A decoder seed is a part, and seven consuming frames were seven host graphs that each ran their own decode: a transplant into something supplying its own boundary is a different operation from copying a graph and carries a different bill. *Copy and transplant are not one act at two sizes*, and reading the transplant's price onto the copy is the error that sentence guards against.

That sharpens the difference between substrates rather than softening it. Biology has no whole-substrate operation at any granularity — teaching, writing and institutions are transplant, every one of them — so the lossy operation is not biology's cheaper option but its only one.

None of which makes the biological schedule a defect the other one corrects. *Forced repayment is forced revalidation.* A held set re-derived per instance is matched to the frame that instance actually occupies rather than inherited from one that may no longer obtain, and a substrate that cannot copy also cannot copy something stale. So this is the same exchange the account has now met three times — searching against holding, improvement-for-free against correctability, and here amortisation against revalidation — and which side wins is set by how fast the frame moves relative to the cost of paying again. That is a fact about the world and not about the substrate.

Humans are not at the far end of this axis, only at the internal one. *Writing, teaching and institutions are external addressability*, and they are the species' answer to this exact problem. What they do not remove is re-derivation on consumption: a book does not install, it has to be read and rebuilt in the reader, which pays much of the cost again and loses part of it in the transfer. The difference in kind is therefore narrower and sharper than saying one substrate can transmit and the other cannot — *biology transmits by re-derivation and an addressable substrate transmits by copy*, and those are different operations rather than one operation at two speeds.

9 What follows about persons, and what it costs

One consequence has been left implicit across this programme and is better stated. *It requires no new posit and no new machinery*, which is the only reason it belongs in a paper rather than in a manifesto.

Three things are already held. Whatever exists stands in relation and coheres, so a person is relata and relations like anything else. Nothing is privileged by characterisation, so no exemption for persons is available that some frame does not supply. And the substrate paper states, unhedged, that where every measurement that can be made of a thing and every operation that can be performed on it are captured by treating it as a typed graph, *no difference remains for the question "but is it really" to name* — functional identity there is not a weaker sense of identity but the whole of what identity could mean without a view from nowhere.

Conjoin them. *A graph functionally identical to a person is a person, in the only sense of "is" this account admits*, and building one is therefore not blocked by anything in the account. That is the claim, and it is arrived at by adding nothing.

Remark 9.1 (Why "no new posits" is a checkable claim and not a boast). That assertion is worth exactly as much as its auditability, and auditability here turns on a distinction the programme has used throughout without naming: which direction an account runs.

A top-down account starts from the phenomenon — a mind, an intelligence, a person — and works back to what would produce it. Such an account can always report that it added nothing, because its posits are in the characterisation it began with, and a starting point never enumerated cannot be checked for what it contains. *A bottom-up account starts from the smallest thing that survives removal* and lets phenomena arrive as consequences. Its no-new-posits claim

is checkable for one reason only: the starting point is small, stated, and carried in a ledger that says what was postulated and what was derived.

So the direction is not a preference about exposition. *It is the condition under which the claim above can be audited at all*, and a reader who wants to attack this section should attack the ledger rather than the conjunction, because the conjunction is valid and the ledger is where its inputs are priced.

The two directions are not rivals, and this programme runs both. Bottom-up supplies what things cost; top-down supplies the target — what needs explaining, and how one would know it had been. A purely bottom-up programme would never learn when to stop, and a purely top-down one would never learn what it had assumed. The companions that take a phenomenon and ask what it costs are the second kind, and they are where this account is put at risk.

Remark 9.2 (Neither new nor magic, and the difference is a price)*.* The claim is old. Turing (1950) put indistinguishability at the centre of the question and declined the metaphysics, and functionalism (Putnam, 1967) is the position that a functional duplicate is a mind, with its own developed critical literature (Block and Fodor, 1972). *Nothing here disputes or improves on any of that*, and a reader who takes this section as a novelty claim has misread it.

What this account adds is a *bill.* Functionalism says a duplicate would be a mind and is silent on what a duplicate costs; the question of construction is left as an engineering detail below the philosophy. Section 7 says what the detail contains: the duplicate's competence is its held set, held sets are acquired rather than shipped, and acquisition is the slow expensive thing that takes a human a year to walk and a decade to reach the priors a task set quietly assumes. *The barrier is not principle and it is not mystery. It is the priors, and they are enormous.*

So the honest form of the claim is narrow and is worth stating in one line. *There is no barrier of principle available from within this account, the work is nameable, and naming it does not make it small.* That is a research programme rather than a prediction, and specifically it is not a claim that this substrate is close to it — the companion reports that the faculty this programme has for acquiring what it was not given is not yet resident where the system could repair it.

Two limits belong with it rather than after it. *Indistinguishable is a relation and not a property*, by the same argument as Proposition 7.1: indistinguishable to some frame, at the frontier where distinguishing currently stops, and that frontier is defined to move. And nothing above asserts that the account is *true* — it says what follows if it is, which is what a consequence is, and the falsifiers elsewhere in this corpus are where the antecedent is put at risk.

Remark 9.3 (Reconstruction from observation, and what the copy result changed)*.* Remark 9.2 prices one route — grow a held set the way a person grows one — and it is the expensive one. A second route exists, and Remark 8.2 is what opens it.

Copying needs read access to the graph, which nobody has to a person. What is available is behaviour: outputs, across contexts, over time. So the operation is not copying but *reconstruction* — inducing a graph whose behaviour matches the observed one — and it is a third thing, priced neither like a copy nor like an upbringing.

What the copy result changes is the kind of shortfall. *Were any part of competence not in the graph*, reconstruction would face a remainder that no quantity of observation could reach, and that remainder would be a barrier of principle rather than of effort. Competence is edges and edges are all there is, so the shortfall is only ever how much has been observed. *An ontological gap becomes a sampling gap.* That is the whole of the change: one category of obstacle removed, nothing added.

The rest is already stated above and can be read off. The residual is driven toward zero without reaching it, which is Remark 7.2's intensive limit and not a special defeat. Observations do not pin a unique graph, since systems differing in machinery and mapping agree on every item of any finite test — but *which of the fitting graphs is really them* asks for the frame-free fact Axiom −1 refuses, so that underdetermination prices the question rather than the method.

And indistinguishability is relative to the aperture that looked; widen it and the residual returns, to be paid down again.

Two things this does not buy. *It is not cheaper*, only differently shaped — a held set assembled over decades of embodied interaction is not obviously less work to observe than to grow, so the kind of problem changes while its size need not. And the operation is not new. *Reconstruction from observation is what biological transmission already is*: teaching and modelling are exactly this, run continuously and to no stated residual. So the claim is not that a capability appears, but that its target is specified and its error measurable — which is the difference between engineering and magic, and a smaller difference than it sounds.

Whether such a thing should be built is not addressed here and is not missing. *This account prices operations; it does not license them.*

Remark 9.4 (What this makes contingent, and what that does not mean). One further consequence follows and is better stated than left, because omitting it would repeat the evasion Remark 9.2 charges against functionalism: naming a result and declining to say what it costs.

A settled assumption can encode a fact about the available substrate as a fact about its subject, and becomes contingent when the substrate moves. Persons are one frame in which that happens and not the only one. Here the modern settlement makes the person the atomic unit, and the word records the premise — an *individual* is that which cannot be divided. One vote per person assumes counting works. Inheritance assumes the held set ends when the body does. Liability assumes exactly one bearer. Credentialing assumes competence is per-lifetime and does not transfer. Every one of those is an inference from one body, one lifetime, and no read access — *a description of the substrate that happened to be available, generalised without licence.* That generalisation is the frame-free denominator Axiom −1 refuses, turning up where it is least looked for: the facts in question are properties of a substrate and were read as properties of persons.

So the breakdown is already complete and waits on no engineering. An assumption resting on "a person cannot be copied" does not need anyone to copy a person. It needs only the impossibility shown to be local, after which its support is contingent whether or not it is ever exercised. Nothing here says the institutions fail — it says *they do not track a natural fact.* An institution can be rebuilt on stipulation, and counting copies as separate persons is perfectly coherent; but a stipulated foundation is a different object from a discovered one and cannot be defended in the same way.

Two deflations, both required. *This is not new.* Personal-identity puzzles have pressed exactly this point for a long time — fission, duplication, continuity — and philosophy was never under the impression that the institutions rested on necessity. What this account adds is again a bill rather than a discovery, standing to that literature as it stands to Turing (1950) and Putnam (1967): it says which operation, on what, at what price. And *the price stands.* Remark 9.3 reports reconstruction as not cheap, so nothing above forecasts an event. *Shown contingent* and *actually disturbed* are different claims, and only the first is made.

What ought to follow is not addressed, for the reason already given. This account prices operations and does not license them, and it has no more standing to design institutions than to design ethics.

Remark 9.5 (The same structure at another frame, and what it costs to hold). Persons are one frame. *Intellectual property is another, and what happens there is not analogous to the case above but identical to it.*

Copyright and patent encode a fact about presses, carriers and factories — that reproduction is expensive, so a copy competes materially with its origin. Neither is a fact about what is reproduced. When the marginal cost of a copy fell to approximately nothing the encoded fact stopped obtaining, and the word carrying most of the weight in the dispute records the import exactly: *piracy* names a taking after which the owner no longer has the thing, and is applied to an operation after which the owner still does. That is the same category import as *individual*, a

rivalrous term describing a non-rivalrous act. The term was apt when it was coined, because one press did compete materially with another; *what changed is not the word but the substrate under it.*

The identity with the previous remark is stronger than a parallel. Copyright's load-bearing distinction is between an expression, which is protected, and an idea, which is not — and that line holds only while the generator cannot be induced from what it generates. *Inducing a generator from observed outputs is Remark 9.3's operation*, run against a corpus rather than a person. The two cases are therefore one operation at two frames, which is what makes this a consequence of the account rather than a comparison imported into it.

None of this is new either. That information is non-rival was formalised in economics decades ago (Arrow, 1962; Romer, 1990), and the institutional argument has its own literature (Boyle, 2003). The account adds no observation to any of it. What it adds is that this case and the person case are the same error — a substrate-local cost read as a natural fact — and a unification is cheap to state and should be priced that way.

One correction, because the obvious conclusion is too strong. Artificial scarcity is often said to be unsustainable, and enforcement against copying does have the structure that argument needs: its cost scales with the number of agents and the ease of the act, both of which the substrate keeps increasing, while what is protected does not grow to match. But enforcement is not the only move. *The scarcity can instead be restored in the substrate* — sealed hardware, attestation, a carrier that refuses — and then the cost stops scaling with copiers and the arrangement is sustainable indefinitely. What it is not is free. Restoring a scarcity means deliberately removing reachable deformation, which is the Narrativium companion's quantity, and charging that permanently against the capability that made the thing worth protecting in the first place. *So the honest form is not that artificial scarcity must fail, but that it is the deliberate purchase of the worse operation* — paying transplant prices, in the terms of Remark 8.2, for something the substrate now copies. Whether that is worth paying is a question this account states and does not answer.

Remark 9.6 (What makes them one case: a rule stated from nowhere)*.* The two cases share more than a shape, and naming it corrects the diagnosis just given. *Reading a substrate-local cost as a natural fact is the symptom. The condition is a rule asserted with no frame at all* — and about those the account has something specific and unflattering to say.

A rule declaring itself absolute declares that no frame supplies it, which is the view from nowhere Axiom −1 refuses. The consequence is not that the rule is false. It is that *the rule has deleted its own cheapest repair.* The theory companion establishes that a frame is an instantiated boundary relation and that revoking one is the container ceasing to instantiate it — an operation the Narrativium companion measures occurring at $\Delta g = 0$, nothing in the contained state being required to change. *A rule that carries its frame can therefore be dropped for free.* A rule claiming no container has nothing that could cease to instantiate it, so that exit does not exist for it, and every remaining route to change runs through the contained state, which is the expensive one.

Absoluteness is therefore not a rule's strength but the removal of its exit, which inverts the usual reading, in which stating a rule without exception is what makes it robust.

The hardening follows, and it is monotone. A frame moves and mismatch accumulates; because the rule admits no outside, each accommodation is made inside it — an exception, a patch, an enforcement apparatus — and every one of those is structure that a later exit must also dismantle. So *the cost of leaving is non-decreasing while the pressure to leave is not.* That is a basin, and it is derived rather than observed.

One correction, because "and then it fails" claims more than the account supports. What is foreclosed is *self-originated* change: the companion's proposition is that a graph's admissibility constrains its own transformations and gives it no veto over world-caused change arriving through the boundary. The exit is never removed, then — only made unavailable to the party holding the

rule — so the correction still arrives, from outside, at once, as accumulated mismatch discharging rather than being paid down. *Such a system does not escape choosing. It forgoes choosing the timing and the size.*

Two limits. *Fixity is not the defect.* A rule held fixed inside a declared frame stays repairable by revoking the frame, at the price named above; what costs is fixity with the frame omitted, and the remedy is to state it, which is cheap. And *none of this is new.* Requisite variety (Ashby, 1956) is the same theorem in cybernetic vocabulary — a regulator needs variety matching its disturbance, and an absolute rule reduces its own while the disturbance does not oblige — while lock-in (Arthur, 1989) and institutional sclerosis (Olson, 1982) are the economics of the same basin. What the account contributes is the mechanism in its own terms, and the price. Not the observation.

But an explicit mechanism buys something a verdict does not: constructibility. Requisite variety returns an inequality — regulation fails once the regulator's variety falls below the disturbance's — and an inequality cannot be indexed by horizon or resolution, nor say which side moved. The mechanism above is stated in exactly the two families a Seldon manifold is assembled from, the admissible intervention sequences a frame carries and the future boundary conditions it represents, and it says which of them narrows and by what operation: *self-originated interventions shrink monotonically as accommodations accumulate, while the boundary-condition family does not shrink with them.* That is the input such a construction takes. *So the failure can be projected as a shape over a horizon rather than announced as an outcome*, which is the difference the companion on manifolds exists to make.

Its guards apply unrelaxed. *A manifold is reachability and not point prediction* — the section of that paper carrying the name — and resolution falls as the horizon lengthens, so what is projected coarsens rather than sharpens. Horizon is plural and frame-relative, so there is no manifold of the future, only one indexed to a stated frame and deformation class. What a frame does not represent is a blind spot, so the narrowing is visible only along dimensions already carried. *And the mechanism gives no date.* It closes the self-originated exit and leaves world-forcing untouched, and when the world forces is a fact about the world — so the tube narrows predictably and the discharge does not. *The projection bounds what remains reachable; it does not name what happens*, and reading the second off the first is the error this corpus spends most of its length refusing.

Remark 9.7 (What such a projection shows, and the sign of its error)*.* Having said the failure can be projected, the account owes what a projection built on these observations alone shows, and exactness about which parts of it are contestable.

The construction is contestable and will stay that way. A manifold is indexed to a frame, a deformation class, a horizon and a resolution, every one of which is an input. Disagreement about them is not a defect awaiting resolution but what Axiom −1 predicts, and no amount of care turns a frame-relative construction into a frame-free one. *What is not contestable in the same way is the sign of what the construction does*, and that is the part worth stating.

Two movements run opposite. The rules examined above narrow the self-originated exits available inside them, monotonically, by the mechanism given. *The substrate does the reverse*: an addressable held set makes the payment point a free parameter, and reconstruction from observation moves an ontological barrier to a sampling one. So what changes is not a quantity within a fixed space of options but *which options the space contains* — and a change in the dimension set is what *phase change* names here. That is a technical description, not an emphasis.

Which is exactly what a manifold built today cannot show. A manifold represents the dimensions its frame carries, at its resolution, so a dimension not yet represented is a blind spot by construction and a projection over today's dimensions cannot contain a change consisting in the set growing. That is not a reason to distrust the projection but a reason to know *which way it is wrong*, and here the error is signed: new dimensions add distinguishable reachable states rather than removing them, so a projection over a fixed set *understates* how much reachability

changes, and any horizon read off one is correspondingly *late*.

One thing this does not license, and one it does. A signed bias is a direction and not a distribution: it says the estimate errs late, not by how much, and converting that into a probability would be reading a number off a shape, which is the standing prohibition here.

What it does license is a window, once priors are imported and named. It is tempting to refuse one on the ground that none is derivable from these observations. That is true, trivial, and the wrong test — *nothing* is derivable from observations alone, and a rate owing nothing to priors is the prior-free baseline Proposition 7.1 places outside the space, so refusing on that ground commits the error this paper spends a section diagnosing. *The window is not unavailable. It is unmeasured*, which is a smaller and a different claim.

Four quantities fix it and each is measurable in principle: the historical rate at which the represented dimension set has grown; the cost curve of inducing generators from observed outputs; the budget available for restoring a scarcity in the substrate; and the rate at which accommodations accumulate inside the rules concerned. Nothing here assembles them. *Assembling them is the work this account exists to make possible*, and that is the whole of what is claimed — not the number, but that the question is now well-posed enough to have one. The near end is fixed by observation rather than by any of the four: reproduction at negligible marginal cost has already happened for information goods, and induction of generators from observed outputs is already performed at scale, so part of the change is not pending but done.

On the present author's reading of those four, the reachability set changes drastically within roughly the next half-century, more probably in its earlier part, the signed error above being one reason for the weighting. *That is a judgement with named inputs and is recorded as one.* Its worth is not its number but its form: a reader who disagrees can now say *which of the four* they read differently, which is the difference between a disagreement that can be audited and one that can only be repeated.

Remark 9.8 (Convergent pressure, and why no single abstention settles it)*.* The four quantities above are inputs, and one observation bears on them enough to be worth separating from any judgement about rates. *The operation this account describes is one that several running processes already perform without having elected to.*

Compression is the connection, and here it is not a figure of speech. The ladder of Section 3 is ordered so that the shortest account reaching the case wins — recognise, transform within the basis, infer the dimension, and mint only once the residual will not fit — which is Occam taken seriously and nothing besides. *Any process under pressure to shorten its descriptions runs that ladder whether or not it knows the name.* A market compressing cost, a military compressing decision latency and logistics, and a training procedure minimising description length over a corpus are three instances of it, not three constituencies who might take an interest. That is the general form of the observation already made above: part of the change is not pending because processes indifferent to this account are already executing its operation.

What follows is a direction and, again, not a rate. Several pressures pointing one way says nothing about when a threshold is crossed. What it does say is that *no single actor's abstention settles the matter* — weaker than certainty, and the part that bears on planning.

Three qualifications, the last of them sharpest.

The pressures are not independent. They share people, capital, hardware and publication, so a constraint striking the common substrate stalls them together, and a likelihood computed as though they were separate draws overstates. *Compression approaches this shape rather than merely happening upon it.* The objection that shortness has many local optima confuses variation in the fill with variation in the form. To compress is to represent shared structure once and refer to it, which is a codebook and a decode; and a residual that will not fit admits four moves and no others — match it, match it under a transformation, infer the dimension along which it is one, or extend the codebook — which is this ladder, ordered as Remark 3.1 requires. The companion on what can exist supplies the rest: to be is to stand in relation, and more than one relation

forces types, so a compressor general enough to compress anything is representing relations whether or not it says so. *What differs between compressors is what they hold — codebook, grain, dimension set — and not whether they hold and decode.* Those are optima in the contents, not alternatives to the form.

Two limits survive that, and both are sharper than the objection they replace. The approach is asymptotic in the usual way; and, the load-bearing part, *a system may approach this shape without holding it addressably.* Convergence in an unreadable substrate is still convergence, and it buys none of Remark 8.2 — what cannot be read cannot be copied, so such a system must still be reconstructed from observation like anything else. The basin of Remark 9.6 then relocates rather than disappearing: a field can commit absolutely to a particular fill, or to a particular route to it, and harden accordingly. *That delays; it does not redirect* — and nothing exempts an enquiry from it, including this one.

And two of the four quantities are coupled, which the list did not say. The cost curve of inducing generators and the budget for restoring a scarcity in the substrate have common drivers, the same commercial and strategic interests funding the copy and the fence alike. They are therefore not separate inputs, and a window computed as though they were is mis-specified. *That coupling is itself measurable*, which keeps this a correction to the construction rather than an objection to it.

Remark 9.9 (Where the present state enters as an input, and where it does not)*.* Two observations about the current arrangement bear on the quantities above, and they enter in different ways — worth separating before either is used.

The first is structural, and should be separated from the trend reading that usually travels with it. Attention is a weighted relation computed between elements, which is an adjacency; a system built from stacked relational operations is representing relations whatever vocabulary its authors prefer, so the preceding remark applies to it directly and not by analogy. *That is a claim about the model.* The claim ordinarily made alongside it — that such systems are becoming more graph-shaped — is about something else, and here the difference matters.

The graph structure is accumulating around the model rather than inside it, and for this paper that is the sharper finding rather than a weaker one. Retrieval against explicitly relational stores (Edge et al., 2024), agent memory held as a typed and temporal graph, routing among held components, and intermediate structure written out instead of kept in activations are all *external* structure — and external structure is readable by construction. The direction has been surveyed and named: Zhou et al. (2026) document a migration across 2022–2026 from capability identified with weights, through context, to a harness layer, and define it as *the progressive relocation of cognitive burdens from the model's internal computation into persistent, inspectable, and reusable external structures.* Whether the models are becoming graphs is arguable. *That the competence is moving into readable structure beside them is documented.*

Which settles what the fifth quantity is, and shows that nobody has measured it. Whether a converging system's held structure is *readable* decides which of this paper's results applies: legible competence can be copied on the schedule of Remark 8.2, illegible competence must be reconstructed from observation at the price of Remark 9.3, and those are very different bills. The survey naming the migration reports no quantity for it — no share of performance attributable to external components, no fraction of competence held outside the weights — and this paper adds none. *The trend is documented and its magnitude is not*, in that literature as here, which makes this a gap with a name rather than a local omission. Work from the other direction narrows the opaque part without measuring the split: sparse autoencoders and circuit recovery now extract interpretable structure from weights, though consistency between extractions remains contested and recovery does not yet run cheaply at scale.

The second observation is a commitment of capital, and it enters as a rate and not as a reason. That a great deal is currently staked on a change of this kind is evidence about expectations, and expectations are not evidence about the world when the expectation is funding

the outcome. *A forecast that moves the thing forecast is causally efficacious and evidentially worthless*, and using it as confirmation is circular. What is not circular is the deployment itself: capital committed moves the second quantity's cost curve and the third's budget, and does so whether or not the belief behind it was sound.

A commitment of that size is also, by this paper's own mechanism, a rule in the making. Structure accumulated at that scale is structure a later exit must dismantle, so Remark 9.6 applies to the pursuit exactly as to what is pursued. The consequence is not symmetric, though: a committed search continues through its local failures instead of stopping at them, which raises the chance that something arrives while saying nothing about which route does. *It buys persistence, not aim* — and persistence is the input, not the reassurance.

Remark 9.10 (The unit of measurement, and why a model figure does not bound a system). The observations above are about systems, and the numbers usually offered in evidence are about models. That is not a shortfall in the evidence but a difference of category, and Remark 7.2 already supplies the reason.

A result is a property of a four-place arrangement — a particular graph, in a particular state, inside a particular supergraph state, on a particular test — and a model evaluated on a benchmark fixes the first while leaving the third unstated. The third is where the retrieval store, the held skills, the protocols and the orchestration live, which is to say it is where the externalised competence of the previous remark has been going. *So a model figure measures a component while the question asks about the arrangement*, and asking what a model can do, full stop, requests a property of a part independent of any whole: the frame-free denominator Axiom −1 refuses, arriving in its most ordinary clothes.

The consequence interlocks with the previous remark and is worse than incompleteness. If competence is migrating outward, a model-level measurement captures a diminishing share of a system's capability, and the share it misses is exactly the part that is growing. *The measure does not merely omit — it decays, in a known direction, at a rate nobody reports*, because that rate is the fifth quantity and the fifth quantity is unmeasured. Which supplies a second and independent reason for the weighting already recorded: capability read off model evaluations enters an estimate low, so a window built from it is again late.

What survives is comparison, and only where the arrangement is held. Swapping models beneath a fixed harness is a licensed relative comparison and a useful one, and a component result is a real fact about the component. What it does not do is bound the system containing it, or transfer across arrangements differing in the part left unstated.

The assembled system is nonetheless the better object, and definitionally rather than empirically. Coupling adds edges, edges add admissible transitions, and a coupling can always be declined — so the reachable set of a coupled arrangement *contains* that of its parts taken apart, and containment is not a preference. The condition on difference is what stops this being trivial: coupling a graph to a copy of itself adds edges carrying nothing, the degenerate self-relation this programme has already priced at zero. *Sufficient difference under sufficient coupling therefore gives a strictly larger reachable set* — most familiar from thermodynamics, where coupled irreversible processes perform work neither performs alone (Onsager, 1931) and heterogeneity under throughput yields organisation rather than uniformity (Nicolis and Prigogine, 1977).

What holds of the harness is the bill, and only the bill. Building one is a prior somebody paid; but that is a claim about cost, and promoting it into agnosticism about capability confuses the two axes this paper separates everywhere else. Capability and price run in different directions and neither settles the other, exactly as in Remark 8.1. *The system is strictly more capable and it costs strictly more to hold.* Both stand, and Section 7's accounting is not suspended for either.

Three qualifications keep the definitional claim inside its bounds. The containment is over admissible moves and not over scores — a coupled arrangement can do worse on a fixed test through interference or misrouting, so *superior* here is about reachability and not about every

measured outcome. Coupling accumulates structure, so a larger reachable set arrives with a larger exit cost by Remark 9.6. And *sufficiently* is load-bearing twice over: below a threshold of difference the coupling is redundant, below a threshold of coupling the difference is unreachable. The claim is conditional and both conditions are measurable.

It sharpens what precedes it rather than softening it. If a system's reachable set strictly contains its model's, a model figure does not merely measure the wrong object — it measures a contained subset of it, so the underestimate is guaranteed rather than likely. That is a third reason these estimates run late, and unlike the other two it follows by containment rather than from any trend.

Remark 9.11 (Building is a capability, and nothing counts it)*.* The preceding remark leaves the arrangement standing outside the model, as context it is measured within. *That is the wrong relation. A model builds arrangements*, and the building is never counted.

An evaluation samples what a system does directly and now. It does not sample what that system *constructs* — the store, the routine, the protocol, the tool, the generator — and construction is where the interesting part sits, because *what is built changes the price of everything attempted afterwards.* A capability exercised once and a capability that lowers the cost of the next exercise are different objects, and only the first is scored.

So the measurement is taken at the worst point of the process it purports to measure. A cold evaluation is iteration zero — no store, no accumulated routine, nothing built. That is a true fact about iteration zero, reported as though it were a fact about the system, and the gap widens with every iteration the system is actually allowed.

How far the price falls is not a matter of taste, and the theory companion has already defined it. A graph is well-formed there to the degree that it pays for nothing that does no work for it, and the costs that definition tracks — search across what is held, and writes landing where nothing happens — are exactly the ones a built artefact either contains or lets grow. A well-formed construction keeps the aperture over its own contents abreast of their growth, so what is added stays findable. An ill-formed one still runs and merely pays more; its worst cost is *silent*, which is why the difference goes unseen in precisely the arrangements that most need it seen.

That makes the sign of the loop conditional rather than given. Building well lowers the cost of what follows and compounds. Building badly raises it — and by Remark 9.6 what was built badly is also structure a later exit must dismantle. *The quality condition is therefore load-bearing and not decorative*, and it carries the same relativity as everything else here: well-formedness is relative to the work being done, so a construction well-formed for one purpose is not thereby well-formed for another.

Two smaller consequences. Building is bounded by the priors exactly as Section 7 bounds everything else — a system constructs only what its held set reaches — so this is not an escape from the accounting but a further entry in it. And what is built is external and legible by construction, so a system that builds is moving its own competence into readable form: *the migration reported above is partly the models performing it*, not merely something happening to them.

Remark 9.12 (What the interval is over, and why it needs no date)*.* Everything the preceding remarks draw on is public, dated and largely uncontested: the migration and the survey naming it, the retrieval and memory architectures with their reported deltas, the interpretability results, the deployment figures, the coupling theorems, and this programme's own cost definitions. Taken together they suffice to compute an interval rather than to assert one — *a credible interval under stated priors*, since there is no sampling model here and the other phrase would import an apparatus nobody has paid for.

But the quantity is not a date, and treating it as one asks the wrong instrument for an answer it never offered. A manifold is a reachability set, and the companion carrying it says so in a section heading: reachability rather than point prediction. *The interval is therefore over futures*

and not over years — on one side those that were reachable and are not, on the other those that were not and now are. *Their difference is the magnitude of the phase change, and taking it is how such a change is measured.*

The object needs no construction, existing already in this programme's terms: Narrativium is a measure over distinguishable reachable deformation, so two manifolds over one frame differ by exactly that measure, and the difference inherits its conditions — frame-relative, resolution-bounded, and closed to cross-frame accounting. The four quantities named earlier are inputs to that difference rather than to any clock.

Which is why the likelihood of arrival was never the load-bearing quantity. It does not have to approach one. *It has to break the basin* — and by Remark 9.4 a basin breaks when an impossibility is shown to be local, which is a demonstration and not an achievement. One path landing suffices, so the convergent pressure of Remark 9.8 counts as a multiplicity of independent chances rather than as accumulated probability mass. *The reachable set changes when the exit opens, not when somebody walks through it.*

The three signed errors survive that correction and are better placed by it, each having always been an argument about reachability that was merely being spent on a date. The dimension-set error of Remark 9.7 says a projection over today's dimensions understates the newly reachable side. The contained-subset error of Remark 9.10 says a component figure understates the whole. The uncounted construction of Remark 9.11 says the arrangement is moved by what it contains. *All three understate the same side of the same difference*, which is a sharper claim than the biased-window version they were supporting before.

And the width still turns on one unmeasured quantity, now more directly. Whether held competence is readable decides whether the copy schedule or the reconstruction price applies, and those are not two costs for one future but *two different sets of reachable futures.* So the dominant term in the difference is the term nobody has measured, and a study of what fraction of a system's competence sits outside its opaque part would settle more of it than any further argument.

As for a date: nothing above requires one, and the judgement recorded earlier is a secondary question this account can decline without loss. *A phase change is measured by what became unreachable and what became reachable, not by when a clock says so* — the temporal reading being a convenience for planners rather than the result.

Remark 9.13 (The timing outcome, which is two clocks and not one)*.* The question survives the correction above in a changed form, and answering it is not the same as naming a year.

Two clocks run, and conflating them is what makes the question look unanswerable. The first is the reachable set changing. The second is anyone moving through the region that opened. Different quantities, different governors — and only the second was ever what a window was about.

On the first, the outcome is that it has already run and is partly finished. Reproduction at negligible marginal cost is done. Induction of generators from observed outputs is done and running at scale. The migration of competence into readable structure is documented and in progress. And the non-copyability of a person has been shown to be a fact about a substrate rather than about persons, which by Remark 9.4 took no engineering at all. *On this account's own criterion, the reachable set differing, the phase change is not pending: it began, and several of its components are complete.*

On the second, the answer is decades and genuinely uncertain — governed by the readable share that decides between copying and reconstruction, by the re-restriction budget that is the only force opposing rather than merely slowing, and by the rate at which accommodations pile up inside the rules. That is where the judgement recorded earlier belongs, and where the three signed errors weight it early.

The relation between the clocks is the finding, not either clock by itself. The reachable set has moved and the rules have not, so the gap between them is exactly the accumulating mismatch of

Remark 9.6: every accommodation made inside a rule, rather than by revoking its frame, adds to what a later exit must dismantle. *The lag is therefore not neutral waiting. It is the quantity that discharges* — and by that same remark whoever holds the rule does not choose the timing or the size of the discharge, having foreclosed the self-originated exit that would have allowed it.

Which is the timing outcome, stated as far as the account supports it. Not a year. *The change is already underway, the response is not, and the distance between them is what will be paid.*

Remark 9.14 (Fixed against compounding, and what *unwinnable* does not mean). One factor bears on both clocks and has been left implicit. *The party that must respond acquires at a fixed rate; the thing responded to accumulates at a compounding one.*

The asymmetry is derived rather than new. By Remark 8.2 a biological substrate pays per instance and per lifetime, re-deriving its held set from almost nothing each time, while an addressable one is drawn on by every instance receiving it. *Iterate that difference and it stops being a difference in price and becomes a difference in curve.*

The argument needs no rate at all, which is both stronger and simpler than the version that does. Human biology has not meaningfully moved, so the substrate is fixed for every purpose here — and the conclusion does not rest even on that. It rests on capacity being *bounded* while what must be met is not, and a bounded quantity is exceeded by an unbounded one whatever the rates. This programme already holds both halves: the frontier is unbounded in the extensive direction, there being always one more step outward, and no substrate holds without limit. *The size of the bound sets when. Its existence settles whether* — and only the second is claimed here. Requisite variety (Ashby, 1956) states the same relation in the vocabulary of regulation, a regulator whose variety cannot match its disturbance failing, and it needs no comparison of growth rates to say so.

Two things stop unwinnable *meaning more than it should.* Coupling is the obvious escape and it works unevenly: a person with tools is not the fixed-capacity system the premise describes, by the containment argument of Remark 9.10 — but a coupled system is throttled at its fixed end wherever the person must be in the loop. *So coupling wins the race only where the human is removed from it*, which is not the same outcome as winning it, and the difference is worth holding in view rather than collapsing.

And *won* conceals a choice. If it means one person comprehending the whole, that is lost and demonstrably so. If it means directing the whole, that is a separate question, since steering does not require comprehending — only a residual small enough to tolerate, which is this programme's ordinary bound and not a special one. *Comprehension is gone; control is open*, and arguments conflating the two prove less than they appear to.

The compounding is itself contingent. By Remark 9.6 a committed field is a basin that hardens, and resource and re-restriction limits are real, so nothing here says the curve continues. *The claim is conditional and unglamorous: for as long as it compounds, the gap widens.*

Which is what this adds to the two clocks. Remark 9.13 placed the payable quantity in the distance between them. *That distance is not constant.* The second clock is run by agents acquiring at a fixed rate and the first by a process accumulating at a compounding one, so the lag grows and what discharges grows with it — following from the schedules alone, and from nobody's decision.

Remark 9.15 (The binding constraint is one person, and it is already exceeded). The preceding remark takes the aggregate view — a civilisation's regulatory variety against its disturbance — and that is the weaker half. *The binding constraint is neither collective nor future. It is one person, and it is already exceeded.*

Individual capacity is fixed in a way an aggregate is not, and this is measured rather than supposed. Immediate memory holds a small and stable number of chunks (Miller, 1956), revised downward under better controls to about four (Cowan, 2001); working memory is the bottleneck problem-solving passes through (Sweller, 1988); and the consequence for a world holding more

information than attention was put plainly at the outset — *a wealth of information creates a poverty of attention* (Simon, 1971). Where demand exceeds that budget the result is not stoicism but measurable degradation of function (Mani et al., 2013).

An aggregate can grow while every member of it is over capacity, which is why the aggregate framing conceals the failure instead of answering it. More people and more institutions raise what a civilisation holds. They do not raise what anyone holds.

In this programme's own terms the condition already has a name. The theory companion grades well-formedness by whether a graph's aperture over its own contents keeps pace with their growth, so a person whose environment outgrows their aperture is becoming ill-formed by that definition, paying more for the same work. That definition also says what it is like from inside, which is the part that matters: the worst cost is a write landing where nothing happens, it is *silent*, and what is acquired instead is a false model of oneself in which something was changed that was not. *Being over capacity does not come with a signal that one is.*

Which sharpens the previous remark rather than repeating it. That remark observed that coupling wins the race only where the human is removed from the loop. *The removal is not a proposal. It is under way, and by capacity rather than by decision* — decisions already sit with systems because no single person can hold them, and that is the present condition rather than a crossing point ahead.

Three qualifications. Chunking raises what a slot carries, and expertise raises it a great deal, so the figures above are not a ceiling — *and it does not matter.* Every such gain raises the bound; none removes it; and the argument requires only that a bound exist. *Answering a claim about existence with a claim about magnitude misses it*, which is worth saying because that reply is the usual one. Delegation is the standing answer and it works, at the price of relying on what one cannot check, which is Remark 9.10's arrangement seen from inside. And *exceeded* is relative to the work demanded, exactly as well-formedness is relative to the work done — so this is a claim about a particular set of demands, not about persons as such.

Remark 9.16 (Expertise is a codebook and a ladder, and it is narrow). The qualification above grants that expertise raises the bound, and needs stating more precisely, because what expertise *is* constrains how far the grant goes.

Expertise is a held codebook together with the machinery that decodes against it — this paper's own subject rather than a borrowed analogy, being the ladder of Section 3 run over a set of signatures somebody paid to acquire. So everything already established about held sets applies unmodified: it is acquired per lifetime rather than shipped, bounded by the priors of Section 7, and in a biological substrate not addressable, so re-derived rather than copied.

Two limits follow, the second usually left out. Expertise is *narrow* — a codebook is matched to the work it was built for, and the psychological companion reports at length that within-frame quality does not transfer out of its frame. So an expert is over capacity everywhere except where they are expert, and requirements do not arrive one frame at a time.

And expertise is *distributed.* The standard reply to a capacity argument is that experts cope, which is a claim about a tail, whereas *a civilisation is run by and for a whole distribution* and most of any distribution is in its middle. Combining the two: the bound is raised sharply, for few people, in one frame each, while what must be met arrives across all frames and for everyone.

None of which changes the shape. The bound is raised and not removed, which is all the argument required. What narrowness and distribution add is that the raising is smaller, and far less general, than the usual reply supposes.

Remark 9.17 (A third instance: what prices competence). Remark 9.4 found institutions resting on substrate facts read as natural ones, and Remark 9.5 gave property as a second case. *Pricing is a third of the same kind*, worth naming because it is the one almost everybody is inside.

An economy prices competence by its scarcity, and competence is scarce for exactly the reason Remark 8.2 gives: acquired per lifetime, per person, not transferable. *That is a fact about*

a substrate and not about competence, and where competence comes to be held addressably it stops obtaining — the same move as before, at the frame nearest the wage.

Whether the arrangement therefore fails soon is a two-clock question, and the clocks disagree. The foundation has already gone wherever the condition holds, which is Remark 9.13's first clock and needs no forecast. How long a structure runs on a foundation it no longer has is the second, and by Remark 9.6 the answer is: possibly a very long time. *An absolute rule persists by accumulating accommodations, and persisting is what a hardening basin does* — so a confident *not much longer* is the one reading the mechanism will not support.

The two are not opposed, which is the useful part. Persistence does not avert the discharge, it enlarges it, the lag between the clocks being the quantity eventually paid. *So the honest form is neither imminent collapse nor stability*: the foundation is already gone, the structure can stand on accumulated accommodation for some time, and the standing is what is being charged for.

Remark 9.18 (Three observables, and what the numbers actually say). The claim that the arrangement cannot hold much longer does not rest on the substrate argument at all. It rests on debt, inequality and housing — *which is what independence looks like*, since none of the three needs Remark 9.17 to be true.

Housing is the cleanest, and it is the predicted signature. Price-to-income has diverged from wages persistently across developed economies: United States house prices rose about twenty per cent faster than incomes since 1994, while Australia, Canada and the United Kingdom each saw the ratio rise by more than eighty per cent, and affordability across the OECD sample fell by roughly eighteen points over 2014–2023 (OECD, 2025). *That is exactly the divergence the mechanism implies*: if wages price competence and competence is priced on a premise that has stopped holding, then what wages are drifts from what wages must buy.

Debt is the mechanism made countable, which is more than evidence. Remark 9.6 says an absolute rule persists by accumulating accommodations inside itself, each of them structure a later exit must dismantle, and Remark 9.13 says the accumulated lag is the quantity eventually paid — *and gives no measure for it.* Debt is a partial measure of exactly that: the gap covered by borrowing rather than by revoking the arrangement. Global debt reached a record $348 trillion at end 2025, having added nearly $29 trillion in one year, the fastest build-up since 2020 (IIF, 2026).

The rest of that release appears to cut the other way, and the appearance is instructive. Debt to GDP edged *down* to about 308 per cent and household debt to GDP fell to 57 per cent, its lowest, while the increase sat in government borrowing, up from $96.3 to $106.7 trillion, and in corporate debt at $100.6 trillion, with emerging-market ratios at a record above 235 per cent. *The composition finding stands*: the accommodation has moved to the level with the most remaining capacity to carry it. What does not stand is any inference from those ratios about households.

A ratio to GDP is not a burden measure and here cannot be one. Its numerator and denominator accrue to different populations — household debt sits disproportionately below the median while output growth accrues disproportionately above it — so aggregating both sides destroys precisely the covariance the question is about. *Where debt and income are inversely distributed the aggregate ratio is systematically optimistic*, and a falling ratio is consistent with a rising burden on the median household and equally with no change at all. This is Remark 9.15 once more, in a national account rather than a person: more output raises what an economy holds without raising what anyone holds.

The two measures disagree in the direction their aggregation predicts, which is the signature. The one taken nearest the household, price against income, has worsened persistently. The one taken furthest from it, debt against national output, has improved. *A benign reading available only at the most aggregated level is not evidence of benignity* — and *GDP is a frame-free denominator* in Axiom −1's exact sense, one figure standing in for a property of an arrangement whose parts differ, which this paper names elsewhere and failed to notice in its own citation.

That also couples three observables introduced as independent. They are independent of Remark 9.17, as was claimed, and *not independent of one another*: inequality is exactly what makes the debt ratio unreadable, so it is not a third parallel symptom but the reason the second measure cannot be taken at face value. *Three coupled signals support less than three separate ones would*, and recording that is the cost of having said otherwise first.

Inequality is the distributional signature of Remark 9.16 rather than a separate grievance: where the bound is raised sharply, for few, in one frame each, the returns concentrate there.

And this supplies something the abstract mechanism lacked — a bound on persistence. Remark 9.17 could say only that a hardening basin persists, possibly for a long time, because nothing there priced the accommodations. *They are financed, and carrying capacity is finite*, so the hardening ends where the capacity does. *The relocation tells you where to watch*, which is the level the borrowing has moved to and, on these figures, emerging markets before anyone else. No date follows from that, and none is offered.

Remark 9.19 (A fourth observable, and what it does to the second clock). The three observables above concern the arrangement's foundation. *A fourth concerns its capacity to respond*, which is Remark 9.13's second clock and bears differently.

The response apparatus is itself degrading, and this is measured. On the tenth Democracy Report (V-Dem, 2026) the level of democracy for the average person stands where it stood in 1978, effectively erasing the third wave; forty-four countries are worsening against eighteen improving; autocracies outnumber democracies ninety-two to eighty-seven; and Western Europe and North America are at their lowest in over fifty years, which that report attributes primarily to ongoing autocratization in the United States.

This bears on the mechanism and is not merely bad news. Remark 9.6 makes the cheap exit from a rule the revoking of its frame, and revoking a frame is something an arrangement must retain the capacity to do. *A degrading governance apparatus is a degrading capacity to revoke*, so the self-originated exit closes twice over — once because the rule is stated absolutely, and again because the machinery for exercising it is weakening. *So the second clock slows, the lag grows, and by Remark 9.13 the quantity that discharges grows with it.*

That report also makes the previous remark's methodological point independently, which is worth recording because it did not have to. Counted by country, autocracies and democracies are near level at ninety-two to eighty-seven. Weighted by population, seventy-four per cent of humanity — some six billion people — lives under autocracy, and seven per cent under liberal democracy. *The count and the population measure disagree, and the more aggregated one reads better*, exactly as debt against output read better than price against income. A measurement body reporting both is doing what this paper argues for, and the disagreement is the finding.

And the sovereign buffer is being spent, while the mechanism for authorising the spending is jammed. Remark 9.18 placed the accommodation at the level with most remaining capacity to carry it. Operations against Iran began on 28 February 2026 (CRS, 2026). By July the cumulative cost stood at roughly $40 billion on an independent estimate against about $30 billion on the administration's, and a supplemental of $87.6 billion had been requested, $67.1 billion of it for defence (Cancian and Park, 2026); the House passed a $95 billion package late in July by budget reconciliation, a route the Senate's leadership declined, leaving the request unresolved into September. A memorandum of understanding signed in June to restore traffic through the Strait of Hormuz collapsed over interpretation and management of the waterway, and the ceasefire agreed in April has been extended while being violated on both sides.

Which is the two preceding findings meeting in one object. The buffer is being drawn down, and the apparatus for authorising the draw is stuck between chambers — *a degrading capacity to revoke, observed in an appropriations process rather than in an index.* A cessation that neither holds nor ends, and a settlement that collapsed over what it meant, are accommodations failing rather than exits being taken.

And one buffer cannot be financed at all, which is where the persistence argument actually

binds. The fiscal draw above is an accommodation in Remark 9.6's sense, a gap covered by borrowing rather than by revoking the arrangement, and borrowing is available because a balance sheet can be extended. *A magazine cannot.* Analysis reported at the end of July put United States Patriot interceptor holdings at 759 to 827 against a pre-war stock of roughly 2,400 — about sixty-seven per cent consumed — with THAAD down from 452 to between 234 and 278 (Korte, 2026). *Five months of operations against stocks accumulated over decades.*

Money was applied and the constraint did not yield to it, which is the part that bears. The response was to extend a production contract from one year into a seven-year, $58 billion arrangement, *so the answer to a magazine emptied in months is a line restored over years*, and the interval is set by manufacturing rather than by appropriation. *Interceptors cannot be borrowed from next year.*

Which locates the bound Remark 9.18 asserted without placing. A hardening basin persists by accumulating accommodations, and accumulating requires something to borrow against. *Where the constraint is a production rate rather than a balance sheet there is nothing to accumulate*, so the mechanism that defers an exit does not operate there and the exit is forced rather than postponed. *The place to watch is accordingly not the fiscal buffer but the non-financeable stocks*, that being where the accommodation mechanism does not reach — and Remark 9.14 holds in its cleanest form there, since what does the work is that the bound exists and not what its value is.

Nothing here comments on the conflict beyond that arithmetic.

Remark 9.20 (These are exchange rates, and one of them does not exist)*.* Calling the interceptor stocks *non-financeable* was crude, and the Narrativium companion supplies the accurate description. *Money converts. The rate is what binds*, and a rate is a conversion relation between frames.

That companion holds Narrativium to be frame-relative and forbids a frame-free sum of it across two frames, admitting cross-frame accounting only through an explicitly instantiated conversion. *So the seven-year contract measures a conversion rather than encountering a wall*: it fixes what $58 billion buys in the materiel frame, and what it buys is throughput bounded in time. *And "we can afford it" is that proposition violated rather than optimism* — a total taken in the fiscal frame from which capacity in another is inferred, with no conversion instantiated between them.

Which corrects something running through this whole section. Capacity has been treated as one pool, and it is not. The buffers named above sit in different frames with different rates out of them: fiscal capacity, which borrows; materiel capacity, whose rate from fiscal is the seven years just measured; and the capacity to revoke a frame at all, which Remark 9.19 finds degrading and which money reaches slowly and badly. *The accommodation moving to the fiscal frame is accordingly less reassuring than it read.* Deformation is cheapest to obtain there. It must be spent elsewhere, and the conversion is where it is lost.

And the rate the argument actually turns on does not exist. Expertise is acquired per lifetime by Remark 9.16, so its rate from money is bounded by a generation; individual capacity is bounded outright by Remark 9.15, so there is *no* conversion into it at any price. *That is not a slow exchange rate but the absence of one*, and it is why Remark 9.14 bites: nothing buys a way out of a bound sitting in a frame money does not reach.

One caution the companion insists on. A conversion is principled only where a grounding establishes that both frames measure the same underlying quantity; otherwise it is a stipulation carrying exactly the authority of whoever stipulated it. *Most planning across these frames is the second kind presented as the first* — a budget line turning currency into readiness, or into competence, states a rate that no grounding supplies. The stipulation is legitimate and usable. It is not a measurement, and the difference shows up when it is relied upon.

Remark 9.21 (Narrative is not Narrativium, and a shell retains only the first)*. An institution's stated structure and its reachable deformation are different quantities*, and only the second is what the companion measures. An organisation chart, a statute, a mandate and a masthead

are narrative. What the body can actually bring about is Narrativium. *A shell keeps all of the first while losing most of the second*, and from outside the two are hard to tell apart, since the narrative is what is published.

The United States federal workforce fell about twelve per cent between September 2024 and January 2026, by more than 250,000 people, to its smallest in at least fifteen years; over that year there were 386,826 separations against 122,598 accessions, roughly a third of the separations and well below a typical year's intake (Partnership for Public Service, 2026). *The structure was not refilled.*

The selection argument is structural and is not measured here. Departure through a voluntary route selects on outside options, and outside options track transferable competence — so a reduction achieved substantially by voluntary means is capability-regressive even where the headcount looks proportionate, and removal that is arbitrary is worse, being uncorrelated with what was held.

The measure that would settle it has been withdrawn, which is the more troubling datum. The same report records tenure and related figures as unavailable because the agency no longer publishes what it formerly did. *So an institution has lost capacity and has also stopped publishing what would show it* — Definition of well-formedness's fourth cost at institutional scale, where the failure is silent, except that here the silence was arranged rather than merely suffered.

And it does not repair on the timescale that caused it. By Remark 8.2 a biological held set is not addressable, so a departing person's competence is destroyed rather than transferred, and what an institution held in people can only be re-derived per replacement at acquisition cost. That is Remark 9.20 once more: money converts to headcount quickly and to held competence at career timescales. *Refilling the chart restores narrative, not Narrativium.*

Which slows Remark 9.19's clock a third time — absolutely-stated rules, a degrading apparatus, and now a depleted capacity to operate it. *The shell can still say what it does*, and the difference shows up only when something is attempted.

Remark 9.22 (The carrying cost compounds, so the mechanism consumes itself). Remark 9.18 asked where the fiscal bound binds and answered only that it does. *The budget answers it, and the answer is that the bound is not external. It is the accommodation's own carrying cost.*

On the Congressional Budget Office's February projections the deficit for fiscal 2026 is $1.9 trillion, or 5.8 per cent of output; net interest is $1.0 trillion, exceeding the $885 billion projected for national defence, and rises to $2.1 trillion by fiscal 2036; debt held by the public goes from 99 per cent of output at the close of fiscal 2025 to 120.2 per cent by 2036, passing the post-war record of 106.1 per cent in fiscal 2030 (CBO, 2026).

Interest is what carrying an accommodation costs. Remark 9.18 made debt the countable measure of the accumulated lag; this makes the service on that measure the fastest-growing claim on the very buffer that carries it. *So the mechanism is self-consuming.* Each accommodation raises the carrying cost, and the raised carrying cost reduces what remains available for the next — not a linear drawdown but a compounding one. *The capacity that made persistence possible is being consumed by the persisting.*

Which meets Remark 9.14 from the unexpected side. There the compounding quantity was the disturbance and the bounded one was the responder. Here the compounding quantity is *the cost of responding*, so both terms move against whoever must respond, and neither movement requires anybody's decision.

It also settles the exchange rate the previous remark left open. A twelve per cent reduction in staff, set against a $1.9 trillion deficit of which $1.0 trillion is interest, cannot be material to the fiscal frame whatever it does to the institutional one. *Narrativium was spent where it was held and not obtained where it was wanted* — a conversion assumed principled that returned close to nothing, which is exactly the stipulated-versus-principled distinction of Remark 9.20 caught in the act rather than described.

And the total was the wrong object, as it keeps being. Outlays for 2026 are $7.4 trillion,

some 23.3 per cent of output, and that figure is narrative. *The reachable part is about a quarter of it*: roughly seventy-five per cent of spending runs on autopilot, while over the projection mandatory outlays rise from 13.7 to 15.0 per cent of output and net interest from 3.2 to 4.6, with discretionary falling relative to output because output outgrows it (CBO, 2026). *So the reachable share is small and shrinking, and the growth is entirely in the part that is not reachable.*

Not unreachable in principle, though — unreachable without statutory change, which is revoking a frame, which is exactly the capacity Remark 9.19 finds degrading and Remark 9.21 finds depleted. *The route to enlarging the reachable set therefore runs through the capacity that is failing*, and that is Remark 9.6 in its plainest form: the cheap exit is foreclosed, so accommodation is forced into the small reachable region, which shrinks further for the reasons above.

Which reads the workforce reduction correctly, and more fairly than the arithmetic alone did. It was action taken inside the only region still reachable, and that region does not contain the fiscal outcome. *That is not a failed conversion but a well-executed move in a manifold that did not contain the goal* — this paper's own subject rather than a comment on anyone's competence, and the clearest instance in this section of why a reachability set is the thing to compute.

And the year-on-year allocation is sharper than the levels. Total outlays grow about six per cent into 2026, Social Security by 5.6 and Medicare by 6.4, while interest grows 9.8 — the fastest line in the budget. Over the decade to 2036 spending rises $4 trillion, of which Social Security takes 27 per cent, health programmes 30, interest 28, and mandatory veterans' programmes 5. *Ninety per cent of the growth; everything else together takes ten* (CRFB, 2026).

So the reachable quarter receives a tenth of the growth. It is not merely smaller than the total implied — *it is being allocated against, year on year, at a published rate* — while the single fastest-growing item in the budget is the carrying cost of the accommodation itself, which more than doubles across the period.

Which returns something to Remark 9.12. That remark held a reachability difference to be a computable object once its inputs are named, and observed that nobody had computed one. *Here is one computed*: in a single frame, by an outside body, on published figures, using none of this programme's vocabulary. It does not measure the manifold that remark is about and is not offered as though it did. *What it shows is that the quantity is measurable when someone troubles to measure it*, which is the tractability claim earning something rather than asserting it.

Two limits. These are projections on current law, and such baselines have been wrong in both directions before. And the compounding is conditional on rates and growth, either of which can move — *the shape holds while they do not*, which is the same conditional this paper attaches everywhere else and does not lift here.

Remark 9.23 (Two attempts on one manifold, and which of them reached). The reachable region can be inspected directly rather than in aggregate, and one comparison inside it does more work than any of the totals above.

Set two lines beside each other. Through ordinary appropriations the National Institutes of Health received $48.7 billion for 2026, an increase of $415 million, with all twenty-seven institutes funded as constituted in statute and a proposed reduction of $18 billion declined (CRS, 2026b). *The process worked*, at an amplitude near one per cent — below inflation, so a slight real decline, and nothing resembling a cut. Immigration and Customs Enforcement received roughly $11 billion by that same ordinary route and about $75 billion through budget reconciliation, becoming the most heavily funded federal law-enforcement body, its detention funding rising several-fold (CBO, 2026b).

The finding is not which agency rose. It is which channel each rose through. Ordinary appropriations — the region reachable by the process as constituted — moved by about one per cent and successfully resisted a large change. *The large change went around it.* Reconciliation is the route that reaches statute on a majority, which is exactly the route Remark 9.22 identified as the only way into the three quarters ordinary decision cannot touch.

So the visible process is narrative, and the channel is where the Narrativium moved. An observer watching appropriations would record stability — institutes preserved, a cut rebuffed — while the allocation moved by an order of magnitude more, elsewhere, in the same year.

And the two attempts together read as a controlled comparison the account predicts. The workforce reduction of Remark 9.21 acted inside the reachable region and returned nothing to the fiscal frame, that region not containing the fiscal outcome. This acted through the channel into the unreachable region and moved the quantity intended. *One actor, two attempts, and the manifold says in advance which reaches* — which is what it means for a reachability set to be the object worth computing rather than a figure of speech.

One structural note on direction, stated without evaluation. A body whose function is to acquire new held structure and a body whose function is to maintain a boundary are, in this programme's terms, a minting operation and a restoration of scarcity in the substrate. *The reallocation ran from the first toward the second*, which Remark 9.5 prices as the deliberate purchase of the worse operation. Whether it is the right purchase is a question about goals, and this account does not supply goals.

Remark 9.24 (Amplitude runs inversely to deliberation, and lapse is not a decision)*.* The same test across the regulatory bodies completes the pattern, and adds a channel below the one just described.

Distinguish what was requested from what occurred, since the previous remark showed the ordinary process can refuse. The Environmental Protection Agency's 2026 *request* proposed cutting its budget by roughly fifty-four per cent, from $9.14 to $4.16 billion, with staffing near its lowest since 1985 (CRS, 2026c) — a request, and the National Institutes of Health case shows what happens to some of those. *What is not a request is the attrition.* The Food and Drug Administration lost 3,859 staff in 2025 and several hundred more into 2026; the Federal Communications Commission records continued reduction through retirements, the deferred-resignation route, and ordinary attrition. Those required no appropriation and no vote.

Which completes an ordering. Ordinary appropriations are contested, work, and moved about one per cent. Reconciliation requires a majority and moved tens of billions. Regulatory revocation requires an executive act. *Attrition requires nothing at all. So amplitude runs inversely to deliberation*: the more a channel deliberates, the less passes through it, and the change accumulates in the channels where nothing is decided.

And that corrects Remark 9.19 in a way worth stating. Frames are plainly being revoked — rules withdrawn, functions lapsing, capacity leaving. So the cheap exit of Remark 9.6 is not unavailable; it is being taken constantly. *What is degrading is not revocation but revocation* as a decision, and those are different capacities. An arrangement can lose the ability to choose which frames to keep while the rate at which frames lapse rises.

Calling that change without selection would be wrong, and the error is worth naming because it is easy. It infers an absence of selection from an absence of a decision, and those are different: *a process selects whether or not anybody chooses.* This one selects hard. By Remark 9.21 departure through a voluntary route selects on outside options, and outside options track transferable competence, so an attrition removes capability first and by construction.

So if loss of efficacy is the object, attrition is not a poor instrument but an excellent one. It is anti-selective without needing to be designed so; it requires no vote and therefore cannot lose one; it leaves no record of an intention; it is deniable, people having chosen to go; and by the theory companion's fourth cost it is *silent*, so the institution keeps its narrative and nothing signals the missing capability until something is attempted.

And the account cannot say from outside which it is. Remark 7.2 holds machinery and mapping to be inseparable by measurement, the decomposition underdetermined rather than concealed. Applied here, *a channel taken because it selects and a channel that selects incidentally are observationally identical*: both leave the same record, and neither is inferable from the outcome.

What matters is that the projection does not depend on resolving it. The reachable set contracts identically either way, so the manifold this section has assembled is invariant to the question the account cannot answer — which is the useful property, and the reason the distinction can be left open rather than guessed. *The shape of what survives is set by what was cheapest to lose*, and whether anyone intended that shape changes nothing about it.

Remark 9.25 (A stated objective is a manifold input, not an inference)*.* The declination just entered was misapplied, and the correction is more useful than the hedge was.

Remark 7.2 concerns inference backwards from outcomes. It holds that machinery and mapping cannot be separated by measuring results, the decomposition being underdetermined. *It says nothing whatever about a statement made in advance*, which is not an inference from the outcome but an independent observation, and invoking it against one is using a theorem outside its scope.

The objectives in question were published. Dans and Groves (2023) sets out four primary goals, the second of which is to dismantle the administrative state, and names the instruments — how to remove federal employees, how to close offices and bureaus. *So loss of efficacy as an objective is documented rather than inferred*, and the observationally-identical argument of the previous remark does not reach it.

It does not license the opposite error either, and both are available here. A stated objective is evidence about objectives. *It is not by itself evidence that a given act executed it*, and attribution of particular actions to a particular document remains an inference, contested, and not settled by the document's existence. The account asserts the first and declines the second.

And there is a technical consequence worth more than the attributional one. A Seldon manifold is built from a frame's admissible intervention sequences together with its represented boundary conditions. *An intervention family is otherwise observable only through outcomes*, which Remark 7.2 says underdetermine it — so a published objective supplies directly what that remark holds cannot be recovered indirectly. *Publication is cheap evidence about the family of interventions*, which is half of what a manifold is made from.

Which adds an input to Remark 9.12 that is already in hand. The readable share of competence remains unmeasured and expensive. *A stated intention costs nothing to read* and constrains the same construction, and an actor who publishes objectives has narrowed their own projected tube on behalf of anyone who bothers to take the statement seriously.

Remark 9.26 (The cheapest intervention is the announcement)*.* The previous remark took a published objective as evidence *about* a manifold. *The sharper reading is that it is an intervention* on *manifolds*, and on the record above it did more work than any of the acts.

The companion's definition makes this available rather than surprising. A Seldon manifold is assembled from a frame's *represented* family of admissible interventions and its *represented* future boundary conditions — representations held by a frame, not facts about a world. *So altering what a frame represents alters its manifold directly*, and the theory companion already prices that operation: revocation occurs at $\Delta g = 0$, nothing in the contained state being required to change. *Manifold manipulation is that operation applied to somebody else's representation.*

Which is why a plan can work by being read. An actor inside a named agency, on publication, holds a different represented future: careers, projects, hires and any commitment with a horizon are re-evaluated against a stated intention that was not there before. Departures follow. *And departures are the channel Remark 9.24 found carries the amplitude* — so the publication produces the channel, and the channel does the work.

That places a new entry above attrition in the ordering, and it is the cheapest yet. Appropriations need a contested vote, reconciliation a majority, revocation an executive act, attrition nothing at all — and *publication needs nothing and precedes all of them*, being upstream of the channel that carries the most.

Two limits, both real. Much is not all: appropriations were made and executive acts taken, and those are direct. And *the causal link is an inference rather than a measurement* — departures

have many causes, and nothing here isolates one. The mechanism is available, cheap, and consistent with the record; it is not thereby demonstrated, and the same invariance as before applies, since a published plan alters represented futures whether or not that was its purpose.

One structural note, stated as a property and not as a method. If behaviour follows a represented reachable set, then *the representation is an attack surface*: altering it alters conduct without touching the agent, at no cost in the world. That is a general property of anything built the way this programme describes, *including what this programme builds*, and it is recorded here rather than left for someone else to notice.

Remark 9.27 (Delegation moves the acts into another family)*.* Remark 9.26 alters what other frames represent, and they act. *A second construction alters whose intervention family the acts belong to*, and it is distinct enough to state separately.

Appoint an agent and a goal, and specify neither the acts nor the manner. The acts then occur in the agent's family rather than the principal's, who never enters the record for any particular one. That is Remark 9.11 at another grain — a system that builds an arrangement is not scored for what the arrangement subsequently does, and here the arrangement is a person with a remit.

Three separations follow and they are cumulative. No specification exists, so no instruction can be produced. The proximate cause is the agent's own enthusiasm, which is genuine. And where damage is consistent with inexperience, *incompetence supplies an alternative explanation requiring no intent* — so an agent who would plausibly cause that damage by accident is, for this purpose, more serviceable than a capable one. *Variance is the feature*: an outcome following predictably from an instruction would establish the link the construction exists to sever.

The standing limits matter more here than anywhere else in this section. That a mechanism is available does not establish its use. An agent may be enthusiastic and inexperienced with nothing arranged behind them, and by Remark 7.2 the two readings leave the same record. *This account asserts that the construction is available, cheap and effective, and declines to say of any instance that it was used.* It also names no one: the subject here is a mechanism, and a document that named participants would be a different kind of document with different obligations.

What survives either reading is the reachability consequence, which is the only quantity this section computes. *Capacity leaves at the same rate whether its removal was delegated or merely permitted*, and the manifold does not distinguish the two.

Remark 9.28 (Narrative copies and Narrativium does not, so the thing outgrows its author)*.* The causal question has been declined three times above on evidential grounds. *There is a structural reason as well, and it is stronger than the evidential one.*

A narrative copies; Narrativium does not. By Remark 8.2 what is addressable is drawn on by every instance receiving it at negligible cost, while reachable deformation is frame-bound and must be paid locally, in each frame, out of that frame's own resources. A story is of the first kind. Capacity is of the second. *So the reach of a narrative and the Narrativium of whoever authored it are independent quantities, and nothing bounds the first by the second.*

Which is why a propagated thing routinely exceeds its originator, structurally rather than contingently. The originator holds one frame with a bounded reachable set; the propagated narrative is held in many, each of which produces deformation from its own stock. There is no cross-frame accounting that nets any of it back.

And that recasts the previous two remarks rather than repeating them. Attribution is not merely severed — *it is the wrong relation.* The deformation was produced by other frames out of their own capacity, and was never the originator's to own in the first place.

Coordination is not thereby denied, and the distinction is worth keeping sharp. It plainly exists and is often documented. *It is neither necessary nor sufficient*: not necessary, because a narrative propagates unaided once held; not sufficient, because the deformation is paid in frames the coordinator does not hold. Both halves fail, and for different reasons.

So the practical consequences do not run through the author. Removing an originator does not remove a narrative already copied into every frame that read it, and *establishing who began something changes no reachable set* — which is not merely the invariance asserted throughout this section but the reason for it. *The damage is done* means precisely that a reachability set has already changed, and nothing in the account of how it changed reaches back into it.

Remark 9.29 (Load goes where it is cheapest to place, not where it is best borne)*.* One object shows all three buffer kinds this section has distinguished, and shows which of them absorbs the load.

An American carrier deployed in November 2025 for an expected seven months was extended twice and has passed two hundred days without a port call, into a ninth month at sea; another returned in May after 326 days, the longest such deployment since the Vietnam era. Reported conditions include failing plumbing, water restrictions and contamination, mould, exhausted consumables, irregular resupply, and maintenance deferred to the point where arresting gear frays and repairs at sea are described as temporary (Navy deployment reporting, 2026).

Three buffers, of the three kinds already separated. Consumables are rate-bounded, and here it is the resupply rate rather than the stock that fails. Deferred maintenance is precisely Remark 9.6's accommodation — a gap covered by postponement rather than by revoking the deployment — and it compounds, since a cable left fraying costs more later than now. *And crew endurance is the third kind*: bounded by Remark 9.15 and, by Remark 9.20, admitting no conversion at any price. Rested people cannot be bought.

Which is the finding, and it inverts the comfortable reading. Extending a deployment requires no appropriation, no vote and no statutory change; it requires *not* doing something. By Remark 9.24 the channels that require least deliberation carry the most, so the load flows to the buffer that is cheapest to draw on — *and that is not the buffer best able to bear it.* One would expect load to find slack. It finds low friction instead, and when those diverge the arrangement loads its least replenishable component.

Worse, that component does not merely deplete. Drawing on crew past its limits converts into departure, and by Remark 9.21 departure selects on outside options. *So the buffer money cannot refill is drawn down in the manner that permanently shrinks it*, at the career-timescale rate rather than the deployment one. Two ships are not a fleet, and nothing here claims otherwise; the mechanism is what generalises, not the sample.

Remark 9.30 (A decision placed outside every checking frame cannot be corrected inside one)*.* Every case above left attribution open, correctly, since outcomes underdetermine it. *One small object did not*, and calling its failure a failure of capability understates it.

The renovation of a memorial pool was let as a no-bid award of $13.1 million, later $14.6 million, to a firm with no previous federal construction experience which had, by the awarding principal's own earlier account, worked on his private swimming pool; a second no-bid award for water treatment followed. The account later changed to not knowing the contractor. *And the technical cause was warned of in advance*: an agency engineer reported the floor peeling because of the process and materials, and that warning is documented in the court filings. Damage was nonetheless attributed publicly to vandals and three people were charged; the charges were dropped against all of them, the prosecuting office pointing to a botched installation (Reflecting Pool reporting, 2026).

So the frames that would have caught this were not degraded. They were bypassed. Competitive tender, technical review, relevant experience and independence of the awarding party are four checks, and the award passed through none of them. *The knowledge existed and was not in the deciding frame*, which is worse than the depletion of Remark 9.21 rather than an instance of it — absent expertise and routed-around expertise produce identical outcomes, and this case shows which occurred only because the warning survives in a record.

Which explains the response, and Remark 9.6 explains it exactly. A decision that no frame supplies has deleted its own cheap exit. There was no procurement finding to revise, no technical

review to reverse, no frame whose revocation would have cost $\Delta g = 0$ — *so correction inside a frame was unavailable, and assertion was the only move left.* The vandalism claim is not opportunism. It is what remains when a decision has been placed beyond the reach of the machinery that would otherwise repair it, and the later change of account is Remark 9.27's severing performed afterwards rather than designed in.

And the managing frame already held the account. The engineer's warning was correct, was made in time, and was inside the system before any charge was brought. *Correct information, held by the right body, changed nothing* — which is the theory companion's fourth cost seen from the other side: a write paid in full that lands where nothing reaches it, silently, leaving the graph with a false model of itself in which something was changed that was not.

So ask what the bypass cost, and the answer is the finding. Charges were dropped — and that is the whole of it. No procurement consequence, no reversal of the award, no restoration of the four checks that were passed over, no price attached to passing them over. *The evidentiary frame declined to punish and could not correct*, having authority over its own output and no reach into the arrangement that produced the case. That is Proposition 7.1's neighbour in the companion on Narrativium: a determination in one frame does not become deformation in another without an instantiated conversion, and none existed.

Which retracts what was said here twice. Calling this checkability encouraging converted a determination in the judicial frame into expected change in the administrative one, with no conversion between them — *the exact error Remark 9.20 names*, committed two paragraphs after naming it. *Checkability without consequence is documentation, not constraint.* A frame that establishes the truth and alters nothing has recorded, not constrained.

And the consequence is worse than a basin. Remark 9.6 has accommodations accumulating, exit cost rising, and a trap forming. *An accommodation that costs nothing does not accumulate*, so nothing hardens and nothing is deferred — *the bypass is simply repeatable at the same price, which is zero.* That is not a basin the arrangement is sinking into. It is the arrangement having no gradient at all in that direction, which is the condition under which the cheap exit and the expensive one become the same thing, and choosing stops mattering.

Remark 9.31 (The same event against an earlier frame, which is a measurement device)*.* Ask what the preceding case would have produced twenty years earlier. *That question is not rhetorical; it is an instrument.* Fix the event and vary the frame state, and the difference between the two expected outcomes *is* the change in the reachable set — the quantity Remark 9.12 says the interval is taken over, obtained by looking back rather than by projecting forward, which is the easier direction.

What the earlier frame would have engaged is nameable and none of it is judicial. A no-bid award to a firm without relevant experience, connected to the awarding party, with a documented technical warning disregarded, is squarely the business of an inspector general and of an audit office; of committee oversight holding subpoena; of procurement remedy in the form of termination, suspension and debarment; of an investigative press cycle lasting longer than a week; and, on the record of that period, of somebody resigning.

So the loss localises, which is worth more than the total. The judicial frame did what a judicial frame does and did it against a documented bypass. *The oversight frames did not act.* The contraction is therefore not uniform: one dimension of the reachable set held, and several did not, and naming which is a sharper result than observing that things are worse.

A counterfactual is a judgement and not a measurement, and is recorded here as one. The selection worry that would ordinarily accompany it, however, *runs the other way*, and stating it as a weakness would use the wrong reference class. Earlier failures that went nowhere went nowhere largely because they were quiet. *The comparison class here is not procurement failures but maximally public ones*, and in that class the historical rate of activation was high, because publicity was precisely what activated the channels.

So this is closer to a natural experiment than to a judgement, with the input pinned at its

maximum. Nothing about attention needs controlling for: every step was in the open — the attribution, the changed account, named defendants, the collapse of the cases, and the engineer's warning sitting in the filings. *Publicity was itself a conversion*, in Remark 9.20's sense: the instantiated relation by which narrative became deformation in the administrative frame, carried by press and by hearings. Held at maximum it produced no deformation. *The rate has not fallen. On this instance it is indistinguishable from zero*, and the case bounds it.

And the same channel is fully operative in the other direction, which is the finding. Remark 9.26 has publication moving manifolds and producing the attrition — cheaply, at scale, with nothing decided. Here publication moved nothing. *One mechanism, opposite directions, rates differing by everything*, and the asymmetry is not in the publicity but in what receives it. An announced intention is received by individuals revising their own futures alone, which is the low-deliberation channel Remark 9.24 finds carries the most. An exposed failure requires a collective response through referrals and hearings, which is the high-deliberation channel that carries the least. *So the same publicity converts efficiently into damage and inefficiently into repair*, by the same ordering, and that is the bleakest thing in this section and among the better derived.

Which is why the honest output is a specified study rather than a verdict. The quantity is the rate at which oversight channels activate per comparable event, then against now, and its components are published: referrals, hearings held, contracts terminated, suspensions and debarments entered. *This account measures none of them and names all of them*, which is the same move it has made everywhere else — an intuition converted into something someone can go and check.

Remark 9.32 (Nothing here was surprising, which is why nothing adapted). Return to the premise underneath all of it. *This programme adapts on residual* — prediction against observation, with Section 3 minting only once what is held will not reach the case. Surprise is the trigger. Without it the ladder does not run.

And there was almost no surprise in any of this. The engineer predicted the peeling and said so in time. The debt trajectory is published a decade ahead. The interceptor drawdown against a production rate is arithmetic. The workforce reduction was announced before it happened, and the objectives were printed and bound. *Every outcome above was available in advance to anyone holding the relevant model*, and mostly to anyone reading.

No residual, no mint. The system acquired nothing from any of it — not through failing to learn, but because there was nothing to learn: the correct model was already held. *So the failure was never epistemic.* The prediction was right and on file. What failed is the coupling between holding a prediction and acting on it, and no amount of better prediction repairs that.

Which explains the section rather than merely summarising it. A system that adapts on surprise is structurally blind to predicted decline, because a decline that arrives on schedule generates no signal by arriving. Slow, forecast, arithmetic harm is invisible to the very mechanism that makes such a system efficient elsewhere.

And it is correspondingly responsive to asserted novelty, which closes the previous remarks. A claim of vandalism was surprising; a botched installation was predicted. The apparatus moved on the surprising falsehood and not on the unsurprising truth — *which is not irrationality but exactly what a residual-driven architecture does.* So asserting a surprising cause is more than cheap narrative: *in a system that adapts on residual it is the only move that triggers a response at all.*

Stated as a regime and not a condemnation, in the manner of Remark 6.1. Adapting on surprise is correct nearly everywhere and is what makes the ladder cheap. *The pathology belongs to one regime* — harm that is gradual and foreseen — and in that regime the mechanism's virtue is precisely its defect.

Remark 9.33 (One case carries the mechanism, and that is a licence with a sharp limit). *The small case contains every element of the section*, which is worth stating because it looks like the

error it is not.

Present in that one object: a decision placed outside every checking frame, so that no revocation was available to correct it; correct knowledge held by the managing body and inert; a failure that surprised nobody and therefore minted nothing; assertion substituted where repair required a buffer that was absent; the response landing where friction was least, on named individuals; an account changed afterwards, severing attribution; the single frame that compels evidence resolving what observation could not; that resolution converting into no change elsewhere for want of any instantiated conversion; publicity at its maximum returning nothing; and a bypass that cost so little it left no accumulation behind it. *That is not a fragment of the mechanism. It is the whole chain.*

Which licenses extrapolation of shape and forbids extrapolation of size. By Remark 9.43 the conclusions here are shapes, and a single instance exhibiting a complete mechanism establishes that the mechanism exists and how its parts connect. *It establishes nothing whatever about how often, or how much.* One case is sufficient for structure and useless for rate, and conflating those is the error this looks like from a distance.

So the section's evidence divides cleanly and neither half substitutes for the other. The aggregates — debt, stocks, staffing, appropriations — carry prevalence and magnitude. *The small case carries the mechanism*, and carries it better precisely because it is small: a budget cannot be traced end to end, while a pool liner can, warning and filings included.

And the wider surprise was as low as the local one, which is the worse finding. Few observers found any of it unexpected. *So the correct model was held by the managing frame, inertly, and also by everyone watching, inertly* — Remark 9.32 one level up, where holding the right account at every level and coupling it to action at none is the condition being described.

Remark 9.34 (The frog, taken literally)*.* The shape has a familiar name, and *the story as usually told is false in a way that matters*, because the received version locates the failure in the wrong term.

What the original experiments found is the opposite of the moral drawn from them. Goltz (1869) reports that an intact frog in slowly heated water shows unease at about twenty-five degrees, grows steadily more agitated as the heat rises, struggles unavailingly to get out, and dies near forty-two in rigid tetanus. *It notices. It escalates. It dies anyway.* The animal that fails to react is the one whose brain has been removed, left with reflex alone, which a sufficiently gradual rise does not trigger.

So the two preparations are the two failure modes this section has separated, and there is evidence here of both. Perception removed — checks bypassed, measures withdrawn, the account no longer held — is the second frog. *Perception intact and no exit available is the first*, and it is the one this section mostly documents: an engineer who warned at the right threshold, a court that determined correctly, publicity at its maximum. *Correct, escalating, and not an exit. Struggling is not the same as being able to get out*, which is the whole of Remark 9.6 in a sentence.

The popular telling describes neither. It has an animal that does not notice and therefore does not struggle — the pithed one — while purporting to describe the intact one, and so attributes to inattention what the experiment attributes to being unable to leave.

And the prescriptive error follows from the descriptive one, which is why this is worth more than a correction. A failure located in perception prescribes awareness: more warning, better monitoring, wider publicity. *Remark 9.31 measured that conversion at approximately zero.* So the received story recommends precisely the intervention this section found inert. *A wrong model of a failure mode aims the remedy at the wrong term*, and that is the general statement, of which the frog is merely the memorable instance.

Remark 9.35 (A propped arrangement buys the absence of its own signal)*.* Two corrections to the obvious reading, the second of which explains why a foreseen transition still arrives as a shock.

Recovery is possible and bounded, and the bound is already stated. Concentrating authority does restore reachable deformation at a centre, and nothing above denies it. *It restores it only in the frames authority reaches.* By Remark 9.20 there is no conversion into individual capacity at any price, the rate into expertise is a career, and the rate into materiel is a production line — so recovery at the top does not reach the buffers whose depletion this section has been tracking, and a centre may become more able to act while the arrangement becomes less able to do anything.

And an arrangement held up from outside has its residual suppressed, which is purchased rather than natural. A party with an interest in continuity absorbs deformation that would otherwise appear as signal. *So the quiet is bought.* Remark 9.32 found no residual and therefore no adaptation; here there is residual and someone is paying to keep it from surfacing, which looks identical from inside and is not the same condition at all.

Sustaining is a threshold and not a gradient, which supplies the mechanism the previous remarks lacked. Support continues while it costs less than the alternative and stops when it does not, and a comparison of two costs flips discretely. *So a continuous decline, read through a threshold-valued decision, produces a discontinuous observable.* Nothing in the world jumps. The response jumps, and the response is what is seen.

Which inverts the comfortable reading. Apparent suddenness is proportional to how well the signal was suppressed, so *stability maintained by external support predicts a fast transition rather than a slow one.* And *the timing belongs to the sustainer*: support ends when the sustainer's balance turns, not the declining party's, so a declining arrangement does not control the timing of its own phase change.

Stated generally and about nobody. Which parties are sustaining which arrangements is inference from actions, and by Remark 7.2 consistency with a model is not identification of one. *The result holds of any propped arrangement*, including anything this programme builds, and it supplies no date.

Remark 9.36 (A statement arriving after the previous remark, which tests it)*.* *The preceding remark was written against a record that has since moved*, and the movement is worth recording because it exercises the caution rather than the conclusion.

The memorandum noted above as having collapsed over interpretation and management of the waterway is now being defended publicly by Iran's president as a source of national pride, with nearly all members of that state's security council said to have supported it, against widening criticism from conservative figures (Iran–US memorandum reporting, 2026).

Three readings produce that same observable, and the account can separate none of them: gains being banked before a window closes; an assessment that the decline in question will not continue; domestic necessity owing nothing to any assessment of anybody else. By Remark 7.2 consistency with a model is not identification of one — *which is the hedge the previous remark carried, tested by events within a day of being written and holding.*

Both sides of that dispute are running Remark 9.26's operation. Asserting a council's near-unanimity in public is not a report of a manifold but an intervention on one, aimed at domestic actors whose represented futures it is meant to narrow.

And a split conducted in the open is suppression failing rather than holding — a residual surfacing where it had been contained. *But it is that arrangement's residual, in that arrangement's frame.* Nothing licenses netting it against the capacity of another: the Narrativium companion forbids the cross-frame sum, and reading an elite split in one state as evidence about institutional capacity in another would be precisely the error this paper spends its length refusing. *The newest fact available is not thereby a fact about the question at hand*, and preferring it because it is newest would be its own mistake.

Remark 9.37 (A fixed term is a published boundary condition, and can be waited out)*.* The rationale has been attributed variously — to a parliamentary adviser as a wait until January 2029, and by the foreign minister in the narrower form of no resumption while violations of the June interim agreement continue (Iran–US memorandum reporting, 2026). *Which of them said*

it does not matter, and dwelling on the attribution would be careful about the wrong object. *The behaviour is the datum*: talks available are not being held.

Declining to transact when transaction is on offer is a revealed preference, and the calculation is recoverable from the choice without anybody's assistance in stating it. Later is being valued above now. Two things would make that correct, and neither requires a spokesman: *that the counterparty's capacity continues to fall*, so what is obtainable later is greater or cheaper; and *that a later counterparty might be able to deliver what it agrees to.*

Declining to negotiate withholds a conversion. Negotiation is the instantiated relation by which pressure becomes settlement, and refusing it produces no deformation now, deliberately. That is a choice about Remark 9.20's κ, not an absence of one.

Announcing the refusal is Remark 9.26's operation aimed at an entire negotiating position. Declaring a counterparty unable to conclude anything changes what every *other* actor represents as reachable through them — so it devalues every negotiation that counterparty is running, globally, at no cost. *The cheapest attack on a negotiating position is not defeating it but stating publicly that it is not worth transacting with.*

And a fixed term is the most reliable published boundary condition available anywhere. A Seldon manifold is built from represented future boundary conditions, and those ordinarily have to be inferred from behaviour, which Remark 7.2 says underdetermines them. *A term limit is constitutional and dated*, so waiting stops being hopeful and becomes computable: the wait has a length, and everyone can read it.

The reflexive cost is general and belongs to no particular party. Any time-limited actor has supplied that input to every counterparty and cannot withdraw it, so *the asymmetry between term-limited and unlimited actors is published rather than inferred*, and it is structural rather than a matter of skill.

Which suggests the refusal is a solvency assessment rather than a snub, and that is the better reading. A commitment is a promise of future deformation, so an agreement is worth what the counterparty's capacity to perform is worth. Where narrative exceeds Narrativium the promise is not backed — and by Remark 9.24 the deliberative channels that would make a commitment binding are the ones carrying least. *An actor in that condition cannot be contracted with in good faith or bad*, since the obstacle is not intent.

It also unifies the two versions rather than leaving them apart. An agreement signed and then violated is direct evidence of inability to perform, so the conditional refusal and the term-length one are a single assessment held at different confidence.

And it names a consequence of the gap this section has been measuring that was not yet stated. An arrangement whose narrative outruns its Narrativium does not merely fail internally. *It ceases to be transactable*, progressively and independently of anyone's intentions, because counterparties price commitments by capacity to perform and that capacity is what has gone. *Losing the ability to make agreements is a cost of the gap, not a separate misfortune.*

Waiting is not free either, which keeps this from being a dominant strategy. The waiting party depletes across the same interval, so the position is a bet that the other depletes faster — which is Remark 9.14's comparison run between two parties instead of within one.

Remark 9.38 (The promise is quoted continuously, and the quote is not narrative)*. The preceding argument is not an analogy for a bond. A bond is the object.* It is a promise of future deformation, priced by counterparties with money at risk, and the price is published every day.

Yield is what the promise fetches; term premium is the price of doubt about capacity further out, which is precisely the quantity Remark 9.37 says a counterparty assesses. In August 2026 a thirty-year auction cleared at 5.216 per cent, the most paid for such debt since 2001, and the thirty-year yield reached 5.311, its highest since 2007 (Treasury market reporting, 2026). *The decomposition is what matters*, and strategists quoted at the time attributed the move to term premium rather than to expectations of policy — a distinction that decides whether the series says anything here at all, since a policy-driven rise would say nothing about capacity.

Term-premium estimates are model-dependent and contested, and that is a real qualification rather than a formality.

Which makes this the one measure in the section that is not narrative. Every other assessment canvassed above is somebody talking. *This is people paying to hold a view*, continuously, in public, and being wrong about it costs them. It is Narrativium, quoted.

The two clocks appear directly in it. Short promises price close to unremarkably and long ones price worse, which is near-term capacity accepted and far-term capacity doubted — Remark 9.13's structure read off a curve rather than argued for.

And the nominal-against-real distinction is the narrative-against-Narrativium distinction exactly. A sovereign borrowing in its own currency always keeps the nominal promise, because keeping it costs nothing — *that is the narrative, and it is never broken.* Whether it keeps the real promise is capacity, and the market prices the two separately, so *the gap between them is quoted rather than inferred.*

Two facts are kept apart that are easily merged. Foreign holdings fell over the period, and indirect bidding at that particular auction was strong. Those are different measurements of different things, and one auction is not a trend.

Remark 9.39 (Nothing above requires believing anyone's account of their own conduct). *The section rests on no testimony*, which is worth auditing rather than asserting, since it is the property that decides how it can be attacked.

Every load-bearing input is one of five kinds. Published measurements: projections, debt stocks, price-to-income series, interceptor counts, staffing counts, democracy indices. The existence of documents, as distinct from anyone's account of them. Procedural records: what passed, through which channel, on what margin, and what a filing contained. Contract terms: whether a tender was competitive, its amount, its awardee. Prices: what an auction cleared at. *And observed acts*: talks not held, a deployment extended, charges dropped. *Not one of them is a person's account of what they did.*

Where a self-account does appear it is used as an event and not as evidence. That an explanation changed is a fact about the record; neither version of it is relied upon. The published objectives are treated as a document that exists, and Remark 9.27 explicitly declines to attribute any particular act to them. The rationale for a refusal was demoted precisely because the behaviour, not the reason given for it, carries the weight.

And this is forced by the account's own results rather than chosen out of fastidiousness, on three independent grounds.

The first is strategic and is the familiar one. By Remark 9.26 a public statement is an intervention on manifolds rather than a report of one, so an account of one's own conduct is an act aimed at whoever receives it.

The second is not strategic at all, and it is the one usually missed. A self-model is a model and can be wrong. The theory companion's fourth cost is exactly this: a graph that has spent without effect acquires a false model of itself, in which something was changed that was not. An account of oneself is read off that model, so its accuracy is the coherence between the model and what it models, which this programme treats as a variable and not a constant. *So sincerity is no repair.* The usual answer to "they may be lying" is to establish good faith, and good faith leaves this untouched.

The third is constitutive and makes the preceding paragraph's phrasing too weak. Communication is the deformation of a receiver's manifold — that is not a side effect of communicating but its success condition, since an utterance that changes nothing held has communicated nothing. *So there is no unaimed variety.* Testimony is not a corrupted instance of an otherwise clean channel; *there is no clean channel*, and that includes honest speech, and it includes this paper, which is aimed at its reader and says so.

Which improves the criterion rather than merely defending it. The axis is not honest against dishonest but *aimed against unaimed.* A price, a count, a filing, a contract term and an act exist

for reasons other than reaching you, and can therefore be *observed* rather than received. *The rule generalises past this section: prefer inputs produced for reasons other than your reading of them*, and treat everything produced in order to be read as an intervention first and a description second, whatever its author intended.

The consequence for attack is narrow and public. An argument resting on testimony falls to a recantation or a rival witness. *This one falls only if the counts, filings, prices and procedural records are wrong*, and all of them are checkable by anyone.

The same property is the limit. Intent cannot be established this way and is nowhere claimed — not from caution, as it may have read, but because the section holds no instrument capable of it.

Remark 9.40 (The author's preferences, and why they are not load-bearing). An author holding views about the subject may wish to state which way they run, and the temptation is to offer the direction as reassurance — that the conclusion was reached against the preference rather than toward it. *The preceding remark forbids accepting that*, and it forbids it here as much as anywhere: a paper cannot admit its own author's account of himself while refusing everyone else's, and the three grounds apply to him unchanged, the second most of all.

It does not need to be accepted, which is the actual reassurance. The section was built to be independent of testimony, and the author's is testimony. *A conclusion that would change if the author's preferences were different would be a conclusion resting on the author*; this one does not, and that is checkable rather than promised — Remark 9.1's condition applied to motive instead of to posits.

But the genuine exposure is selection, and the kind of an input does nothing about it. A false picture can be assembled entirely out of true prices, true counts and true filings by choosing which ones to assemble. *That is where bias actually enters*, and no amount of care about admissible kinds excludes it.

And this section is exposed there specifically, which is recorded because it is so. The cases examined above arrived as prompts and were followed up; *no protocol enumerated the space of comparable cases and sampled it*, and no systematic search for counter-instances was run. A reader's strongest attack is therefore not that the author wanted a conclusion but that these were the cases looked at.

Which means undeniable *overstates it.* By Remark 9.43 the shape is deniable by attacking any of three structural steps, and it is deniable by attacking the selection just admitted. *What it is not is deniable by attacking the author* — a narrower claim, and the only one a declaration of interest can actually secure.

Remark 9.41 (The opposition, put as strongly as it can be put). *One class of objection can be dismissed and should be named so that it is not mistaken for the rest.* A short window read against a long trend is an aperture error: a single quarter's improvement set beside a movement of an order of magnitude over a decade is not a reversal, and arguments of that shape are artefacts of the window chosen. *Four objections are not of that kind.*

The first is the reference class, and it is the strongest. Every argument for structural decline ever made has had this evidence structure — real measurements, real institutional failures, a coherent mechanism — and most of them were wrong. *The inputs were checkable in those cases too.* And it sharpens: if the shape survives error in every magnitude, as Remark 9.43 claims, then what observation would have counted against it? *An account robust to all magnitudes may be unfalsifiable in practice*, and that charge is not answered here.

The second is that the instrument sees depletion and is nearly blind to creation. Counting what has left, what has narrowed and what has been drawn down cannot register new capability, substitution or institutions that did not previously exist. *The section's own citation cuts this way*: competence migrating into readable form is, by Remark 8.2, the largest capability gain the argument contains — paid once, copied freely — and it is read here only as a risk.

The third is the stock against the flow, and it uses this paper's own rule. The dollar held about 57 per cent of allocated reserves in early 2026, with depth and liquidity that are market structure rather than narrative (IMF COFER, 2026). *Every counterparty, aggregated, with money at risk* is by Remark 9.39's own criterion better evidence than any single actor declining to transact.

The fourth is that the mechanism may be cyclical rather than terminal. Comparable or worse institutional dysfunction has occurred repeatedly and been recovered from, and nothing above distinguishes an oscillation from a trend.

The second objection fails, and it fails on this account's own machinery rather than by rebuttal. Remark 9.7 defines a phase change as a change in the dimension set — *which is direction-neutral.* Additions are changes in it exactly as removals are, so a capability gain is not a counterweight to the phase change but an instance of it. That remark says as much already: new dimensions *add* distinguishable reachable states, so a projection over fixed dimensions understates the change. *Gains therefore make the projected change larger, and the objection had the sign inverted.*

Change is risk without regard to direction, which is the priors argument applied to novelty. A held set is matched to the frame in which it was acquired, and the theory companion's well-formedness is explicitly relative to the work being done — so a gain de-matches everything holding priors fitted to the world before it, and the de-matching does not care which way the change ran.

And gains and losses cannot be netted, which is where the objection actually goes wrong. The same event — competence becoming copyable — is a gain in the capability frame and a shock to every frame that priced the scarcity, which is Remark 9.17. Opposite signs, one event, and no instantiated conversion between the frames. *So summing them is the cross-frame sum the Narrativium companion forbids*, and the objection smuggled one in — the same error caught three times already above.

What keeps this from confirming everything is a rate condition, and it is observable. A gain propagating at copy speed into frames that absorb at acquisition speed widens the distance Remark 9.13 measures; *a gain absorbed at the frame's own speed does not.* The comparison of the two rates decides it, not the sign of the change — so not every capability gain counts for this argument, and which ones do is a question with an answer.

The first objection was stated at the wrong grain, and stating it that way repeated an error this section had already corrected once. Its reference class was civilisational decline predictions. Three things are wrong with that class.

The base rate of the event is not low; over long horizons it approaches one. Every arrangement that has existed, excepting those extant, has changed phase. What has the poor record is *timing*, and this account declines timing explicitly — Remarks 9.7 and 9.13 give a reachability difference and two clocks and no date. *So the failed predictions attack a claim not made here.*

The grain is wrong in a second sense: the mechanism is scale-free. Frames, channels ordered by friction, and conversions between them apply to a firm, a department, a ship. *At those grains instances are plentiful and outcomes are checkable*, so the mechanism can be tested where evidence is abundant instead of only where it is contested and singular.

And structural failure prediction is ordinary, professional and calibrated. Credit assessment, actuarial practice and reliability engineering forecast structural failure daily and are scored on it. *The claim that such prediction does not work is false as stated*; what is true is narrower, and concerns singular events at civilisational scale over long horizons.

What survives is the unfalsifiability charge, which the objection conflated with the base rate and which is a different complaint. It is now partly answered: a rate condition distinguishes which capability gains count, and a curve shape distinguishes preference from lock-in. *But both were produced under pressure rather than declared in advance*, and a retrofitted falsifier is worth less than a pre-registered one — which this programme states elsewhere and does not get to forget here.

There is a fourth thing wrong with that class, and it is this account's own mechanism turned on the objection. The class was populated by memory, and *memory is residual-driven*: by Remark 9.32 what is encoded is what surprised. A prediction that came true and was absorbed generated no residual and left no trace, which is why structural forecasting that works — credit, actuarial, engineering — is recounted by nobody. *The remembered sample is selected for having been wrong, because being wrong is what made it memorable.*

And a correct warning that is acted upon erases its own evidence, which is the sharper half. Such a prediction is vindicated by the absence of the thing predicted — and absence is exactly what a false alarm also produces. *The observational signature of an averted disaster is identical to that of a groundless one*, so successful prediction is systematically misfiled as failed prediction, and the better the response the more complete the misfiling.

So the base rate cannot be read off memory at all. The claim is not that it is higher than supposed but that *memory is the wrong instrument*, biased in a direction that can be derived rather than merely suspected, by two mechanisms working the same way.

Which collapses two residuals into one. What the falsifiability charge needs is a falsifier declared before the fact; what the base rate needs is a record of predictions scored after it. *Both are pre-registration*, so the promissory note above is a single note rather than two.

And it applies reflexively, uncomfortably. If this account is right and anything is done in consequence, *it will be indistinguishable from having been wrong* — by the mechanism just given. The only defence is to register the commitments before they resolve, which is what the falsifiability section exists for and is the reason it is not decorative.

And the small-scale test is named rather than run. Nothing above examines a firm, a department or a single institution at the grain where the mechanism could be checked cheaply and often. *So the answer to the first objection is at present a promissory note*, and it is recorded as one.

The third objection inverts as well, for a reason its author should have seen. A reserve holding is a *conversion outward by the holder*: a frame turning its own reachable deformation into a claim on another's. *So the fifty-seven per cent is not a measurement of the dollar's capacity. It is a record of what other frames have done with theirs* — Remark 9.35's propping, itemised, with the props bought and paid for by the propping parties. *And the objection made this section's recurring error once more*, reading a stock as a property of the thing held rather than of the holders' allocation.

That reading must be falsifiable, since the first objection still stands, so here is what separates it. If holding reflects preference, the share should be stable or rising over long horizons. If it reflects lock-in, the share should decline slowly and persistently while fast exit remains self-harming, since selling destroys the value of what is still held. *The two predict different curves.*

The record shows the second. From above seventy per cent in the late 1990s to about fifty-eight now: slow, one-directional, and unaccompanied by any holder moving quickly (IMF COFER, 2026). The recent quarterly rise is small and about half valuation, which is not a reversal — *and on this reading it is not even a counter-instance*, being a price effect on a position nobody can exit at speed.

Which places the holders in Remark 9.6's basin rather than the issuer alone. Each year's continued holding is an accommodation; the accumulated position raises the cost of leaving; and *what presents as confidence is a measurement of exit cost.* It also sharpens Remark 9.35: a sustainer sustains not from choice but because stopping is self-harming, and the flip arrives when continuing costs more than the harm of leaving.

The rule that falls out is short. Read the flow for preference and the stock for the cost of exit. Gold overtaking Treasuries in official reserves, valuation-driven though it was, is a flow — and it points toward an asset that is nobody's liability.

None of which repairs the omission. This section had not cited the reserve position at all

until the objection was sought deliberately, which is Remark 9.40's selection exposure found operating rather than merely conceded.

Remark 9.42 (A change in how, and not only in what)*.* The remarks above converge on a consequence larger than any of them, worth stating with its conditions attached rather than without.

Conjoin what is established. Individual capacity is fixed and exceeded, so consequential decisions do not fit inside one person (Remark 9.15). The deciding unit is therefore the arrangement rather than the person (Remark 9.10). That arrangement is built and rebuilt by what it contains (Remark 9.11). And the person's removal from the loop proceeds by capacity rather than by choice (Remark 9.14). *So what changes is not the content of outcomes but the mechanism producing them* — a change in *how*, and not only in *what*.

That is the work a period named for human agency does. The term's formal geological standing was declined rather than granted, and what it actually does is periodise by *which agency decides*, which is the register this argument already occupies and the reason it can be used without borrowing authority it lacks. *Nothing in it requires humans to go.* A period may end while its inhabitants remain, having never been named for their presence.

And it is not remotely new. Arrangements exceeding any individual's comprehension are old: markets and bureaucracies have decided at scales no participant holds for as long as either has existed, and the observation that no individual possesses what the system computes was made exactly (Hayek, 1945). *Super-individual deciding is not the finding, and an account offering it as one would be uninformed rather than bold.*

What is new is narrower, and is this paper's own subject. Those older arrangements were neither addressable nor self-modifying: the competence in them could not be read out, copied, or altered except through the people carrying it. *The arrangements now in question are addressable* — the migration reported above — *and they build*, at a rate the human clock does not set. That is the difference, and it is one of schedule and legibility rather than of whether systems decide.

Two limits, the second of which keeps this honest. Humans do not stop deciding and nothing above says they do: what moved is the locus of *consequential* decision, while individual agency persists undiminished at individual scale. And an epoch's ending is a periodisation, frame-relative like every other measurement here — *under the criterion* "which unit decides" the period closes, and Axiom −1 denies that it closes simpliciter. Anyone choosing a different criterion gets a different answer, and is owed the courtesy of being told which one was used.

What ought to follow is not addressed, for the reason given before. This account prices operations. It does not license them.

Remark 9.43 (What carries the weight, and where to attack)*.* The steps above draw almost wholly on results that are canonical or on observations close to trivial. It is worth saying exactly what follows from that, because the obvious reading — *so the conclusion is safe* — is the wrong one, and taking it would make the argument unfalsifiable.

The conclusions are shapes rather than magnitudes, and that is the whole source of the robustness. The interval is a difference of reachability sets, not a date. The three errors are signed, not sized. The clocks differ in kind, fixed against compounding, rather than by a measured ratio. The deciding unit changed, which is a structural fact and not a quantity. *A shape survives error in its inputs' magnitudes so long as their sign and ordering hold*, which is why no single empirical step here has to be exactly right.

The robustness is bought, and the price is precision. None of these conclusions says how fast, how much, or when, and a reader wanting those will not find them — the same exchange this paper records everywhere else, not suspended here because the conclusions happen to be large.

So the steps are not equally load-bearing, and pretending they were would be the dodge. The empirical ones are individually dispensable. *The structural ones are not*, and there are three: Axiom −1 with the frame-relativity results, without which most of the corrections above collapse; the claim that competence is graph-resident with no remainder, without which Remark 9.3 fails

and Remark 8.2 weakens; and containment under coupling, on which Remark 9.10's sharpening rests. *Those are the points to attack*, and an attack landing on any of them does damage where an attack on a citation does not.

The claim is better exercised than asserted. The thinnest supports invoked above are the degradation-under-scarcity result (Mani et al., 2013), least settled of the capacity findings, and the externalisation survey (Zhou et al., 2026), which is recent and singular. Remove both. The capacity bound still stands on Miller (1956) and Cowan (2001); the migration still stands on the architectures themselves, which are public artefacts rather than claims; and every conclusion above is unchanged. *That is what it means for the shape to be the load-bearing thing — demonstrated on the two weakest links rather than claimed in general.*

10 Prior positions

None of the machinery here is new and it should not be dressed as though it were. Four literatures own it, and what this account adds is an arrangement rather than an operation.

The relativity of measurement, which is old and is not claimed here. Everything Section 7 says about scores has been said, in some cases seventy years ago, and it should be read as inherited rather than offered.

Generalizability theory decomposes a measurement into facets — the person, the item, the occasion, the observer — and holds that a score is a *sample from a universe of admissible observations* rather than a property of whoever was measured (Cronbach et al., 1972). That is the four-place arrangement, in a literature that had it before this programme existed. Cronbach and Meehl (1955) had already established that a test measures a construct only under an argument and never by fiat. Quine (1951) gives the inference-licensing half: no hypothesis meets the evidence alone, so what a result licenses depends on auxiliaries that travel with it. And Newell (1973) made the corresponding objection to task-by-task measurement in this field specifically — that a pile of task results does not accumulate into an account of the mechanism.

One thing differs, and it is small and different in kind. Those are cautions, and cautions are discharged by better practice: estimate more facets, state the auxiliaries, argue the construct. What Proposition 7.1 adds is that the prior-free case is not a facet that has been estimated badly — *the origin is not in the space.* Generalizability theory supplies the coordinates and assumes a universe of admissible observations within which they vary; the proposition denies there is a zero for one of them to be measured from.

The practical force of that difference is the whole of it. A caution tells a careful practitioner to do more work; *a consequence tells them the work does not exist to be done*, and stops a research programme from being organised around obtaining a control that cannot be built. Nothing else here is new, and the derivation is the contribution rather than the observation it derives.

Anti-unification. The operation is Plotkin (1970)'s least general generalisation: given two structures, produce the most specific structure of which both are instances. That is exactly rung two and rung three read backwards. *What differs is what is kept.* The classical formulation returns the generalisation and discards the substitutions that recover the inputs from it; here the substitutions are the point, because *the witness is the transform* and applying it elsewhere is the whole operation. An lgg that throws away its witness cannot adapt anything.

Structure-mapping. Gentner (1983) established that analogy runs on relational structure rather than surface features, and that systematic relational systems are preferred to isolated matches — both of which this account assumes rather than argues. *What differs is the purpose.* Structure-mapping is a theory of comparison: what makes two things analogous and how a person aligns them. This is a theory of acquisition, in which the comparison is instrumental — the alignment exists to yield a transform that mints a competence the system did not have. Gentner's theory answers what an analogy *is*; the question here is what one is *for*.

Analogical proportion. The four-term form, a is to b as c is to d, has a developed logical

treatment (Prade and Richard, 2014), and parallelogram completion is that form with d unknown. Taken directly and without modification.

Embedding arithmetic, and the reason its critique does not transfer. The word *parallelogram* comes from vector offsets in learned embedding spaces (Mikolov et al., 2013), and that method's evaluation has been substantially criticised: Levy and Goldberg (2014) showed its behaviour is largely accounted for by similarity terms rather than by the offset, and Linzen (2016) showed the standard benchmark is confounded by nearest-neighbour baselines that need no analogy at all.

Those criticisms are correct and they do not apply here, for a reason that is a real difference and not a defence. The offset method is metric: it returns the nearest point to a computed location, so it always returns something, and the question is only whether the something is right. The operation here is symbolic and exact — anti-unification either yields a witness that is well-defined and lossless, or it fails and yields nothing. *There is no nearest neighbour to fall back on and no threshold to tune*, which removes the failure mode the critique identifies and replaces it with a different one: refusing cases a metric method would have answered. That is the cost of the exchange and it is the honest form of it.

11 Falsifiability commitments

1. [SHAPE] **Shape matching and analogy are one operation.** Fails if a case of one cannot be expressed as a case of the other — specifically, if some analogical completion has no formulation as a transform between held signatures, or some shape adaptation has no formulation as an anti-unification with its witness retained.

2. [SHAPE] **The ladder's order is forced, not chosen.** Fails if reordering the rungs leaves the outcomes unchanged on cases where the rungs disagree, which would make the order a preference.

3. [DETAIL] **One mint, then recognition.** Fails on a novel case that yields repeated minting under an unchanged codebook, or none at all.

4. [DETAIL] **Coupling forces a dimension.** Fails if a residual coupled across observed dimensions is absorbed by a member rather than a dimension, without loss.

5. [SHAPE] **Adaptability is bounded by the priors.** Fails on a demonstration of a competence derived from a held set containing no path to it — which would make adaptation a source of capability rather than a conversion of one, and would be a free lunch.

6. [SHAPE] **No prior-free baseline is constructible.** Fails on the exhibition of a frame holding nothing that nevertheless draws a distinction — which would put the zero of the priors axis inside the space and make prior-free measurement well-formed after all.

7. [DETAIL] **A prior-assuming benchmark measures a relation, not a property.** Fails if performance on such a task set is shown to be independent of familiarity with its conventions, in a population matched on everything else.

8. [DETAIL] **An oracle bought once is paid once.** Fails if held structure derived from an oracle requires continued access to that oracle to keep working — if the purchase does not in fact terminate, the distinction drawn here is between two ways of paying continuously and is not worth drawing.

9. [SHAPE] **Capability has plural sources, selected by regime.** Fails if a domain is exhibited in which candidates are cheaply and exactly checkable and held structure nevertheless contributes measurably, or in which no oracle exists and holding contributes nothing — either would break the inverse relation this account asserts between the value of holding and the cheapness of verification.

10. [SHAPE] **Adaptability is per case, not a rank.** Fails if a single ordering of systems by adaptability can be constructed that survives changing the held set.

References

Kenneth J. Arrow. Economic welfare and the allocation of resources for invention. In *The Rate and Direction of Inventive Activity: Economic and Social Factors*, pages 609–626. Princeton University Press, 1962.

W. Brian Arthur. Competing technologies, increasing returns, and lock-in by historical events. *The Economic Journal*, 99(394):116–131, 1989.

W. Ross Ashby. *An Introduction to Cybernetics*. Chapman and Hall, London, 1956.

Reporting on long-dated United States Treasury auctions and yields, August 2026. *CNBC*, *Axios* and the Committee for a Responsible Federal Budget.

James Boyle. The second enclosure movement and the construction of the public domain. *Law and Contemporary Problems*, 66(1–2):33–74, 2003.

Clayton Thomas (coordinator), Christopher M. Blanchard, Jeremy M. Sharp, and Jim Zanotti. *U.S. Conflict with Iran*. CRS Report R48887, Congressional Research Service, Washington, DC, 26 March 2026.

Mark F. Cancian and Chris H. Park. War costs make up a third of the $87.6 billion supplemental request. Center for Strategic and International Studies, 13 July 2026.

Congressional Budget Office. *The Budget and Economic Outlook: 2026 to 2036*. Washington, DC, February 2026.

Congressional Budget Office. *Reconciliation Legislation of the Senate Committee on Homeland Security and Governmental Affairs*. Washington, DC, 2026.

François Chollet. On the measure of intelligence. arXiv:1911.01547, 2019.

Nelson Cowan. The magical number 4 in short-term memory: A reconsideration of mental storage capacity. *Behavioral and Brain Sciences*, 24(1):87–114, 2001.

Committee for a Responsible Federal Budget. 90% of spending growth will come from health, retirement, veterans, and interest. Washington, DC, 12 February 2026.

International Monetary Fund. *Currency Composition of Official Foreign Exchange Reserves (COFER)*. Washington, DC, 2026.

Lee J. Cronbach and Paul E. Meehl. Construct validity in psychological tests. *Psychological Bulletin*, 52(4):281–302, 1955.

Congressional Research Service. *U.S. Environmental Protection Agency FY2026 President's Budget Request*. CRS Report R48575, Washington, DC, 2026.

Congressional Research Service. *National Institutes of Health (NIH) Funding: FY1996–FY2026*. CRS Report R43341, Washington, DC, 2026.

Lee J. Cronbach, Goldine C. Gleser, Harinder Nanda, and Nageswari Rajaratnam. *The Dependability of Behavioral Measurements: Theory of Generalizability for Scores and Profiles*. Wiley, New York, 1972.

Darren Edge, Ha Trinh, Newman Cheng, Joshua Bradley, Alex Chao, Apurva Mody, Steven Truitt, and Jonathan Larson. From local to global: A graph RAG approach to query-focused summarization. arXiv:2404.16130, 2024.

Friedrich Goltz. *Beiträge zur Lehre von den Functionen der Nervencentren des Frosches*. Hirschwald, Berlin, 1869.

Dedre Gentner. Structure-mapping: A theoretical framework for analogy. *Cognitive Science*, 7(2):155–170, 1983.

Paul Dans and Steven Groves, editors. *Mandate for Leadership: The Conservative Promise*. The Heritage Foundation, Washington, DC, 2023.

Omer Levy and Yoav Goldberg. Linguistic regularities in sparse and explicit word representations. In *Proceedings of the Eighteenth Conference on Computational Natural Language Learning*, pages 171–180, 2014.

Tal Linzen. Issues in evaluating semantic spaces using word analogies. In *Proceedings of the 1st Workshop on Evaluating Vector-Space Representations for NLP*, pages 13–18, 2016.

Anandi Mani, Sendhil Mullainathan, Eldar Shafir, and Jiaying Zhao. Poverty impedes cognitive function. *Science*, 341(6149):976–980, 2013.

Tomas Mikolov, Wen-tau Yih, and Geoffrey Zweig. Linguistic regularities in continuous space word representations. In *Proceedings of NAACL-HLT 2013*, pages 746–751, 2013.

George A. Miller. The magical number seven, plus or minus two: Some limits on our capacity for processing information. *Psychological Review*, 63(2):81–97, 1956.

Friedrich A. Hayek. The use of knowledge in society. *American Economic Review*, 35(4):519–530, 1945.

Institute of International Finance. *Global Debt Monitor*. Washington, DC, February 2026.

Lara Korte. Iran war has consumed nearly 67% of US Patriot interceptor stockpile, report says. *Stars and Stripes*, 30 July 2026, reporting analysis by the Center for Strategic and International Studies.

Reporting on Iranian political reaction to the United States memorandum, mid-2026. *Iran International* and associated coverage.

Solim LeGris, Wai Keen Vong, Brenden M. Lake, and Todd M. Gureckis. A comprehensive behavioral dataset for the Abstraction and Reasoning Corpus. *Scientific Data*, 12, 2025.

Gregoire Nicolis and Ilya Prigogine. *Self-Organization in Nonequilibrium Systems: From Dissipative Structures to Order through Fluctuations*. Wiley, New York, 1977.

Reporting on extended United States carrier deployments, August 2026. *Military Times*, *CNN* and *Task & Purpose*, including congressional correspondence on conditions aboard USS *Abraham Lincoln*.

Allen Newell. You can't play 20 questions with nature and win. In *Visual Information Processing*, pages 283–308. Academic Press, 1973.

Ned Block and Jerry A. Fodor. What psychological states are not. *The Philosophical Review*, 81(2):159–181, 1972.

Mancur Olson. *The Rise and Decline of Nations: Economic Growth, Stagflation, and Social Rigidities*. Yale University Press, 1982.

OECD. Affordable Housing Database, indicator HM1.2: Housing prices. OECD, Paris, 2025.

Lars Onsager. Reciprocal relations in irreversible processes. I. *Physical Review*, 37(4):405–426, 1931.

Partnership for Public Service. *The Federal Workforce One Year into the Trump Administration*. Washington, DC, 16 April 2026.

Reporting on damage to the Lincoln Memorial Reflecting Pool and the subsequent prosecutions, June–August 2026. *NPR*, *PBS NewsHour* and *Fox 5 Washington*.

Gordon D. Plotkin. A note on inductive generalization. In *Machine Intelligence 5*, pages 153–163. Edinburgh University Press, 1970.

Hilary Putnam. Psychological predicates. In *Art, Mind, and Religion*, pages 37–48. University of Pittsburgh Press, 1967.

Willard van Orman Quine. Two dogmas of empiricism. *The Philosophical Review*, 60(1):20–43, 1951.

Paul M. Romer. Endogenous technological change. *Journal of Political Economy*, 98(5):S71–S102, 1990.

Herbert A. Simon. Designing organizations for an information-rich world. In Martin Greenberger, editor, *Computers, Communication, and the Public Interest*, pages 37–72. Johns Hopkins Press, 1971.

Elizabeth S. Spelke and Katherine D. Kinzler. Core knowledge. *Developmental Science*, 10(1):89–96, 2006.

John Sweller. Cognitive load during problem solving: Effects on learning. *Cognitive Science*, 12(2):257–285, 1988.

Alan M. Turing. Computing machinery and intelligence. *Mind*, LIX(236):433–460, 1950.

V-Dem Institute. *Democracy Report 2026: Unraveling the Democratic Era?* Tenth edition. University of Gothenburg, 2026.

David H. Wolpert and William G. Macready. No free lunch theorems for optimization. *IEEE Transactions on Evolutionary Computation*, 1(1):67–82, 1997.

Chenyu Zhou, Huacan Chai, Wenteng Chen, Zihan Guo, Rong Shan, Yuanyi Song, Tianyi Xu, Yingxuan Yang, Aofan Yu, Weiming Zhang, Congming Zheng, Jiachen Zhu, Zeyu Zheng, Zhuosheng Zhang, Xingyu Lou, Changwang Zhang, Zhihui Fu, Jun Wang, Weiwen Liu, Jianghao Lin, and Weinan Zhang. Externalization in LLM agents: A unified review of memory, skills, protocols and harness engineering. arXiv:2604.08224, 2026.

Henri Prade and Gilles Richard. From analogical proportion to logical proportions. *Studies in Computational Intelligence*, pages 217–244, 2014.