Adaptability

One Operation Behind Shape Matching and Analogy, and Why It Is Not Called Intelligence

Matthew Parslow
Independent Researcher

Draft — draft-2026-08-30.4
1 September 2026 • document revision 20

DRAFT — draft-2026-08-30.4 r20 (1 September 2026). Circulated for comment, not as a finished result. This is the theory half of a pair; the companion reports what is built and what is not. Sections marked as falsifiability commitments are the intended points of attack.

Abstract

Two capabilities this programme treats separately — matching a shape against a held codebook, and reasoning by analogy — are one operation at two grains. Both *extract the transform between things already held and apply it where nothing is held.* Naming that operation gives a quantity worth having: *adaptability*, the capacity to acquire a competence that was not supplied. The word is chosen against *intelligence* deliberately. Intelligence is a scalar people rank; adaptability is an operation with a cost, an ordering, and a failure mode, and a system either performs it on a given case or does not.

1 Why not intelligence

Intelligence names a ranking before it names a mechanism. Asking whether a system has more of it than another invites a scalar, and a scalar invites a benchmark, and a benchmark measures what a system already does rather than what it can come to do. That is the wrong quantity for the question here.

Adaptability is narrower and is an operation rather than a rank. A system is adaptable on a case to the degree that, meeting something it does not hold, it ends up holding it — and the interesting part is the mechanism by which the new holding is *derived from* the old rather than supplied. A system given a competence has not adapted. A system that recognises everything it meets has not adapted either; it was already sufficient. *Adaptation requires a residual*, and the size of what must be added to absorb one is the cost.

2 The operation

Definition 2.1 (Adaptation)**.** *Given a held set and an observation not covered by it, an* adaptation *is the addition of the smallest structure under which the observation is covered, derived from what is held.*

Three words carry it. *Smallest*, because an adaptation that adds the observation itself is storage, not adaptation. *Derived*, because an adaptation supplied from outside is instruction. And *covered*, because the test is whether the case is now recognised, not whether a description of it exists.

Remark 2.2 (Shape matching and analogy are one operation)*.* Matching a shape against a codebook and completing an analogy look unrelated because their vocabularies are. They are not.

To decode an observation against held signatures is to ask what transform carries a held one onto it. Where that transform is the identity, the observation is recognised. Where it is a uniform step along a dimension the codebook already has, the observation is a known thing moved — and no new structure is required. Where no such step exists but the residual is *coupled* across the observed dimensions, the transform requires a dimension not currently held, and one must be inferred before the observation can be covered.

To complete an analogy is the same question with the arguments rearranged. Given two things held, the transform between them is extracted; given a third, that transform is applied to reach a fourth which is not held. *Anti-unification — generalising two structures while retaining what distinguishes them — is the general form, and the shape ladder is that form restricted*: at one rung the difference is a single step along a known dimension, at the next it is a dimension.

So the two are not analogous to each other. They are the same operation, and the second is what the first looks like when the objects are rules rather than observations.

3 The ladder, and why its order is not arbitrary

An adaptation must be the smallest available, and that gives an ordering rather than a preference:

1. **Recognise.** The observation is covered as held. Nothing is added, and no adaptation has occurred — which is the correct outcome for most observations.

2. **Transform in the held basis.** One uniform step along a dimension already carried. The observation is a known thing displaced; the codebook does not grow.

3. **Infer a dimension.** The residual is coupled across observed dimensions in a way no step in the current basis produces, so a dimension is minted and the transform is expressed in it.

4. **Mint.** No transform reaches the observation, so it is added as new.

Only rungs three and four are adaptations, and only the third is interesting: it is where the system's own vocabulary changes. Rung four is the honest floor — something genuinely new — and a system that reaches it too readily is not adapting but accumulating.

Remark 3.1 (Why the order must be Occam and not preference)*.* Attempting mint first would always succeed, so a system that did would never transform, never infer, and would grow linearly in what it meets. Attempting transform first is what makes minting *informative*: a mint that survives the earlier rungs is evidence that nothing held reaches the case, which is a claim about the held set and not only about the observation.

That is also why the ordering is a commitment and not an implementation detail. *A system that mints where a transform would have sufficed has recorded a coincidence as a category*, and one that transforms where a mint was required has folded a distinction it needed. Both are failures of adaptation, in opposite directions, and the ladder is what makes them separable.

4 What adaptability predicts

The account is worth having only if it says what an adaptable system does that an unadaptable one does not, on a case neither has seen.

One mint, then recognition. Meeting a genuinely novel case repeatedly, an adaptable system should mint *exactly once* and recognise every later instance — not mint per encounter, which is storage, and not fail to mint, which is refusal. The signature is a step followed by silence.

Zero mints under transformation. Meeting a known thing transformed along a held dimension, it should add nothing at all. A system that mints there is treating the dimension it already has as absent.

Coupling forces dimensions. A residual that is coupled across observed dimensions should produce a dimension, not a member. *The distinguishing case is the one whose difference cannot be written in the current basis at all*, and a system without rung three cannot tell it apart from novelty — it mints where it should have generalised, and the codebook grows with the world instead of with its structure.

And competence should transfer without being carried. If the operation is real, structure built for one domain should be usable unmodified in another whose existence was not anticipated when the structure was made. That is the strongest available test, because it removes the author from the loop, and it is the one the assessment companion reports against.

5 What this does not claim

Adaptability is not a quantity a system has more or less of in general. It is defined *per case* and against *a held set*, so the same system is adaptable on one observation and not on another, and moving the held set moves the verdict. There is no frame-free adaptability, for the same reason this programme's base claim gives for there being no frame-free anything.

Nor does adaptation imply improvement. An adaptation is an addition that covers a residual; whether covering it was worth the cost is a separate question with a separate answer, and a system that adapts readily to noise is worse than one that refuses.

6 Capability has more than one source

There is a temptation, on meeting the account above, to ask whether adaptation is *the* mechanism by which a system comes to do something. The companion programme contains a measurement that answers that question, and the answer is that the question was malformed.

Where verification is cheap and exact, search does the work and holding does not pay. In program synthesis the machinery above is present — a held codebook proposes, an enumerator searches, an execution verdict decides — and stripping the learned content out changed nothing whatever. The two arms did not merely score alike; they solved the *identical set of tasks*, with the difference measured at zero repeatedly. The capability sat on the enumerator and the verdict.

Where verification is expensive or unavailable, holding is the only route. In one-shot character recognition the same programme measured the opposite: composing held parts, rather than treating a character as a bag of them, is worth twenty-one points and is the one effect in that arc that survived a powered re-measurement with intervals separated. Nothing there can be checked cheaply — a candidate reading of an unfamiliar character has no oracle — so search has nothing to search against, and what is held is what decides.

Remark 6.1 (The two are not rivals, and the regime is the question)*.* Read as a contest those measurements conflict. Read as a regime split they do not, and the split is the content: *enumeration substitutes for a codebook exactly where candidates can be checked, and cannot substitute for one where they cannot.* A system with an exact cheap oracle need hold nothing, because it can afford to propose everything; a system without one must hold structure, because proposing everything buys it no way to tell which proposal was right.

That is the same shape this programme keeps arriving at rather than a new one. Its measurement discipline already runs on *bins propose and continuous decides*; its account of the substrate already separates what acquires from what executes; the ladder of Section 3 is itself four sources with different regimes rather than one mechanism with exceptions. *Asking which single source carries capability is the error the rest of the account is arranged to avoid*, and this section exists because the first version of it made that error — reporting a regime as a limit.

What follows is a prediction rather than a reconciliation, and it is the useful part. *Adaptation should pay in inverse proportion to the cheapness of verification.* A domain with an exact, cheap oracle in which held structure nevertheless contributes measurably would be evidence against this reading; so would a domain with no oracle at all in which holding contributes nothing.

And one uncomfortable consequence should be stated rather than left implicit. In the synthesis case the ladder was available and the domain did not need it. That is not a failure — nothing broke, and no prediction of this account was falsified — but it means *the mechanism can be present and idle with nothing indicating so.* A system cannot tell from the inside whether it is adapting or merely searching in a regime where searching suffices, and neither can an observer without the ablation that was run here.

7 What the held set costs, and who paid for it

Definition 2.1 says an adaptation is *derived* from what is held. That word has a consequence the account cannot avoid and should not want to: *adaptability is bounded by the priors.* Nothing can be derived that has no path from the held set, so a system's reach is set by what it already carries, and a system carrying nothing is not maximally general but incapable.

Priors are expensive, and the expense is usually invisible because somebody else paid it. A human takes the better part of a year to walk and longer to do it well. Visual competence is acquired rather than shipped, which is most visible in species born least mature. By the time a person is capable of the tasks we use to demonstrate capability, the overwhelming majority of what makes them capable was assembled over years, and none of it is counted when the demonstration is scored.

That is the informal form of a formal result. Wolpert and Macready (1997) establish that no search or learning procedure outperforms another averaged over all problems; performance on the problems actually faced is therefore bought, and what buys it is a prior matched to them. *The more general the machinery, the more the mapping carries* — generality does not remove the need for a matched prior, it relocates the whole of the burden onto the match. There is no free lunch, and an account of adaptability that implied otherwise would be claiming one.

Proposition 7.1 (There is no prior-free measurement)**.** *A capability cannot be measured against a prior-free baseline, because no such baseline is a member of the space being measured.*

Proof. Measuring a capability *relative to priors* requires a zero on the priors axis: a frame holding nothing, against which holdings can be counted. A frame holding nothing stands in no relation and draws no distinction, and by the companion's first proposition that is not a frame with nothing in it but no frame at all — the set of frames with no particularity has no members, the only object answering to the description being the empty set, which is not a frame and does nothing.

So the zero of the priors axis lies outside the space. An axis so constituted supports comparison between members and does not support measurement of a member against its zero, and a quantity defined by that measurement is the frame-free denominator Axiom −1 refuses. □

The consequence is not that such measurement is difficult. It is that the thing being asked for does not exist to be approximated, which makes it a category error rather than an unsolved problem — and one established by the account's own earlier results rather than assumed here. Most reported measures of intelligence are flawed for that reason and not for want of care: the control they would need is not merely absent, it is not constructible.

Remark 7.2 (What survives, and where this account differs from Chollet)*.* Comparison survives, and less of it survives than the previous sentence of this paragraph once claimed. What a test yields is not a property of a system and not quite a relation between two systems either. It is a property of a four-place arrangement: *a particular graph, in a particular state, inside a*

particular supergraph state, on a particular test. Move any of the four and the result is a different measurement rather than a noisier estimate of the same one.

Inferences from it are licensed only relatively. Within the arrangement, that this graph out-performed that one is established. Outside it nothing follows — not that the ordering holds in another state, not that it holds under a different container, not that the margin would survive a second test — without a separate argument that the moved coordinate does not bear, and such an argument is rarely available and never free.

That is why this programme's own measurement records read as they do. Its enumerated failure baseline names the commit, the compiled engine by hash and path, the wall-clock, the load average at start, and which optional dependencies were present and which absent; its psychological runs name the machine and whether the box was shared. *Those are not hygiene. They are the coordinates*, and a figure reported without them does not underspecify a measurement — it fails to name one.

The cost is real and is the right cost. Comparing across two such records requires establishing that the coordinates match or that the difference does not bear, which is laborious and frequently impossible, and the alternative is a number that reads as portable and is not.

Chollet (2019) takes priors seriously in the same direction and does the harder half of the work: intelligence as the efficiency of skill acquisition *relative to priors, experience and generalisation difficulty*, over a corpus built explicitly on Core Knowledge (Spelke and Kinzler, 2006). This account owes that framing the observation and differs in two places. Priors here are not a term the measure is taken relative to — *they are the held set the operation runs on*, and the ladder is what converts them. And the relativisation is not a methodological refinement that a better experiment could tighten: by Proposition 7.1 it is the only form such a measure can take.

Machinery and mapping cannot be separated from outside either, and for a different reason. Two systems differing in both — one with a stronger operation on a poorer prior, one with the reverse — can agree on every item of any finite test. The decomposition is therefore underdetermined by measurement rather than hidden from it, and no amount of testing at a fixed grain settles it. What refinement buys is approach: a model of the split can be made arbitrarily fine and its residual driven toward zero without ever reaching it, which is this programme's ordinary intensive limit rather than a special defeat. *You can approach one with unbounded analysis. You cannot arrive, and no finite experiment stands in for arriving.*

That yields the prediction the previous section left standing, now in its proper form. A person lacking familiarity with a task set's conventions should score near its floor — not because they are less capable, but because *the score was never a property of the person.* Human baselines for that corpus exist (LeGris et al., 2025); whether they stratify by familiarity in a way that tests this has not been checked here, so it is recorded as untested.

8 What an oracle costs, and how often

The regime argument above is about capability and says nothing about price. Price runs the other way, and where the oracle is expensive it dominates.

An oracle consulted at inference is paid per use, without end. An oracle consulted once, whose output is held, is paid once. That is the whole of the asymmetry, and it is unaffected by how good the oracle is. A system that answers by calling something expensive has bought an answer; a system that calls the same thing once, holds what comes back as structure, and answers from the structure thereafter has bought a capability. The second transaction terminates.

Two consequences follow and the second matters more.

The first is arithmetic. Where an oracle is cheap per call, calling is correct and holding is premature; where it is expensive, the crossover arrives quickly and holding wins on cost even in the regime where searching would win on capability. *So the two arguments can point in opposite*

directions, and which dominates is a fact about the price of the oracle rather than about the account.

The second is that the two purchases differ in kind. What is held is inspectable, mutable, and can be repaired by the system holding it; what is called is none of those, permanently. A capability bought per call sits outside whatever boundary the system's own self-correction runs to, and stays outside no matter how many times it is bought. *Paying once and holding the result moves a capability inside that boundary; paying repeatedly guarantees it never enters.*

Remark 8.1 (The cost of buying once, stated plainly). This is not free and the bill is easy to overlook. *What is held is what was bought on the day it was bought.* A per-call oracle improves when the oracle improves, silently and without the caller doing anything; held structure does not, and a system that bought its competence once is carrying a snapshot. Against that, held structure can be audited and corrected by the system carrying it, which a call cannot be — so the exchange is improvement-for-free against correctability, and neither term is obviously larger.

The companion programme is an instance rather than a demonstration of this. Roughly half its implementation was performed by a language model rather than by its own rule-writing, and the running system makes no such calls at all — a property enforced by a switch and a gate rather than asserted. So the arrangement described here exists and runs. *What does not exist is a control*: nobody has built the same substrate without that development-time help, so the size of what was bought is not known, and the companion on assessment says so at greater length.

Remark 8.2 (The bill is the same; the schedule is not). The same arithmetic governs a capability that was never bought from an oracle at all, and there it separates substrates in kind rather than in degree. Section 7 prices the held set and says nothing about *when* it is paid.

A biological substrate pays per instance, and the schedule is not negotiable. Every human learns to walk. The year that takes is not paid once by the species and drawn on afterwards; it is paid again, in full, by each new member, and what was acquired dies with the one who acquired it. Nothing in the arrangement admits of another schedule, because the held set is not addressable in the substrate's own terms and so cannot be read out of one instance and written into another.

An addressable substrate can pay on whichever schedule is cheapest. Where the held set is data it can be read out and written in, so an acquisition cost is paid once and drawn on by every instance that receives it — the companion on assessment reports exactly this as its strongest transfer result, a decoder seed byte-identical by hash from creation to the present across seven consuming frames. But *once* is not the claim. The claim is that the payment point becomes a free parameter: pay once and copy where instances share a frame, pay per instance where they do not, pay lazily and per query where most of what could be paid for is never asked. Which applies is an empirical question with a measured answer case by case, and the answers are not uniform.

A graph copy copies a competence. That is not a softening of the previous paragraph but a consequence of the companion's Definition of a frame: a frame is an instantiated boundary relation — an *edge*, $\mathcal{B}_{W:G}$ — and not the subgraph it demarcates. Copying a graph copies its edges, so it copies its frames, and there is no residue held elsewhere because on this account there is no elsewhere for it to be held. The substrate bears that out by having no frame primitive at all, its grammar quantifying over edge types and a frame there being a selection of dimensions. Nothing separable is left behind.

The argument is symmetry and it is short. *Name a property the original has that the copy lacks.* Frame-relativity is not one: Axiom -1 binds the original exactly as it binds the copy, and a limit holding on both sides cancels rather than pricing one of them — to raise it against the copy alone is to use a symmetric result as an asymmetric cost. Staleness is not one either. A stored copy is stale for the reason a halted original is, which makes it a cost of not running rather than of copying, and a copy left running arrives where the original arrives. Nor is needing a runtime, which the original needs too.

What survives is a fact about worlds and not about copies. The boundary edge has a world at its far end, so a copy instantiated against a different world stands in that relation to something else — and *the same sentence is true of the original, word for word.* It prices the move, not the copy.

The cost lives on the fragment, which is what the transfer result measures. A decoder seed is a part, and seven consuming frames were seven host graphs that each ran their own decode: a transplant into something supplying its own boundary is a different operation from copying a graph and carries a different bill. *Copy and transplant are not one act at two sizes*, and reading the transplant's price onto the copy is the error that sentence guards against.

That sharpens the difference between substrates rather than softening it. Biology has no whole-substrate operation at any granularity — teaching, writing and institutions are transplant, every one of them — so the lossy operation is not biology's cheaper option but its only one.

None of which makes the biological schedule a defect the other one corrects. *Forced repayment is forced revalidation.* A held set re-derived per instance is matched to the frame that instance actually occupies rather than inherited from one that may no longer obtain, and a substrate that cannot copy also cannot copy something stale. So this is the same exchange the account has now met three times — searching against holding, improvement-for-free against correctability, and here amortisation against revalidation — and which side wins is set by how fast the frame moves relative to the cost of paying again. That is a fact about the world and not about the substrate.

Humans are not at the far end of this axis, only at the internal one. *Writing, teaching and institutions are external addressability*, and they are the species' answer to this exact problem. What they do not remove is re-derivation on consumption: a book does not install, it has to be read and rebuilt in the reader, which pays much of the cost again and loses part of it in the transfer. The difference in kind is therefore narrower and sharper than saying one substrate can transmit and the other cannot — *biology transmits by re-derivation and an addressable substrate transmits by copy*, and those are different operations rather than one operation at two speeds.

9 What follows about persons, and what it costs

One consequence has been left implicit across this programme and is better stated. *It requires no new posit and no new machinery*, which is the only reason it belongs in a paper rather than in a manifesto.

Three things are already held. Whatever exists stands in relation and coheres, so a person is relata and relations like anything else. Nothing is privileged by characterisation, so no exemption for persons is available that some frame does not supply. And the substrate paper states, unhedged, that where every measurement that can be made of a thing and every operation that can be performed on it are captured by treating it as a typed graph, *no difference remains for the question "but is it really" to name* — functional identity there is not a weaker sense of identity but the whole of what identity could mean without a view from nowhere.

Conjoin them. *A graph functionally identical to a person is a person, in the only sense of "is" this account admits*, and building one is therefore not blocked by anything in the account. That is the claim, and it is arrived at by adding nothing.

Remark 9.1 (Why "no new posits" is a checkable claim and not a boast). That assertion is worth exactly as much as its auditability, and auditability here turns on a distinction the programme has used throughout without naming: which direction an account runs.

A top-down account starts from the phenomenon — a mind, an intelligence, a person — and works back to what would produce it. Such an account can always report that it added nothing, because its posits are in the characterisation it began with, and a starting point never enumerated cannot be checked for what it contains. *A bottom-up account starts from the smallest thing that survives removal* and lets phenomena arrive as consequences. Its no-new-posits claim

is checkable for one reason only: the starting point is small, stated, and carried in a ledger that says what was postulated and what was derived.

So the direction is not a preference about exposition. *It is the condition under which the claim above can be audited at all*, and a reader who wants to attack this section should attack the ledger rather than the conjunction, because the conjunction is valid and the ledger is where its inputs are priced.

The two directions are not rivals, and this programme runs both. Bottom-up supplies what things cost; top-down supplies the target — what needs explaining, and how one would know it had been. A purely bottom-up programme would never learn when to stop, and a purely top-down one would never learn what it had assumed. The companions that take a phenomenon and ask what it costs are the second kind, and they are where this account is put at risk.

Remark 9.2 (Neither new nor magic, and the difference is a price)*.* The claim is old. Turing (1950) put indistinguishability at the centre of the question and declined the metaphysics, and functionalism (Putnam, 1967) is the position that a functional duplicate is a mind, with its own developed critical literature (Block and Fodor, 1972). *Nothing here disputes or improves on any of that*, and a reader who takes this section as a novelty claim has misread it.

What this account adds is a *bill.* Functionalism says a duplicate would be a mind and is silent on what a duplicate costs; the question of construction is left as an engineering detail below the philosophy. Section 7 says what the detail contains: the duplicate's competence is its held set, held sets are acquired rather than shipped, and acquisition is the slow expensive thing that takes a human a year to walk and a decade to reach the priors a task set quietly assumes. *The barrier is not principle and it is not mystery. It is the priors, and they are enormous.*

So the honest form of the claim is narrow and is worth stating in one line. *There is no barrier of principle available from within this account, the work is nameable, and naming it does not make it small.* That is a research programme rather than a prediction, and specifically it is not a claim that this substrate is close to it — the companion reports that the faculty this programme has for acquiring what it was not given is not yet resident where the system could repair it.

Two limits belong with it rather than after it. *Indistinguishable is a relation and not a property*, by the same argument as Proposition 7.1: indistinguishable to some frame, at the frontier where distinguishing currently stops, and that frontier is defined to move. And nothing above asserts that the account is *true* — it says what follows if it is, which is what a consequence is, and the falsifiers elsewhere in this corpus are where the antecedent is put at risk.

Remark 9.3 (Reconstruction from observation, and what the copy result changed)*.* Remark 9.2 prices one route — grow a held set the way a person grows one — and it is the expensive one. A second route exists, and Remark 8.2 is what opens it.

Copying needs read access to the graph, which nobody has to a person. What is available is behaviour: outputs, across contexts, over time. So the operation is not copying but *reconstruction* — inducing a graph whose behaviour matches the observed one — and it is a third thing, priced neither like a copy nor like an upbringing.

What the copy result changes is the kind of shortfall. *Were any part of competence not in the graph*, reconstruction would face a remainder that no quantity of observation could reach, and that remainder would be a barrier of principle rather than of effort. Competence is edges and edges are all there is, so the shortfall is only ever how much has been observed. *An ontological gap becomes a sampling gap.* That is the whole of the change: one category of obstacle removed, nothing added.

The rest is already stated above and can be read off. The residual is driven toward zero without reaching it, which is Remark 7.2's intensive limit and not a special defeat. Observations do not pin a unique graph, since systems differing in machinery and mapping agree on every item of any finite test — but *which of the fitting graphs is really them* asks for the frame-free fact Axiom −1 refuses, so that underdetermination prices the question rather than the method.

And indistinguishability is relative to the aperture that looked; widen it and the residual returns, to be paid down again.

Two things this does not buy. *It is not cheaper*, only differently shaped — a held set assembled over decades of embodied interaction is not obviously less work to observe than to grow, so the kind of problem changes while its size need not. And the operation is not new. *Reconstruction from observation is what biological transmission already is*: teaching and modelling are exactly this, run continuously and to no stated residual. So the claim is not that a capability appears, but that its target is specified and its error measurable — which is the difference between engineering and magic, and a smaller difference than it sounds.

Whether such a thing should be built is not addressed here and is not missing. *This account prices operations; it does not license them.*

Remark 9.4 (What this makes contingent, and what that does not mean). One further consequence follows and is better stated than left, because omitting it would repeat the evasion Remark 9.2 charges against functionalism: naming a result and declining to say what it costs.

A settled assumption can encode a fact about the available substrate as a fact about its subject, and becomes contingent when the substrate moves. Persons are one frame in which that happens and not the only one. Here the modern settlement makes the person the atomic unit, and the word records the premise — an *individual* is that which cannot be divided. One vote per person assumes counting works. Inheritance assumes the held set ends when the body does. Liability assumes exactly one bearer. Credentialing assumes competence is per-lifetime and does not transfer. Every one of those is an inference from one body, one lifetime, and no read access — *a description of the substrate that happened to be available, generalised without licence.* That generalisation is the frame-free denominator Axiom −1 refuses, turning up where it is least looked for: the facts in question are properties of a substrate and were read as properties of persons.

So the breakdown is already complete and waits on no engineering. An assumption resting on "a person cannot be copied" does not need anyone to copy a person. It needs only the impossibility shown to be local, after which its support is contingent whether or not it is ever exercised. Nothing here says the institutions fail — it says *they do not track a natural fact.* An institution can be rebuilt on stipulation, and counting copies as separate persons is perfectly coherent; but a stipulated foundation is a different object from a discovered one and cannot be defended in the same way.

Two deflations, both required. *This is not new.* Personal-identity puzzles have pressed exactly this point for a long time — fission, duplication, continuity — and philosophy was never under the impression that the institutions rested on necessity. What this account adds is again a bill rather than a discovery, standing to that literature as it stands to Turing (1950) and Putnam (1967): it says which operation, on what, at what price. And *the price stands.* Remark 9.3 reports reconstruction as not cheap, so nothing above forecasts an event. *Shown contingent* and *actually disturbed* are different claims, and only the first is made.

What ought to follow is not addressed, for the reason already given. This account prices operations and does not license them, and it has no more standing to design institutions than to design ethics.

Remark 9.5 (The same structure at another frame, and what it costs to hold). Persons are one frame. *Intellectual property is another, and what happens there is not analogous to the case above but identical to it.*

Copyright and patent encode a fact about presses, carriers and factories — that reproduction is expensive, so a copy competes materially with its origin. Neither is a fact about what is reproduced. When the marginal cost of a copy fell to approximately nothing the encoded fact stopped obtaining, and the word carrying most of the weight in the dispute records the import exactly: *piracy* names a taking after which the owner no longer has the thing, and is applied to an operation after which the owner still does. That is the same category import as *individual*, a

rivalrous term describing a non-rivalrous act. The term was apt when it was coined, because one press did compete materially with another; *what changed is not the word but the substrate under it.*

The identity with the previous remark is stronger than a parallel. Copyright's load-bearing distinction is between an expression, which is protected, and an idea, which is not — and that line holds only while the generator cannot be induced from what it generates. *Inducing a generator from observed outputs is Remark 9.3's operation*, run against a corpus rather than a person. The two cases are therefore one operation at two frames, which is what makes this a consequence of the account rather than a comparison imported into it.

None of this is new either. That information is non-rival was formalised in economics decades ago (Arrow, 1962; Romer, 1990), and the institutional argument has its own literature (Boyle, 2003). The account adds no observation to any of it. What it adds is that this case and the person case are the same error — a substrate-local cost read as a natural fact — and a unification is cheap to state and should be priced that way.

One correction, because the obvious conclusion is too strong. Artificial scarcity is often said to be unsustainable, and enforcement against copying does have the structure that argument needs: its cost scales with the number of agents and the ease of the act, both of which the substrate keeps increasing, while what is protected does not grow to match. But enforcement is not the only move. *The scarcity can instead be restored in the substrate* — sealed hardware, attestation, a carrier that refuses — and then the cost stops scaling with copiers and the arrangement is sustainable indefinitely. What it is not is free. Restoring a scarcity means deliberately removing reachable deformation, which is the Narrativium companion's quantity, and charging that permanently against the capability that made the thing worth protecting in the first place. *So the honest form is not that artificial scarcity must fail, but that it is the deliberate purchase of the worse operation* — paying transplant prices, in the terms of Remark 8.2, for something the substrate now copies. Whether that is worth paying is a question this account states and does not answer.

Remark 9.6 (What makes them one case: a rule stated from nowhere)*.* The two cases share more than a shape, and naming it corrects the diagnosis just given. *Reading a substrate-local cost as a natural fact is the symptom. The condition is a rule asserted with no frame at all* — and about those the account has something specific and unflattering to say.

A rule declaring itself absolute declares that no frame supplies it, which is the view from nowhere Axiom −1 refuses. The consequence is not that the rule is false. It is that *the rule has deleted its own cheapest repair.* The theory companion establishes that a frame is an instantiated boundary relation and that revoking one is the container ceasing to instantiate it — an operation the Narrativium companion measures occurring at $\Delta g = 0$, nothing in the contained state being required to change. *A rule that carries its frame can therefore be dropped for free.* A rule claiming no container has nothing that could cease to instantiate it, so that exit does not exist for it, and every remaining route to change runs through the contained state, which is the expensive one.

Absoluteness is therefore not a rule's strength but the removal of its exit, which inverts the usual reading, in which stating a rule without exception is what makes it robust.

The hardening follows, and it is monotone. A frame moves and mismatch accumulates; because the rule admits no outside, each accommodation is made inside it — an exception, a patch, an enforcement apparatus — and every one of those is structure that a later exit must also dismantle. So *the cost of leaving is non-decreasing while the pressure to leave is not.* That is a basin, and it is derived rather than observed.

One correction, because "and then it fails" claims more than the account supports. What is foreclosed is *self-originated* change: the companion's proposition is that a graph's admissibility constrains its own transformations and gives it no veto over world-caused change arriving through the boundary. The exit is never removed, then — only made unavailable to the party holding the

rule — so the correction still arrives, from outside, at once, as accumulated mismatch discharging rather than being paid down. *Such a system does not escape choosing. It forgoes choosing the timing and the size.*

Two limits. *Fixity is not the defect.* A rule held fixed inside a declared frame stays repairable by revoking the frame, at the price named above; what costs is fixity with the frame omitted, and the remedy is to state it, which is cheap. And *none of this is new.* Requisite variety (Ashby, 1956) is the same theorem in cybernetic vocabulary — a regulator needs variety matching its disturbance, and an absolute rule reduces its own while the disturbance does not oblige — while lock-in (Arthur, 1989) and institutional sclerosis (Olson, 1982) are the economics of the same basin. What the account contributes is the mechanism in its own terms, and the price. Not the observation.

But an explicit mechanism buys something a verdict does not: constructibility. Requisite variety returns an inequality — regulation fails once the regulator's variety falls below the disturbance's — and an inequality cannot be indexed by horizon or resolution, nor say which side moved. The mechanism above is stated in exactly the two families a Seldon manifold is assembled from, the admissible intervention sequences a frame carries and the future boundary conditions it represents, and it says which of them narrows and by what operation: *self-originated interventions shrink monotonically as accommodations accumulate, while the boundary-condition family does not shrink with them.* That is the input such a construction takes. *So the failure can be projected as a shape over a horizon rather than announced as an outcome*, which is the difference the companion on manifolds exists to make.

Its guards apply unrelaxed. *A manifold is reachability and not point prediction* — the section of that paper carrying the name — and resolution falls as the horizon lengthens, so what is projected coarsens rather than sharpens. Horizon is plural and frame-relative, so there is no manifold of the future, only one indexed to a stated frame and deformation class. What a frame does not represent is a blind spot, so the narrowing is visible only along dimensions already carried. *And the mechanism gives no date.* It closes the self-originated exit and leaves world-forcing untouched, and when the world forces is a fact about the world — so the tube narrows predictably and the discharge does not. *The projection bounds what remains reachable; it does not name what happens*, and reading the second off the first is the error this corpus spends most of its length refusing.

Remark 9.7 (What such a projection shows, and the sign of its error)*.* Having said the failure can be projected, the account owes what a projection built on these observations alone shows, and exactness about which parts of it are contestable.

The construction is contestable and will stay that way. A manifold is indexed to a frame, a deformation class, a horizon and a resolution, every one of which is an input. Disagreement about them is not a defect awaiting resolution but what Axiom −1 predicts, and no amount of care turns a frame-relative construction into a frame-free one. *What is not contestable in the same way is the sign of what the construction does*, and that is the part worth stating.

Two movements run opposite. The rules examined above narrow the self-originated exits available inside them, monotonically, by the mechanism given. *The substrate does the reverse*: an addressable held set makes the payment point a free parameter, and reconstruction from observation moves an ontological barrier to a sampling one. So what changes is not a quantity within a fixed space of options but *which options the space contains* — and a change in the dimension set is what *phase change* names here. That is a technical description, not an emphasis.

Which is exactly what a manifold built today cannot show. A manifold represents the dimensions its frame carries, at its resolution, so a dimension not yet represented is a blind spot by construction and a projection over today's dimensions cannot contain a change consisting in the set growing. That is not a reason to distrust the projection but a reason to know *which way it is wrong*, and here the error is signed: new dimensions add distinguishable reachable states rather than removing them, so a projection over a fixed set *understates* how much reachability

changes, and any horizon read off one is correspondingly *late*.

One thing this does not license, and one it does. A signed bias is a direction and not a distribution: it says the estimate errs late, not by how much, and converting that into a probability would be reading a number off a shape, which is the standing prohibition here.

What it does license is a window, once priors are imported and named. It is tempting to refuse one on the ground that none is derivable from these observations. That is true, trivial, and the wrong test — *nothing* is derivable from observations alone, and a rate owing nothing to priors is the prior-free baseline Proposition 7.1 places outside the space, so refusing on that ground commits the error this paper spends a section diagnosing. *The window is not unavailable. It is unmeasured*, which is a smaller and a different claim.

Four quantities fix it and each is measurable in principle: the historical rate at which the represented dimension set has grown; the cost curve of inducing generators from observed outputs; the budget available for restoring a scarcity in the substrate; and the rate at which accommodations accumulate inside the rules concerned. Nothing here assembles them. *Assembling them is the work this account exists to make possible*, and that is the whole of what is claimed — not the number, but that the question is now well-posed enough to have one. The near end is fixed by observation rather than by any of the four: reproduction at negligible marginal cost has already happened for information goods, and induction of generators from observed outputs is already performed at scale, so part of the change is not pending but done.

On the present author's reading of those four, the reachability set changes drastically within roughly the next half-century, more probably in its earlier part, the signed error above being one reason for the weighting. *That is a judgement with named inputs and is recorded as one.* Its worth is not its number but its form: a reader who disagrees can now say *which of the four* they read differently, which is the difference between a disagreement that can be audited and one that can only be repeated.

10 Prior positions

None of the machinery here is new and it should not be dressed as though it were. Four literatures own it, and what this account adds is an arrangement rather than an operation.

The relativity of measurement, which is old and is not claimed here. Everything Section 7 says about scores has been said, in some cases seventy years ago, and it should be read as inherited rather than offered.

Generalizability theory decomposes a measurement into facets — the person, the item, the occasion, the observer — and holds that a score is a *sample from a universe of admissible observations* rather than a property of whoever was measured (Cronbach et al., 1972). That is the four-place arrangement, in a literature that had it before this programme existed. Cronbach and Meehl (1955) had already established that a test measures a construct only under an argument and never by fiat. Quine (1951) gives the inference-licensing half: no hypothesis meets the evidence alone, so what a result licenses depends on auxiliaries that travel with it. And Newell (1973) made the corresponding objection to task-by-task measurement in this field specifically — that a pile of task results does not accumulate into an account of the mechanism.

One thing differs, and it is small and different in kind. Those are cautions, and cautions are discharged by better practice: estimate more facets, state the auxiliaries, argue the construct. What Proposition 7.1 adds is that the prior-free case is not a facet that has been estimated badly — *the origin is not in the space.* Generalizability theory supplies the coordinates and assumes a universe of admissible observations within which they vary; the proposition denies there is a zero for one of them to be measured from.

The practical force of that difference is the whole of it. A caution tells a careful practitioner to do more work; *a consequence tells them the work does not exist to be done*, and stops a research programme from being organised around obtaining a control that cannot be built. Nothing else

here is new, and the derivation is the contribution rather than the observation it derives.

Anti-unification. The operation is Plotkin (1970)'s least general generalisation: given two structures, produce the most specific structure of which both are instances. That is exactly rung two and rung three read backwards. *What differs is what is kept.* The classical formulation returns the generalisation and discards the substitutions that recover the inputs from it; here the substitutions are the point, because *the witness is the transform* and applying it elsewhere is the whole operation. An lgg that throws away its witness cannot adapt anything.

Structure-mapping. Gentner (1983) established that analogy runs on relational structure rather than surface features, and that systematic relational systems are preferred to isolated matches — both of which this account assumes rather than argues. *What differs is the purpose.* Structure-mapping is a theory of comparison: what makes two things analogous and how a person aligns them. This is a theory of acquisition, in which the comparison is instrumental — the alignment exists to yield a transform that mints a competence the system did not have. Gentner's theory answers what an analogy *is*; the question here is what one is *for*.

Analogical proportion. The four-term form, a is to b as c is to d, has a developed logical treatment (Prade and Richard, 2014), and parallelogram completion is that form with d unknown. Taken directly and without modification.

Embedding arithmetic, and the reason its critique does not transfer. The word *parallelogram* comes from vector offsets in learned embedding spaces (Mikolov et al., 2013), and that method's evaluation has been substantially criticised: Levy and Goldberg (2014) showed its behaviour is largely accounted for by similarity terms rather than by the offset, and Linzen (2016) showed the standard benchmark is confounded by nearest-neighbour baselines that need no analogy at all.

Those criticisms are correct and they do not apply here, for a reason that is a real difference and not a defence. The offset method is metric: it returns the nearest point to a computed location, so it always returns something, and the question is only whether the something is right. The operation here is symbolic and exact — anti-unification either yields a witness that is well-defined and lossless, or it fails and yields nothing. *There is no nearest neighbour to fall back on and no threshold to tune*, which removes the failure mode the critique identifies and replaces it with a different one: refusing cases a metric method would have answered. That is the cost of the exchange and it is the honest form of it.

11 Falsifiability commitments

1. [SHAPE] **Shape matching and analogy are one operation.** Fails if a case of one cannot be expressed as a case of the other — specifically, if some analogical completion has no formulation as a transform between held signatures, or some shape adaptation has no formulation as an anti-unification with its witness retained.

2. [SHAPE] **The ladder's order is forced, not chosen.** Fails if reordering the rungs leaves the outcomes unchanged on cases where the rungs disagree, which would make the order a preference.

3. [DETAIL] **One mint, then recognition.** Fails on a novel case that yields repeated minting under an unchanged codebook, or none at all.

4. [DETAIL] **Coupling forces a dimension.** Fails if a residual coupled across observed dimensions is absorbed by a member rather than a dimension, without loss.

5. [SHAPE] **Adaptability is bounded by the priors.** Fails on a demonstration of a competence derived from a held set containing no path to it — which would make adaptation a source of capability rather than a conversion of one, and would be a free lunch.

6. [SHAPE] **No prior-free baseline is constructible.** Fails on the exhibition of a frame holding nothing that nevertheless draws a distinction — which would put the zero of the priors axis inside the space and make prior-free measurement well-formed after all.

7. [DETAIL] **A prior-assuming benchmark measures a relation, not a property.** Fails if performance on such a task set is shown to be independent of familiarity with its conventions, in a population matched on everything else.

8. [DETAIL] **An oracle bought once is paid once.** Fails if held structure derived from an oracle requires continued access to that oracle to keep working — if the purchase does not in fact terminate, the distinction drawn here is between two ways of paying continuously and is not worth drawing.

9. [SHAPE] **Capability has plural sources, selected by regime.** Fails if a domain is exhibited in which candidates are cheaply and exactly checkable and held structure nevertheless contributes measurably, or in which no oracle exists and holding contributes nothing — either would break the inverse relation this account asserts between the value of holding and the cheapness of verification.

10. [SHAPE] **Adaptability is per case, not a rank.** Fails if a single ordering of systems by adaptability can be constructed that survives changing the held set.

References

Kenneth J. Arrow. Economic welfare and the allocation of resources for invention. In *The Rate and Direction of Inventive Activity: Economic and Social Factors*, pages 609–626. Princeton University Press, 1962.

W. Brian Arthur. Competing technologies, increasing returns, and lock-in by historical events. *The Economic Journal*, 99(394):116–131, 1989.

W. Ross Ashby. *An Introduction to Cybernetics*. Chapman and Hall, London, 1956.

James Boyle. The second enclosure movement and the construction of the public domain. *Law and Contemporary Problems*, 66(1–2):33–74, 2003.

François Chollet. On the measure of intelligence. arXiv:1911.01547, 2019.

Lee J. Cronbach and Paul E. Meehl. Construct validity in psychological tests. *Psychological Bulletin*, 52(4):281–302, 1955.

Lee J. Cronbach, Goldine C. Gleser, Harinder Nanda, and Nageswari Rajaratnam. *The Dependability of Behavioral Measurements: Theory of Generalizability for Scores and Profiles*. Wiley, New York, 1972.

Dedre Gentner. Structure-mapping: A theoretical framework for analogy. *Cognitive Science*, 7(2):155–170, 1983.

Omer Levy and Yoav Goldberg. Linguistic regularities in sparse and explicit word representations. In *Proceedings of the Eighteenth Conference on Computational Natural Language Learning*, pages 171–180, 2014.

Tal Linzen. Issues in evaluating semantic spaces using word analogies. In *Proceedings of the 1st Workshop on Evaluating Vector-Space Representations for NLP*, pages 13–18, 2016.

Tomas Mikolov, Wen-tau Yih, and Geoffrey Zweig. Linguistic regularities in continuous space word representations. In *Proceedings of NAACL-HLT 2013*, pages 746–751, 2013.

Solim LeGris, Wai Keen Vong, Brenden M. Lake, and Todd M. Gureckis. A comprehensive behavioral dataset for the Abstraction and Reasoning Corpus. *Scientific Data*, 12, 2025.

Allen Newell. You can't play 20 questions with nature and win. In *Visual Information Processing*, pages 283–308. Academic Press, 1973.

Ned Block and Jerry A. Fodor. What psychological states are not. *The Philosophical Review*, 81(2):159–181, 1972.

Mancur Olson. *The Rise and Decline of Nations: Economic Growth, Stagflation, and Social Rigidities*. Yale University Press, 1982.

Gordon D. Plotkin. A note on inductive generalization. In *Machine Intelligence 5*, pages 153–163. Edinburgh University Press, 1970.

Hilary Putnam. Psychological predicates. In *Art, Mind, and Religion*, pages 37–48. University of Pittsburgh Press, 1967.

Willard van Orman Quine. Two dogmas of empiricism. *The Philosophical Review*, 60(1):20–43, 1951.

Paul M. Romer. Endogenous technological change. *Journal of Political Economy*, 98(5):S71–S102, 1990.

Elizabeth S. Spelke and Katherine D. Kinzler. Core knowledge. *Developmental Science*, 10(1):89–96, 2006.

Alan M. Turing. Computing machinery and intelligence. *Mind*, LIX(236):433–460, 1950.

David H. Wolpert and William G. Macready. No free lunch theorems for optimization. *IEEE Transactions on Evolutionary Computation*, 1(1):67–82, 1997.

Henri Prade and Gilles Richard. From analogical proportion to logical proportions. *Studies in Computational Intelligence*, pages 217–244, 2014.